



# Virtual and Augmented Reality


CS-GY 9223/CUSP-GX 6004

<https://nyu-icl.github.io/courses/2022fall-vr-ar>

Prof. Qi Sun  
qisun@nyu.edu | [www.qisun.me](http://www.qisun.me)

# Important Logistics for Nov 7

## 1. Remote guest lecture




Josef Spjut, Research Scientist, NVIDIA Research

<https://research.nvidia.com/person/josef-spjut>



## 2. Workshop/tutorial for ARCore/ARKit development

# Eye Tracking



# Eye Tracking

- necessary for gaze-contingency paradigm:
- foveated rendering, gaze-contingent rendering, gaze-contingent focus...
- interaction
- eye contact
- ...


# Eye Tracking

- many different techniques:
  - electro-oculography
  - contact lens tracking
  - video-oculography
  - pupil / corneal reflection tracking
  - dual Purkinje image




# Eye Tracking – Pupil / Corneal Reflection

- corneal reflection stays constant, pupil center moves relative!




# Eye Tracking – Pupil / Corneal Reflection

- corneal reflection stays constant, pupil center moves relative!



# Eye Tracking

- where am I looking? what am I looking at?



<http://www.getbusymedia.com/small-business-insights-eye-tracking/>



# Latency

- min acceptable: 20 ms
- interactive applications <20 ms  
(say target is 5 ms)

*The latency between the physical movement of a user's head and updated photons from a head mounted display reaching their eyes is one of the most critical factors in providing a high quality experience.*

- John Carmack

# Latency – where does it come from?


- IMU ~1 ms
- sensor fusion, data transfer
- rendering: depends on complexity of scene & GPU – a few ms
- data transfer again
- display: LCD ~60 Hz = 16 ms; OLED <1 ms

# Latency – how bad is it really?


- example:
  - 16 ms (display) + 16 ms (rendering) + 4 ms (orientation tracking) = 36 ms latency total
  - head rotates at 60 degrees / sec (relatively slow)
  - 1Kx1K display over 100 degrees field of view
- in 36 ms, my head moved 1.92 deg ~ 19 pixels = size of thumb at arm's length! too much

# Display Pixel Updates


Raster Scan  
(e.g. electron  
beam in CRT)



Rolling Update  
(most LCDs)




Global Update  
(some LCoS,  
DLP, other)




# Display Pixel Switching - Persistence

- after the display pixel switched states, how long is it on?


Full Persistence  
(most LCDs)



Half Persistence  
(strobe backlight)




Low Persistence  
(OLED, strobing )




example: switch from white to black to white to black as fast as possible




# Rapid relative motion




# Space-time diagrams (static)




# Spatial movement over time




# Pixel-based movement




# Sequential RGB display









# Sequential RGB with eyes moving






# Full persistence




# Half persistence






# Zero persistence



# Full persistence + head rotation






# Post-rendering Image Warp

- also called time-warp by John Carmack
- minimize end-to-end latency
- original paper from Mark et al. 1997, also Darsa et al. 1997
- overview:
  1. get orientation from IMU, perhaps also position
  2. render scene into off-screen buffer (larger than screen)
  3. read latest orientation from IMU
  4. warp rendered image with latest orientation
- 2D image translation v 2D image warp v 3D image warp

# End-to-End Low-Latency Optical See-Through AR Pipeline


---





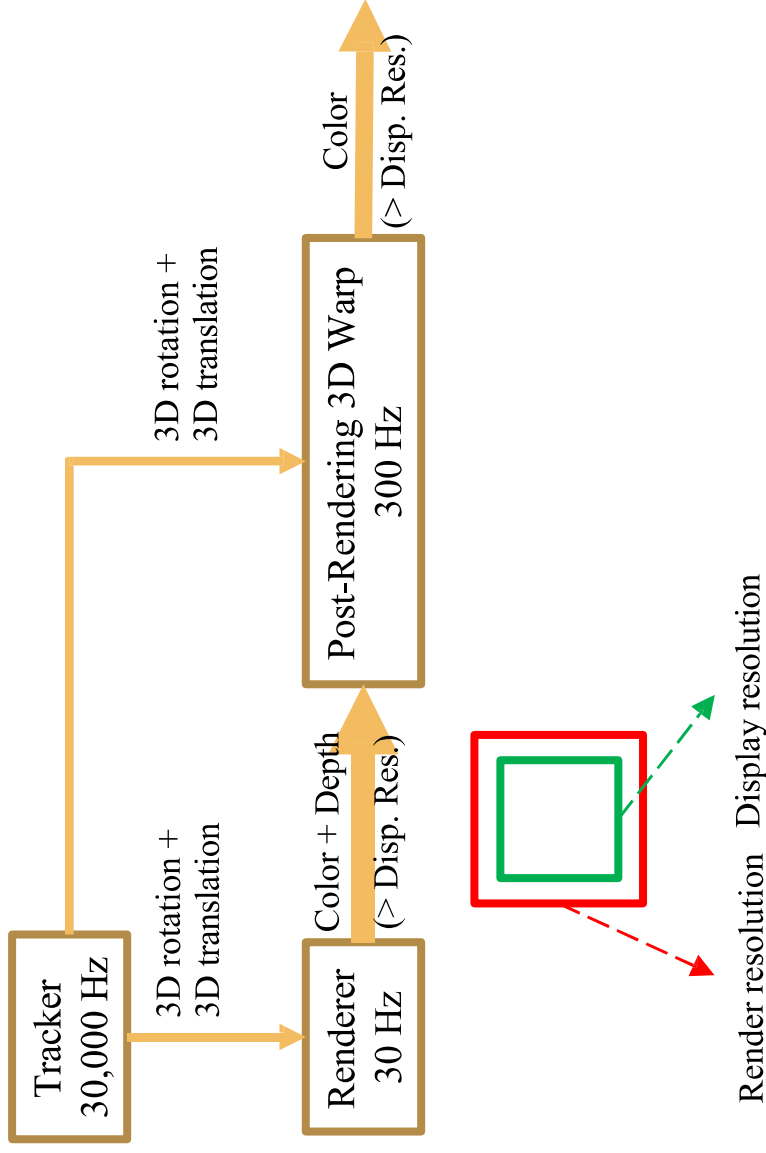
# End-to-End Low-Latency Optical See-Through AR Pipeline

---

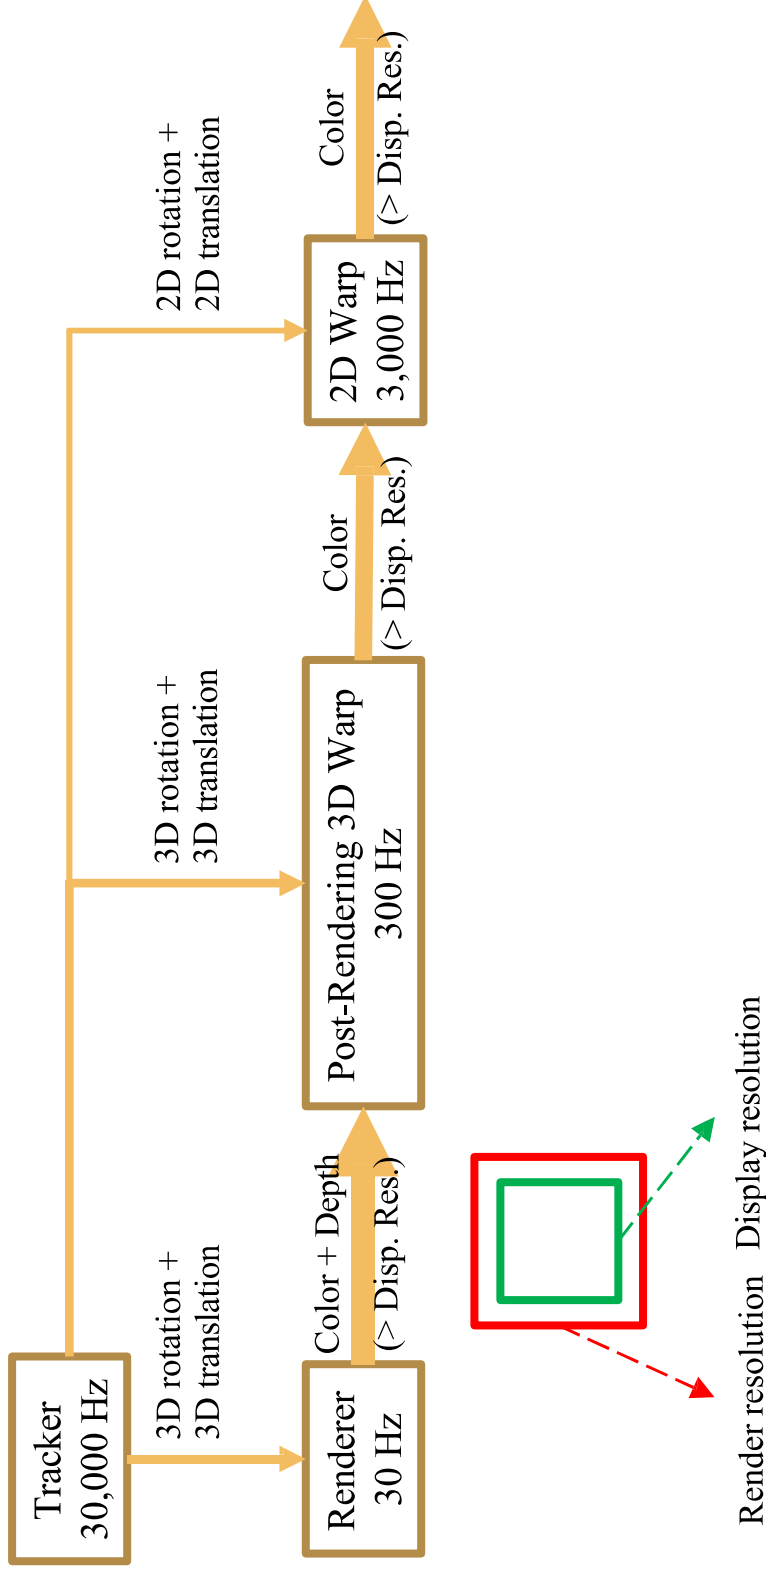


# End-to-End Low-Latency Optical See-Through AR Pipeline

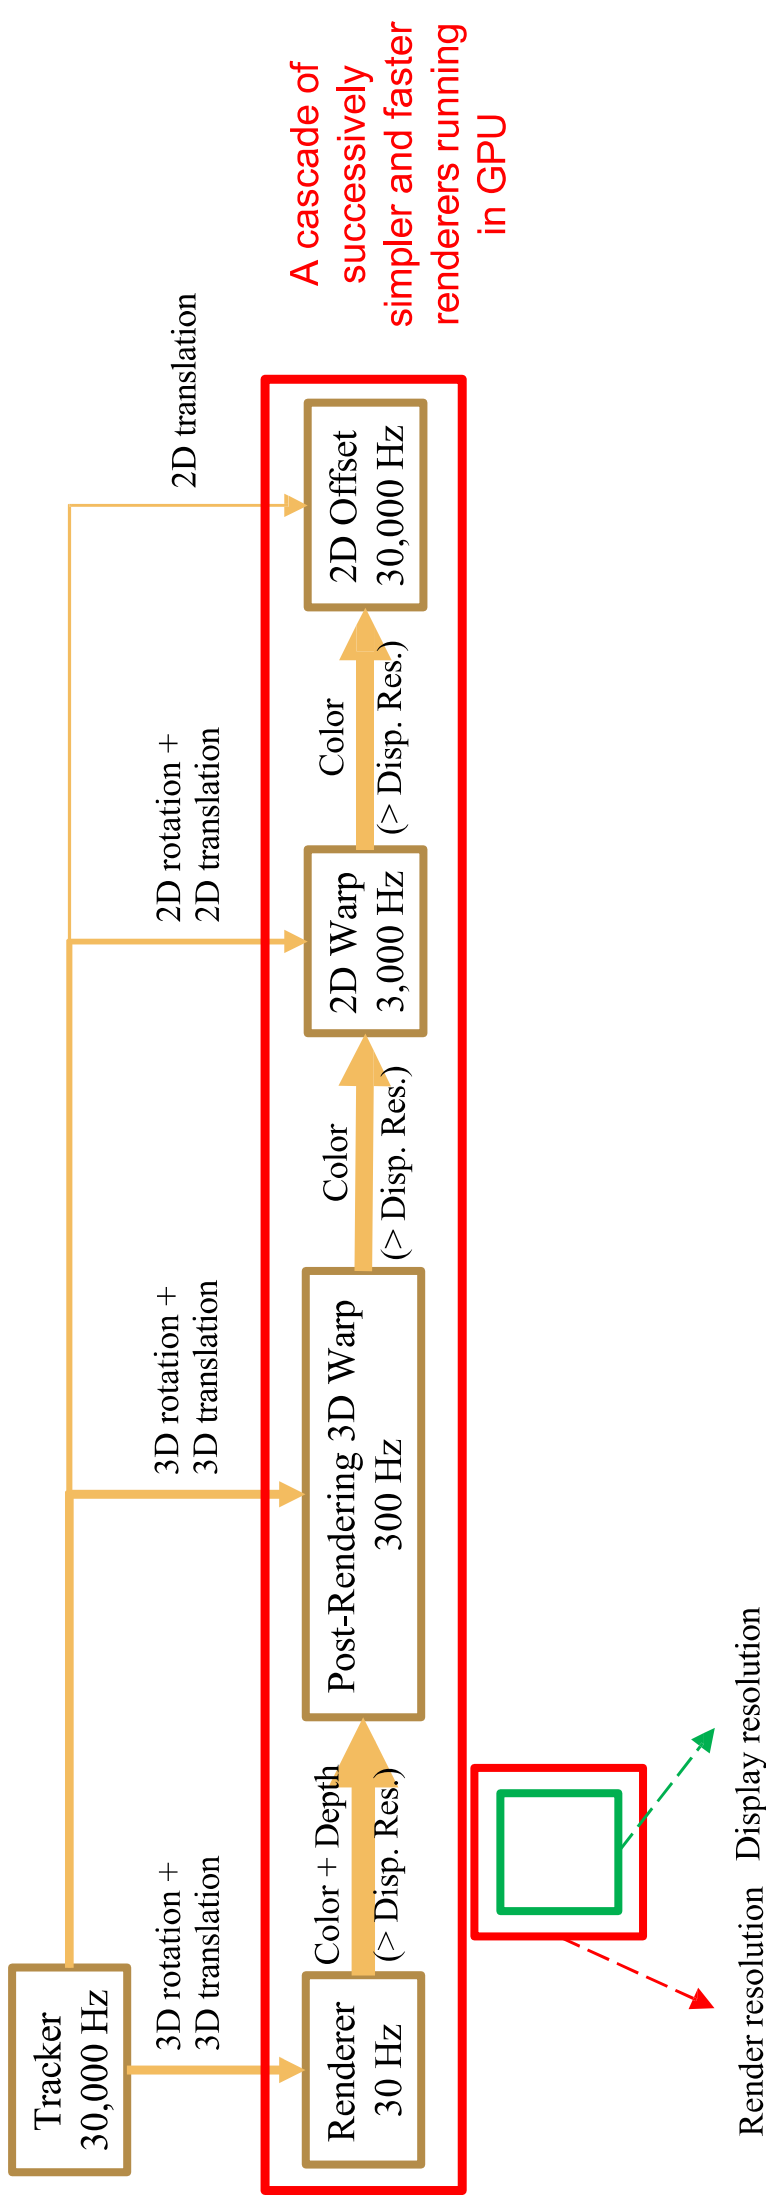
---



# End-to-End Low-Latency Optical See-Through AR Pipeline




# End-to-End Low-Latency Optical See-Through AR Pipeline



# Summary: Latency, Persistence, etc.

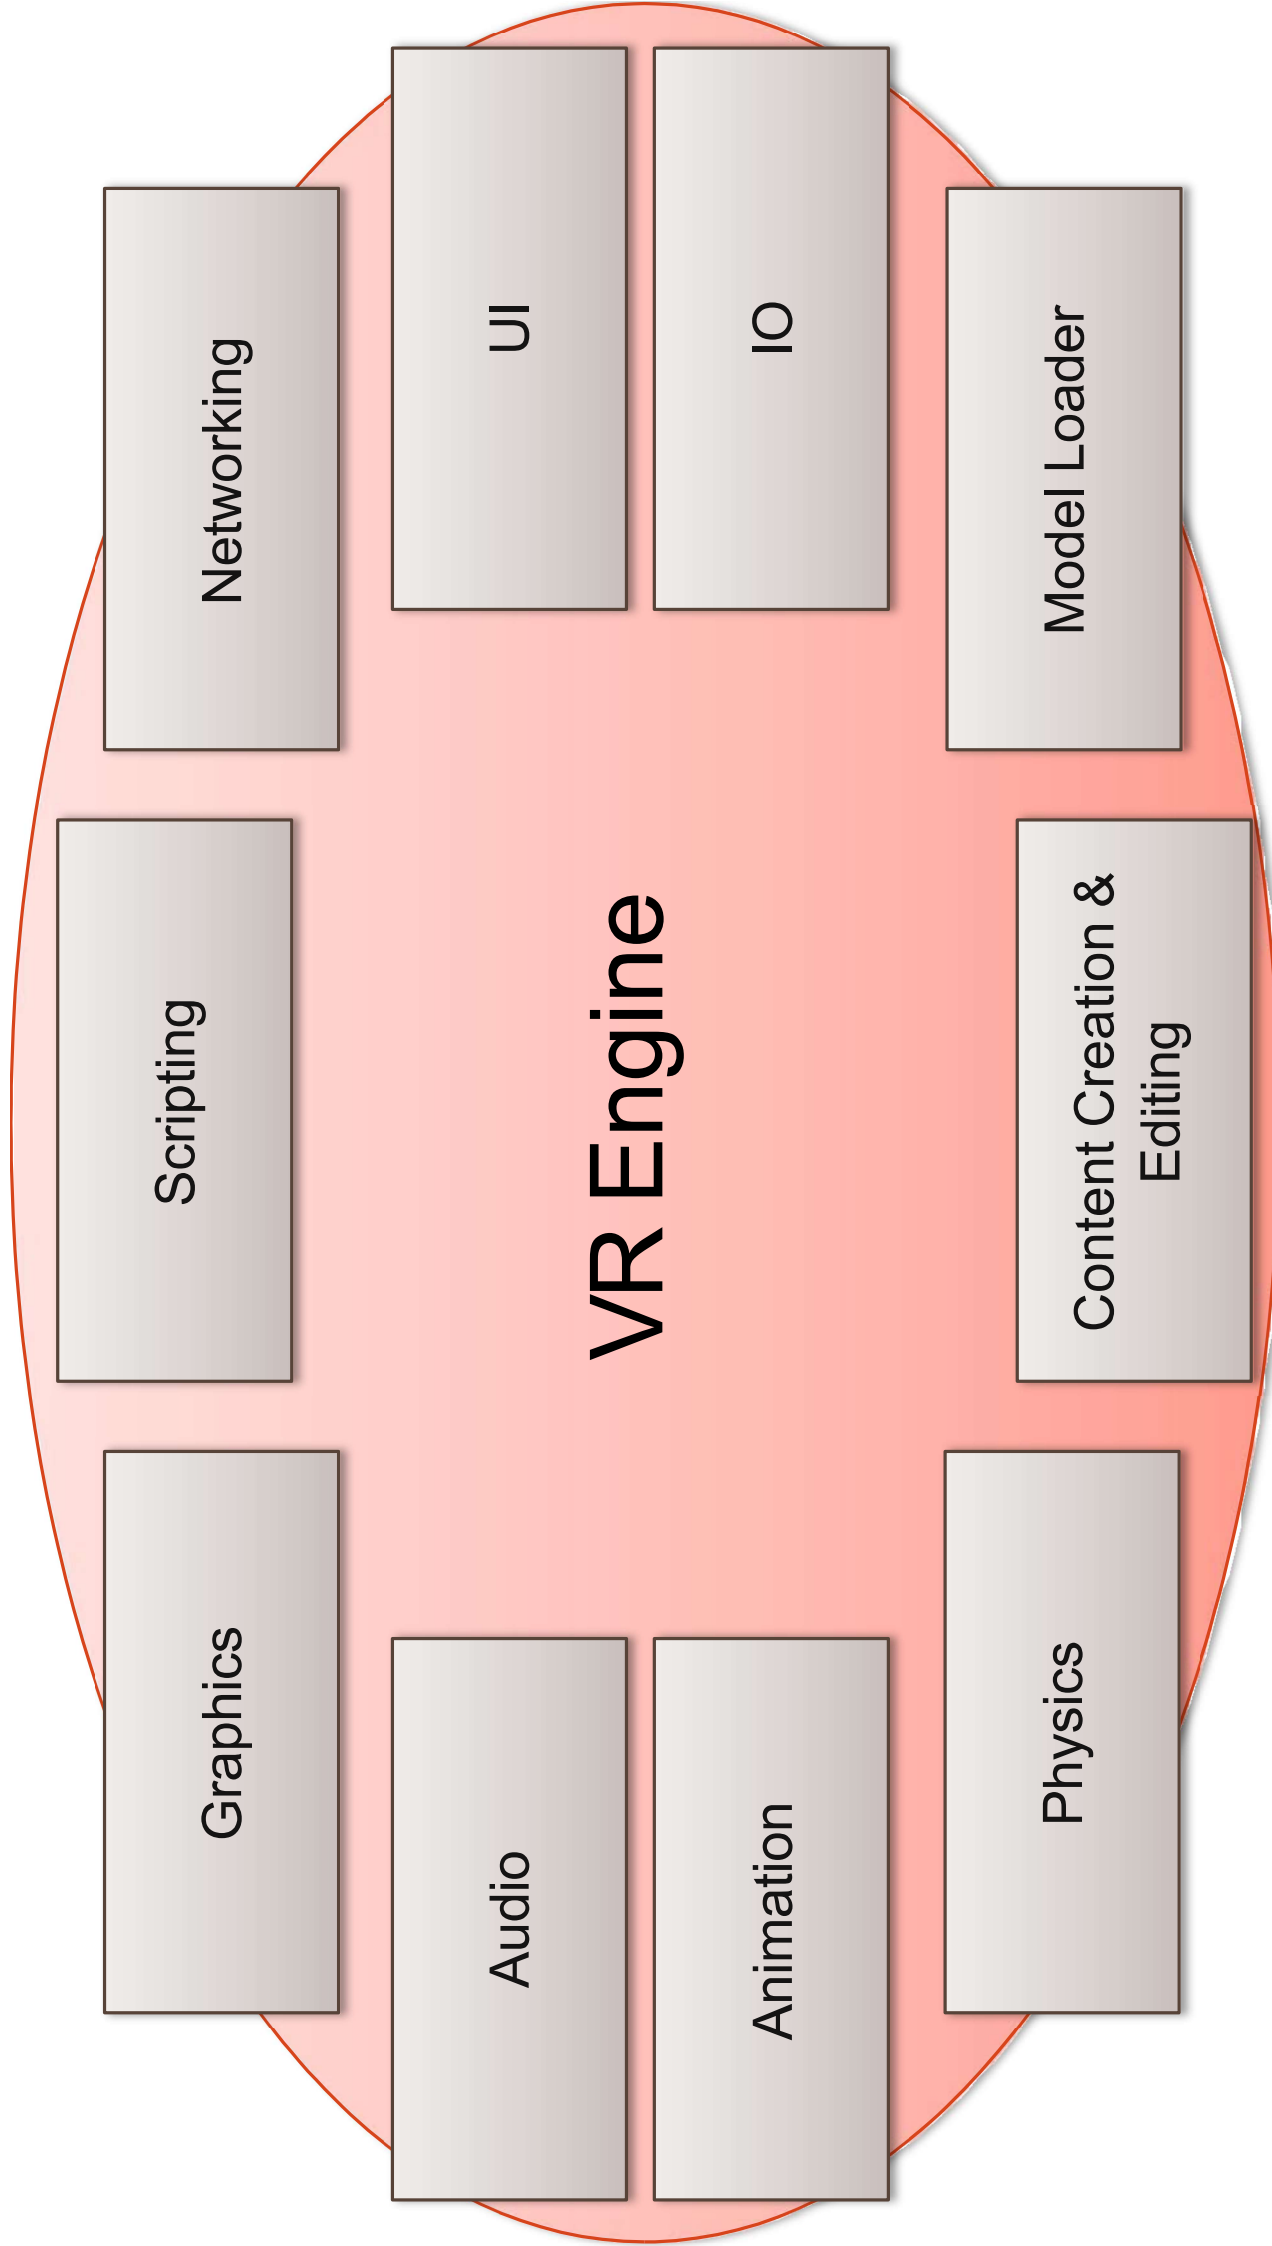
- predictive tracking (e.g. LaValle ICRA 2014)
- post-rendering warp
- design and build really great hardware & algorithms
- use OLED displays or strobing backlights for low persistence
- design some type of a device to actually measure latency!



VR Engine  
(so far)

Graphics

IO



# VR Engines - Audio

- middleware – between audio card and application (e.g. game)
- usually provides functionality for:
  - loading different types of sound files
  - mixing and mastering
  - 3D sound
  - occlusions, echoes, reverberation, ...

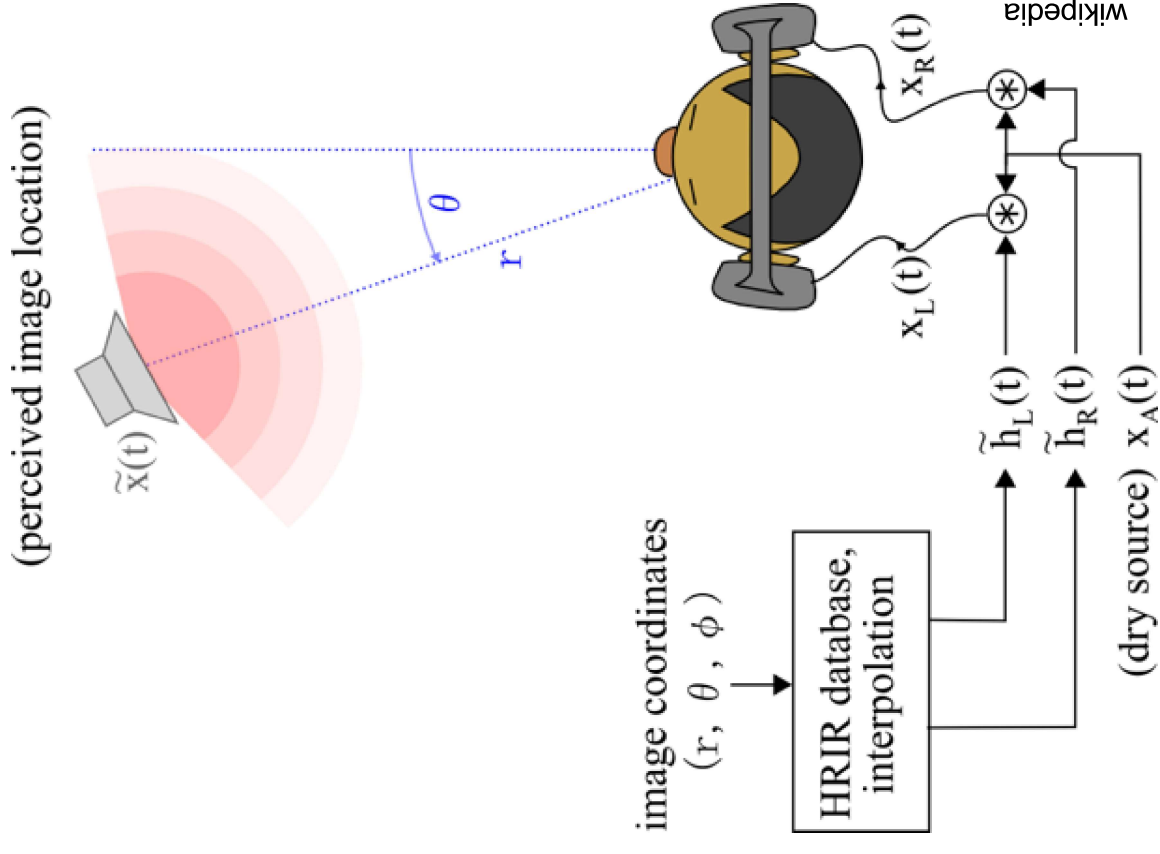


# VR Engines - Audio

- examples:
  - Google Resonance Audio <https://resonance-audio.github.io/resonance-audio/>
  - FMOD - [www.fmod.org](http://www.fmod.org)
  - OpenAL - “OpenAL for sound”
  - SDL – provides basic functionality

# VR Engines – 3D Audio

- start with mono sound  $x_A(t)$
- head-related impulse response (HRIR) model time delay and attenuation via convolution
- basically different temporal shift for each ear
- but HRIR also includes other effects created by shape of ear and other factors



# VR Engines - Physics

- framework to simulate:
  - rigid body dynamics (e.g. collision detection)
  - soft body dynamics (e.g. deformation, cloth, ...)
  - fluid dynamics (water, smoke, fire, ...)

# VR Engines - Physics

- examples:
  - Open Dynamics Engine (<http://www.ode.org/>): free © but limited to rigid body dynamics & collision
  - Bullet Physics (<http://bulletphysics.org/>): free ©, rigid & soft body dynamics, widely used
  - havok (owned by Microsoft) – not free ® but widely used, real-time rigid body dynamics

# VR Engines – User Interface (UI)

- concept is straightforward: widgets, menus, buttons, checkboxes, ...
- types of UIs:
  - non-diegetic – lives in screen space (e.g. player status); doesn't work in VR (no screen space)

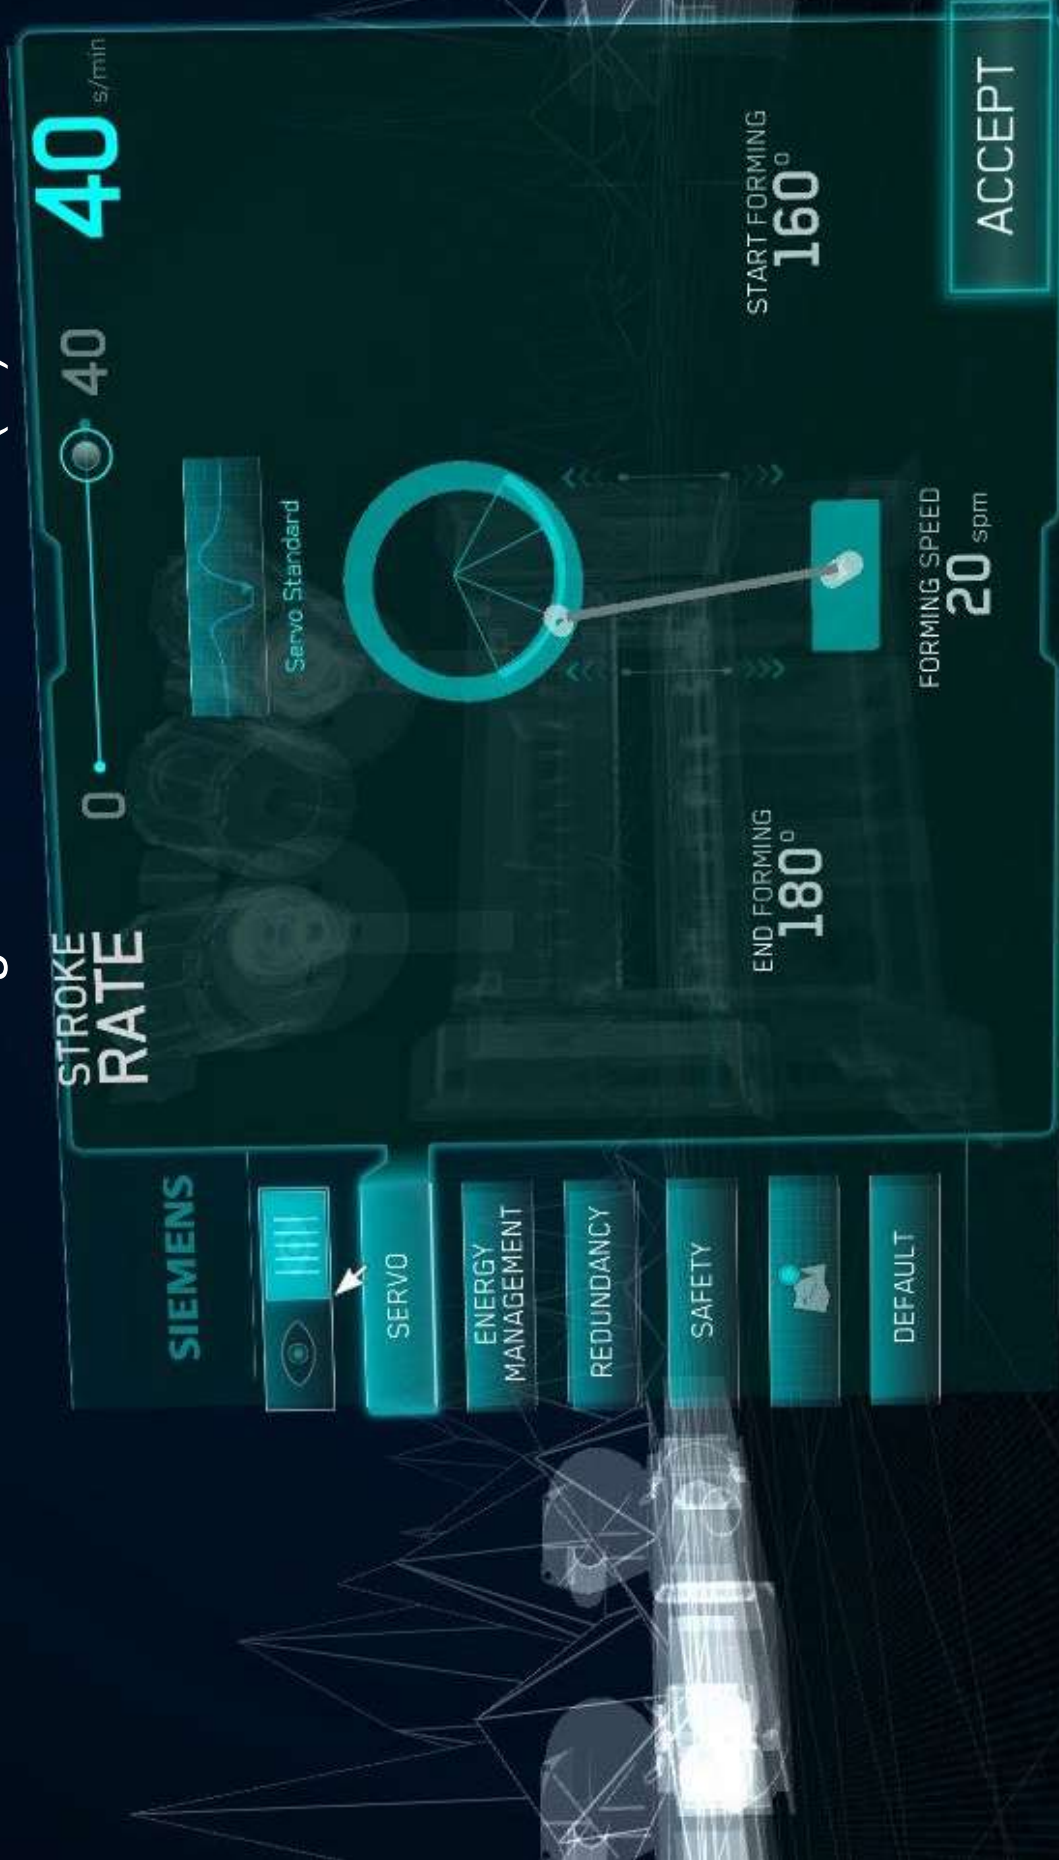


- spatial UI – lives in the virtual world

- diegetic – menus in world



## VR Engines – User Interface (UI)



# VR Engines - IO

- support for interfaces: keyboard, mouse, 3D mouse, standard haptic devices, ...
- VR engine would provide functionality as well (e.g. Unity)

# VR Engines – Content Creation

- 3D modeling programs / Computer-aided Design (CAD):
  - Maya (production)
  - 3ds Max (games)
  - Blender – free
  - SolidWorks – 3D printing & fabrication
  - Tinkercad: free & online



# VR Engines – Content Creation

- what's involved?
  - conceptual design
  - 3D modeling
  - animation and/or simulation
  - scripting behavior and artificial intelligence of characters
  - testing
  - ... many different stages in application development ...

# VR Engines - Scripting

- core engine is usually designed for performance – C++
- developing applications should be easy! the user almost never wants to touch the C++ source but needs flexibility
- provide a script-based interface to allow user to change anything they need for their application
  - create & manipulate objects
  - script behavior
  - change shaders (e.g. change camera or fragment shader art)

# VR Engines - Networking

- manage low-level communication protocols (TCP/IP, UDP, ...)
- ensure that character states, graphics, sound, and everything else is synchronized
- connect to application that's running as client
- network updates, messages, ...

# Popular VR/Game Engines

- Unity: cross-platform, Direct3D (Win), OpenGL (Mac & Linux), iOS & Android support, also came console APIs; personal license is free; seems to be the easiest to use so we'll use it for Lab 6 and HW 6
- Unreal: very popular, lots of awards, unreal engine 4 is free
- CryEngine: popular game engine, just announced support for VR; free for non-commercial use