



# Virtual and Augmented Reality

CS-GY 9223/CUSP-GX 6004

<https://nyu-icl.github.io/courses/2022fall-vr-ar>

Prof. Qi Sun

qisun@nyu.edu | www.qisun.me

# Final Project Presentation

2.5min in total (no more) =  
cover page +  
~.5min on motivation/problem +  
~1 min on system showcasing +  
~1 min on tech details

# Panoramic Imaging and Cinematic VR



Jaunt VR



Jaunt VR

Lytro



Lytro



Google



Nokia



W: 168,36mm / 6.7"

L: 262,95mm / 10.4"



W: 157,83mm / 6.3"

H: 262,95mm / 10.4"

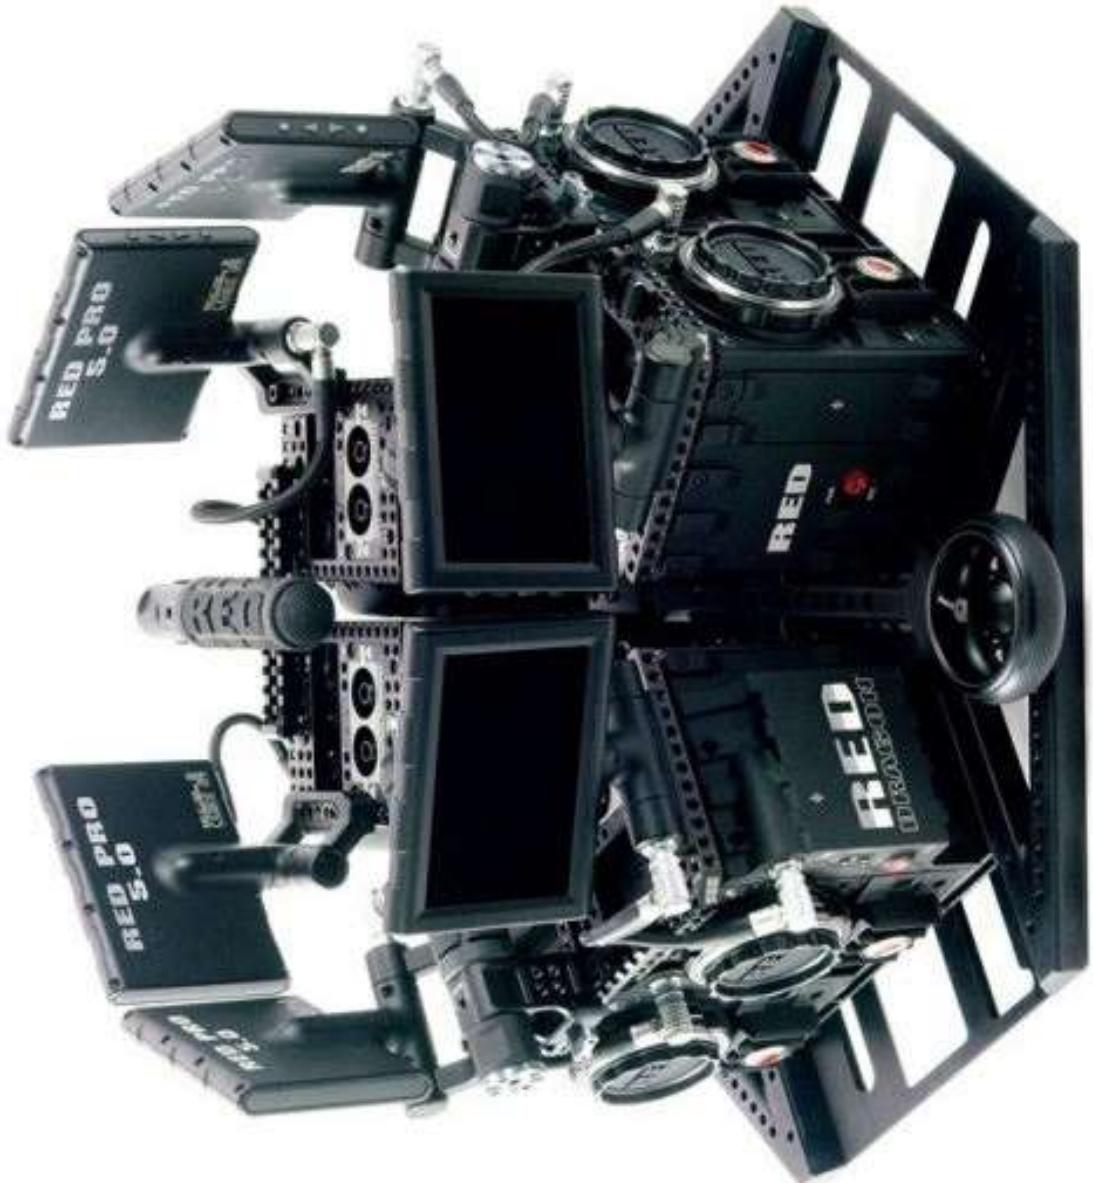
Facebook



see Brian Cabral's SCIENCE talk @talks.stanford.edu



Red



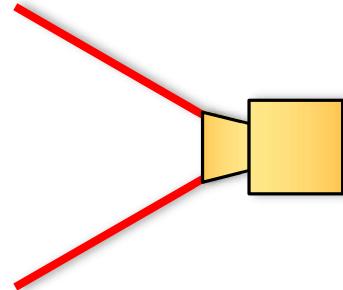
Samsung



# Panorama v Stereo Movie v Stereo Panorama

## Panorama

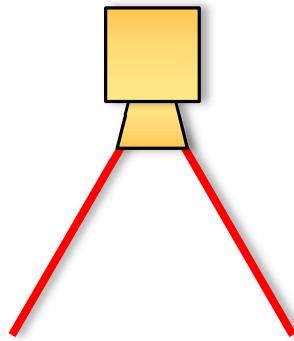
mono & head rotation



# Panorama v Stereo Movie v Stereo Panorama

Panorama

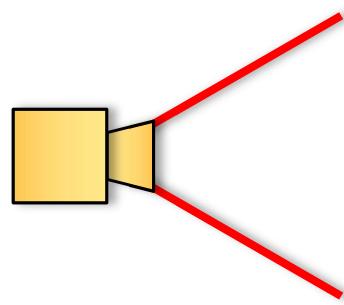
mono & head rotation



# Panorama v Stereo Movie v Stereo Panorama

## Panorama

mono & head rotation



# Panorama v Stereo Movie v Stereo Panorama

Panorama

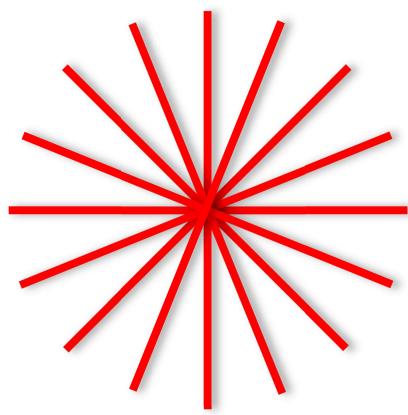
mono & head rotation



# Panorama v Stereo Movie v Stereo Panorama

## Panorama

mono & head rotation

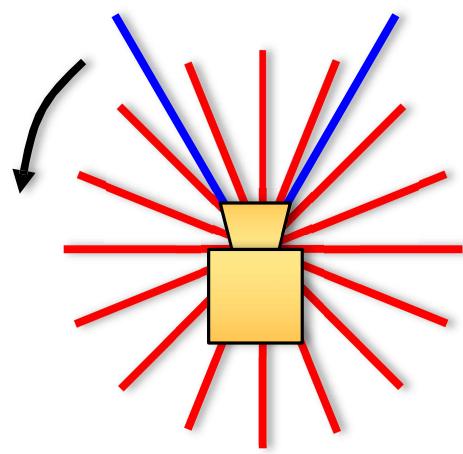


1 center of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

## Panorama

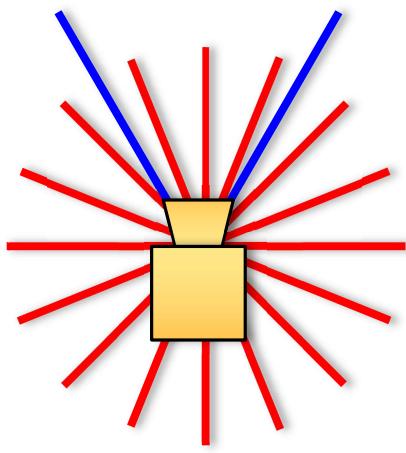
mono & head rotation



1 center of  
projection!

# Panorama v Stereo Movie v Panorama

Panorama  
mono & head rotation



1 center of  
projection!

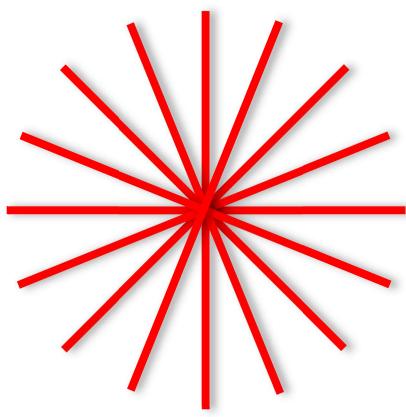
center of  
projection



# Panorama v Stereo Movie v Stereo Panorama

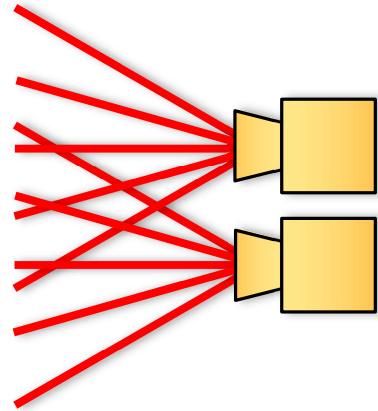
Panorama

mono & head rotation



Stereo

stereo & no head rotation



Stereo Panorama

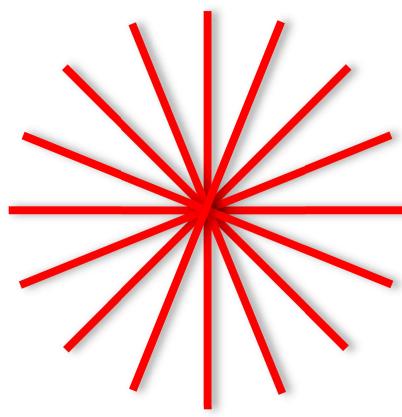
stereo & head rotation

1 center of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

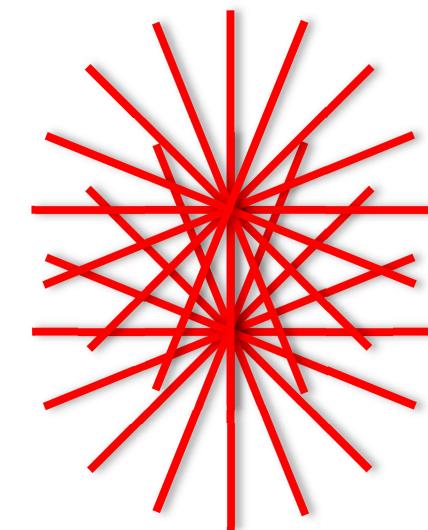
Panorama

mono & head rotation



Stereo

stereo & no head rotation



Stereo Panorama

stereo & head rotation

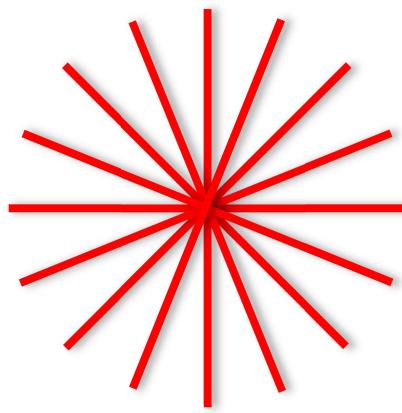
1 center of  
projection!

2 centers of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

Panorama

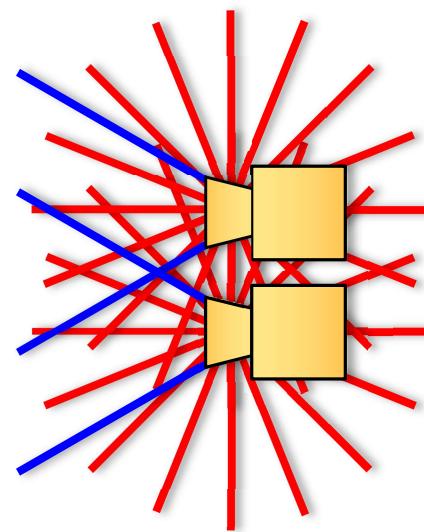
mono & head rotation



1 center of  
projection!

Stereo

stereo & no head rotation



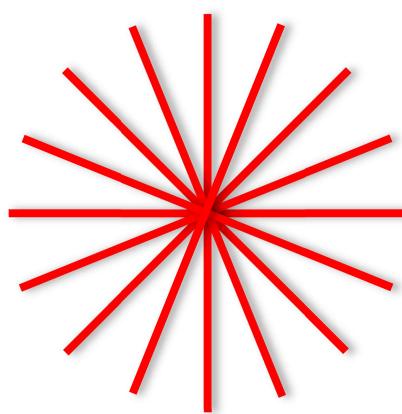
Stereo Panorama

stereo & head rotation

# Panorama v Stereo Movie v Stereo Panorama

Panorama

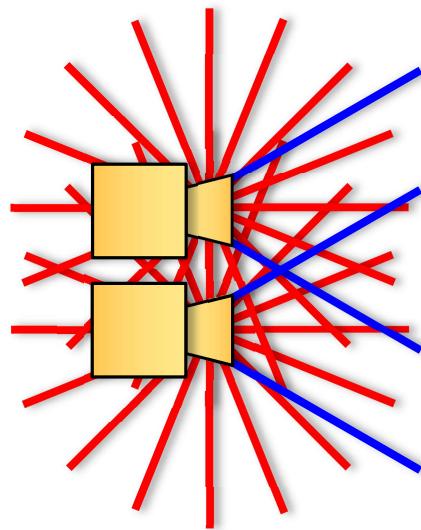
mono & head rotation



1 center of  
projection!

Stereo

stereo & no head rotation



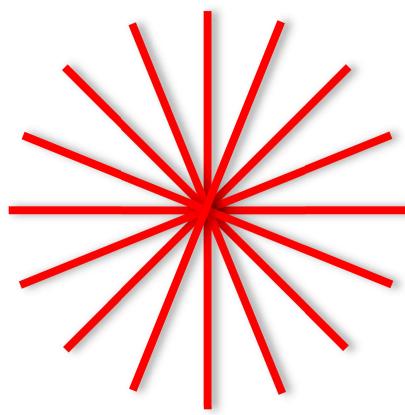
Stereo Panorama

stereo & head rotation

# Panorama v Stereo Movie v Stereo Panorama

Panorama

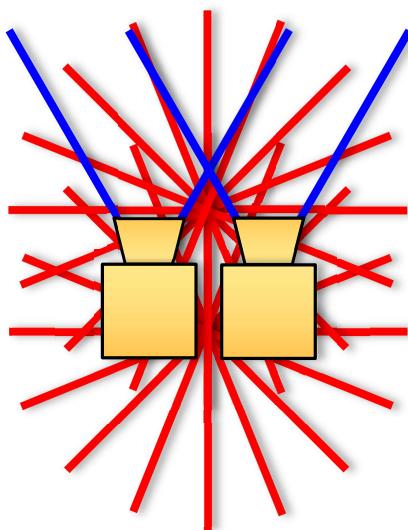
mono & head rotation



1 center of  
projection!

Stereo

stereo & no head rotation



Stereo Panorama

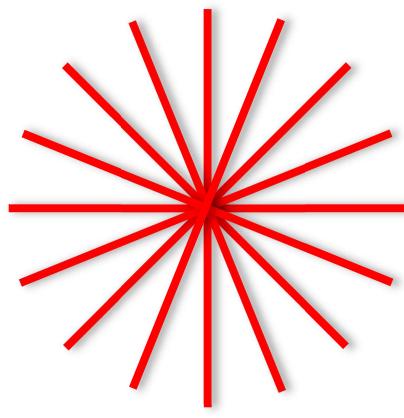
stereo & head rotation

2 centers of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

Panorama

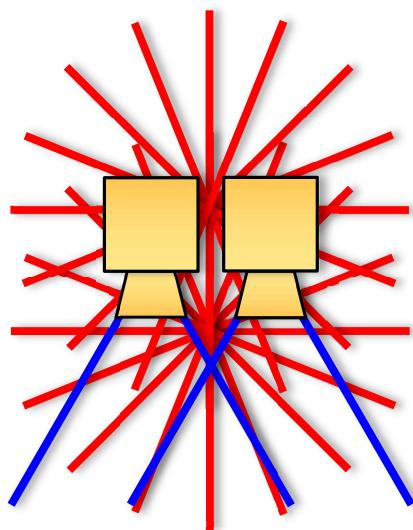
mono & head rotation



1 center of  
projection!

Stereo

stereo & no head rotation



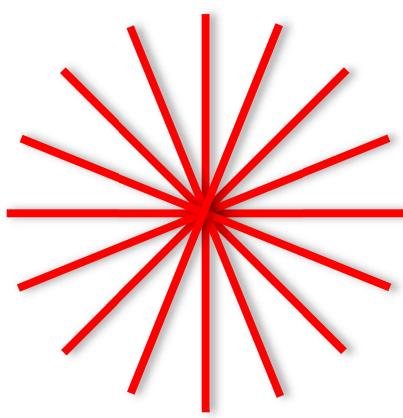
Stereo Panorama

stereo & head rotation

# Panorama v Stereo Movie v Stereo Panorama

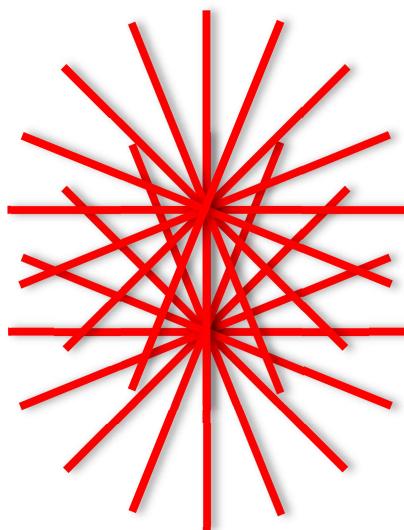
Panorama

mono & head rotation



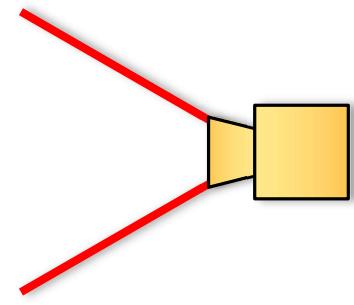
Stereo

stereo & no head rotation



Stereo Panorama

stereo & head rotation



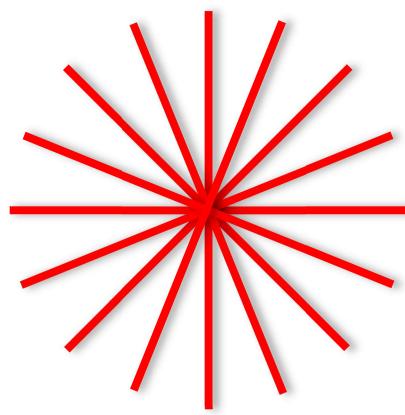
1 center of  
projection!

2 centers of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

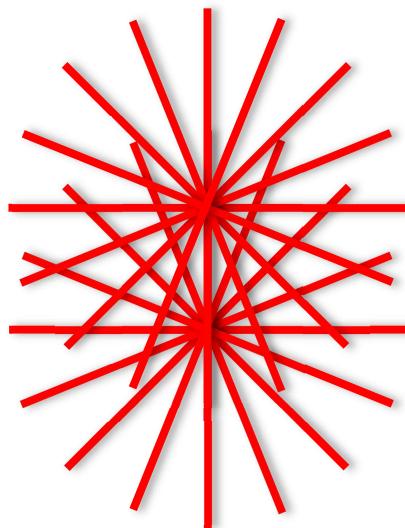
Panorama

mono & head rotation



Stereo

stereo & no head rotation



Stereo Panorama

stereo & head rotation



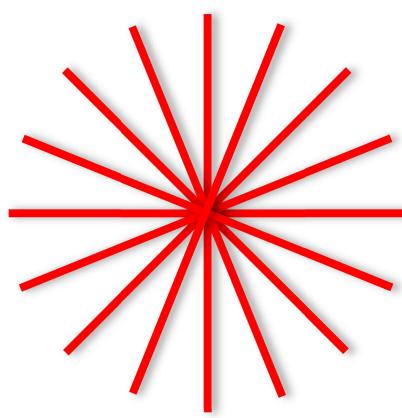
1 center of  
projection!

2 centers of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

Panorama

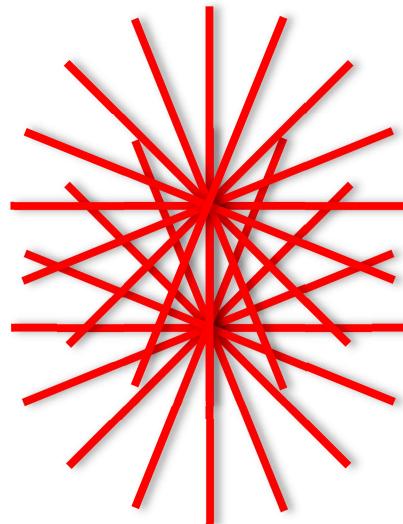
mono & head rotation



1 center of  
projection!

Stereo

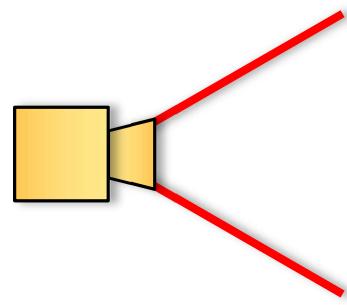
stereo & no head rotation



2 centers of  
projection!

Stereo Panorama

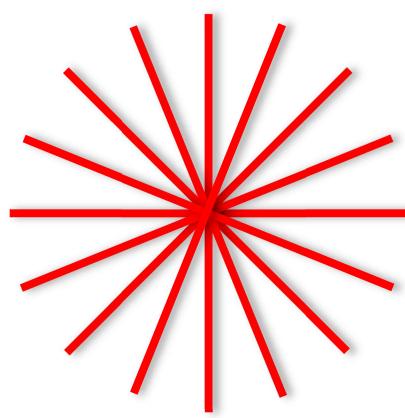
stereo & head rotation



# Panorama v Stereo Movie v Stereo Panorama

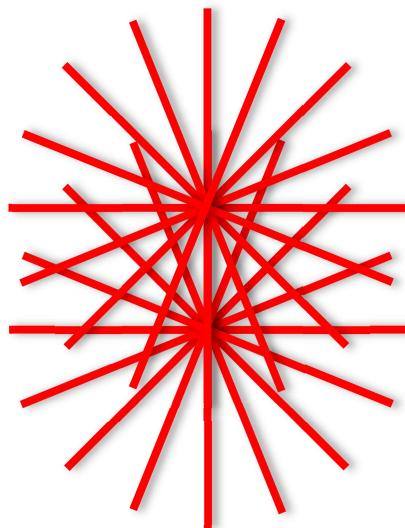
Panorama

mono & head rotation



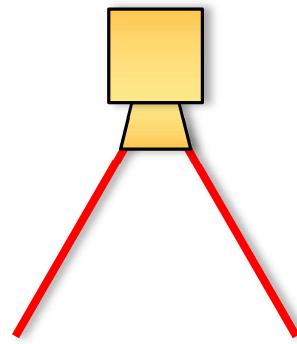
Stereo

stereo & no head rotation



Stereo Panorama

stereo & head rotation



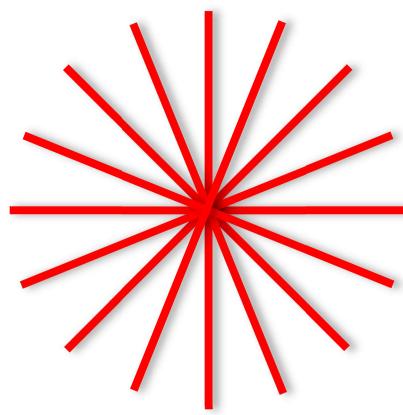
1 center of  
projection!

2 centers of  
projection!

# Panorama v Stereo Movie v Stereo Panorama

## Panorama

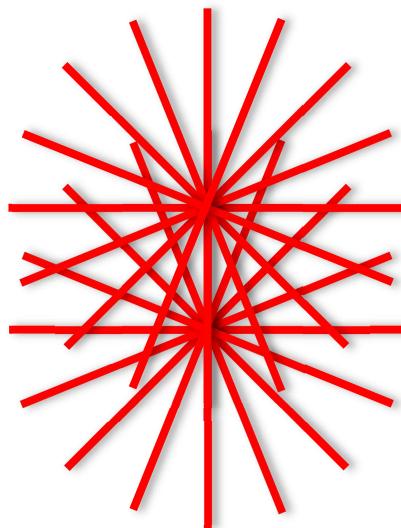
mono & head rotation



1 center of  
projection!

## Stereo

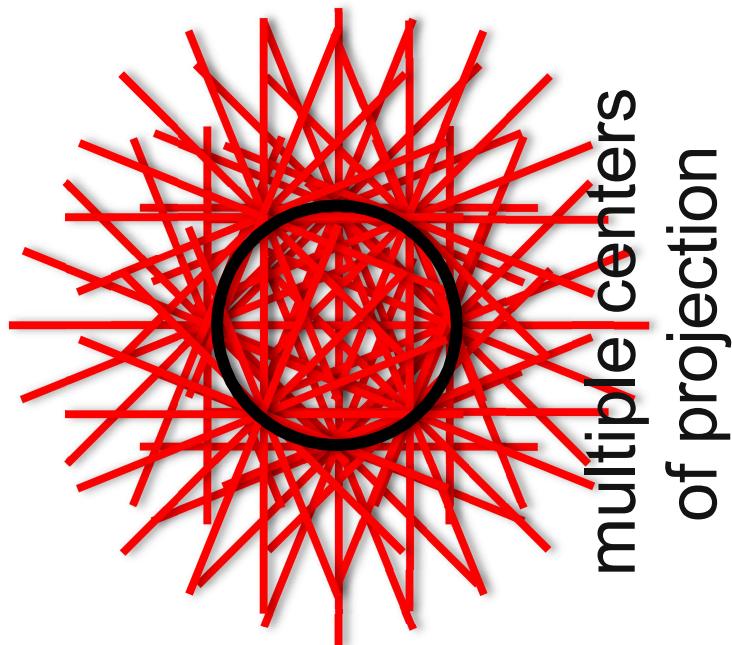
stereo & no head rotation



2 centers of  
projection!

## Stereo Panorama

stereo & head rotation

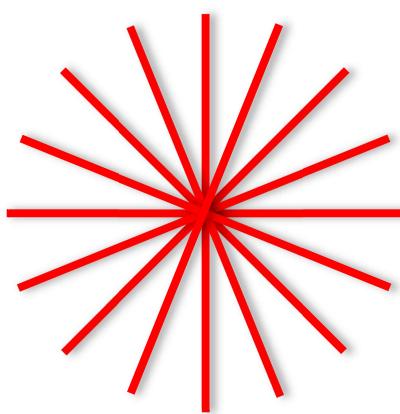


multiple centers  
of projection

# Panorama v Stereo Movie v Stereo Panorama

Panorama

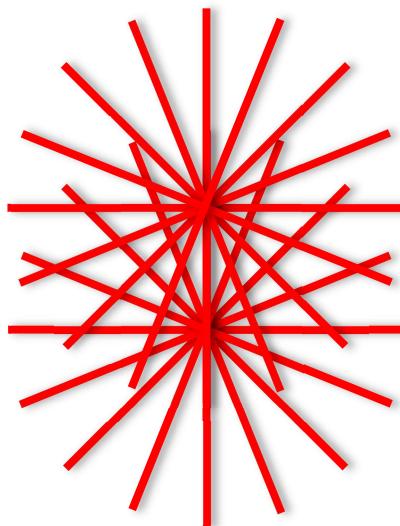
mono & head rotation



1 center of  
projection!

Stereo

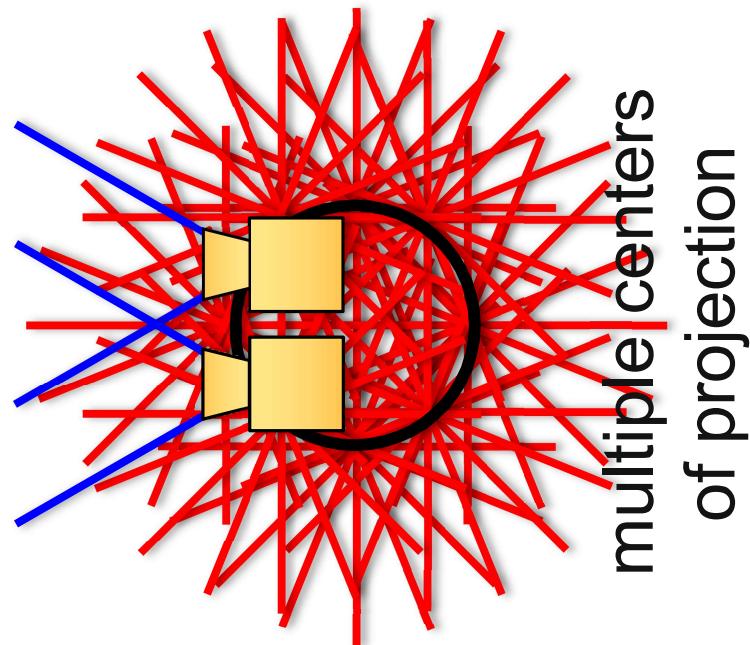
stereo & no head rotation



2 centers of  
projection!

Stereo Panorama

stereo & head rotation

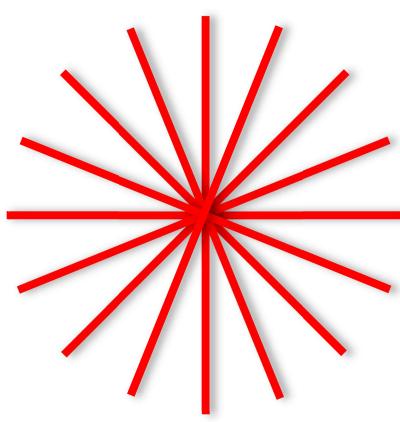


multiple centers  
of projection

# Panorama v Stereo Movie v Stereo Panorama

Panorama

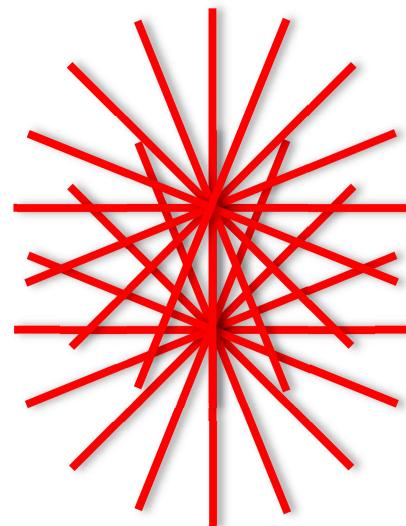
mono & head rotation



1 center of  
projection!

Stereo

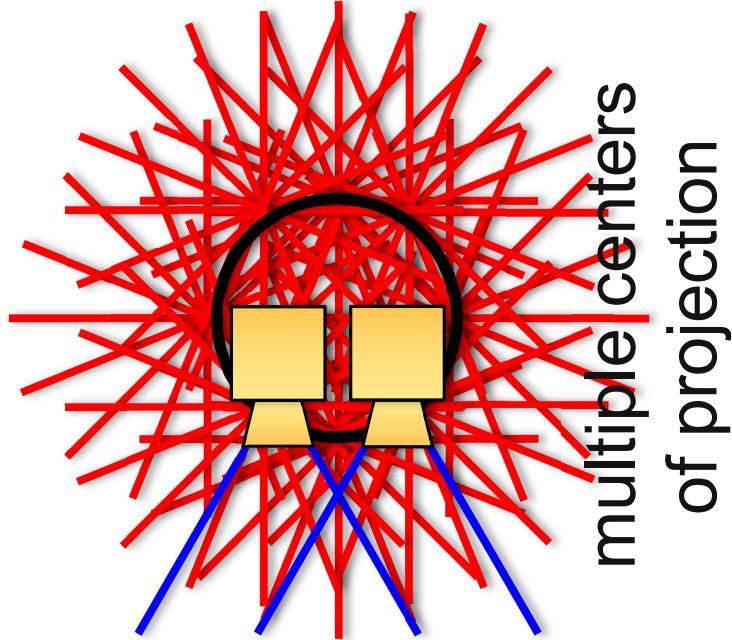
stereo & no head rotation



2 centers of  
projection!

Stereo Panorama

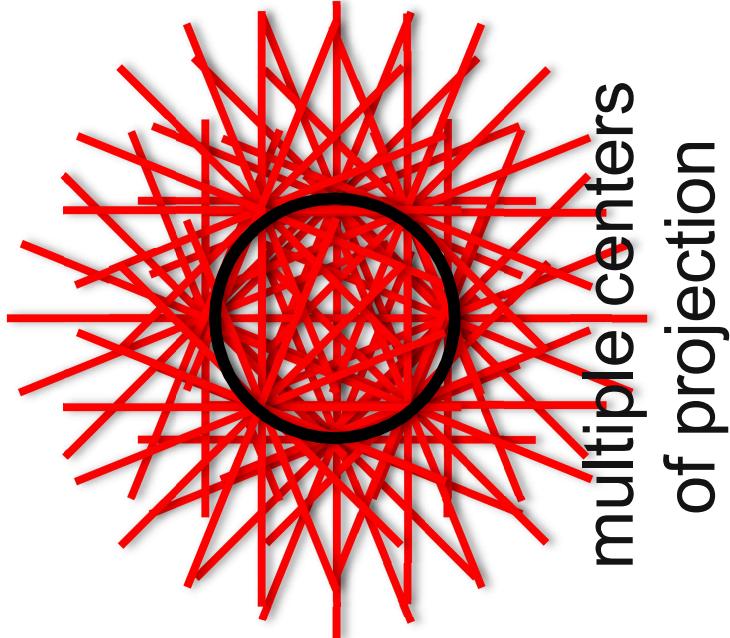
stereo & head rotation



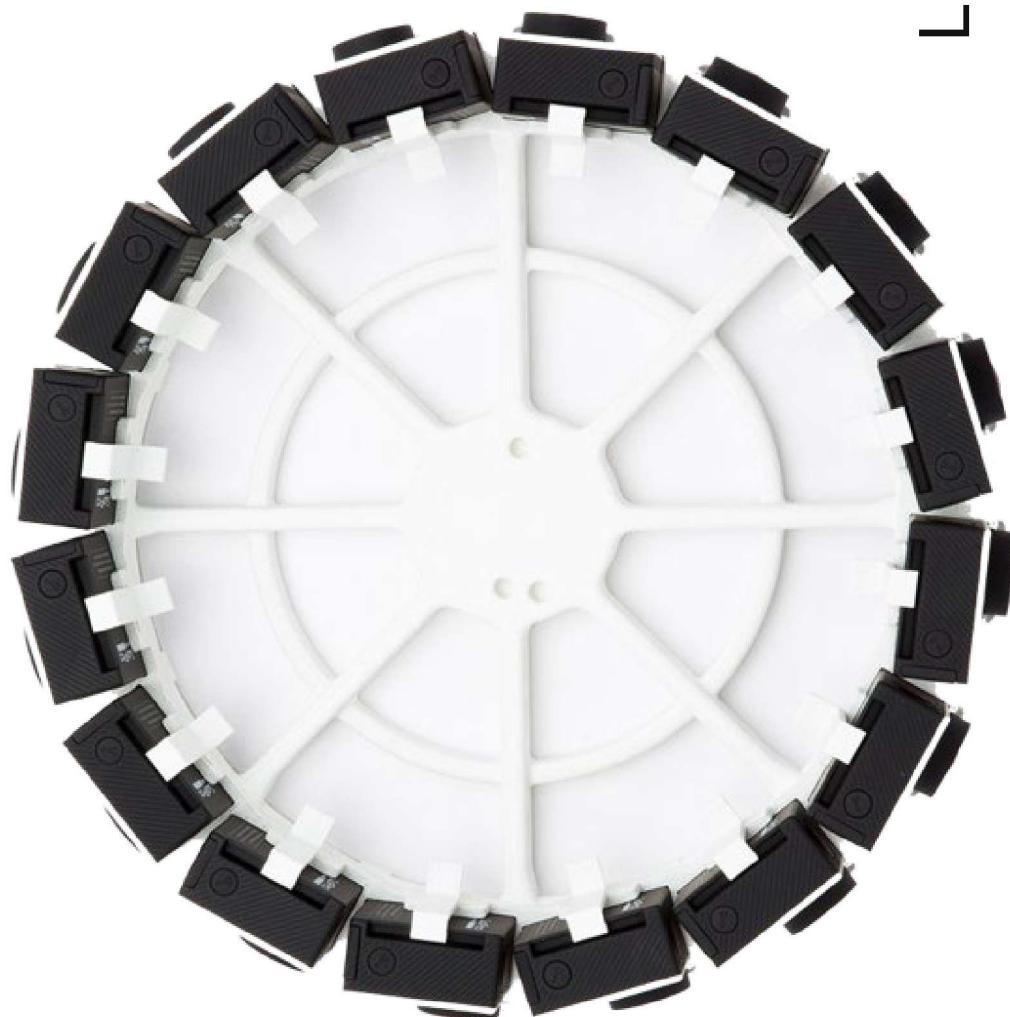
multiple centers  
of projection

# Panorama v Stereo Movie v Stereo Panorama

Stereo Panorama  
stereo & head rotation



Light Field!



# Panorama v Stereo Movie v Stereo Panorama

## Panorama

mono & head rotation



RICOH Theta

## Stereo

stereo & no head rotation



## Stereo Panorama

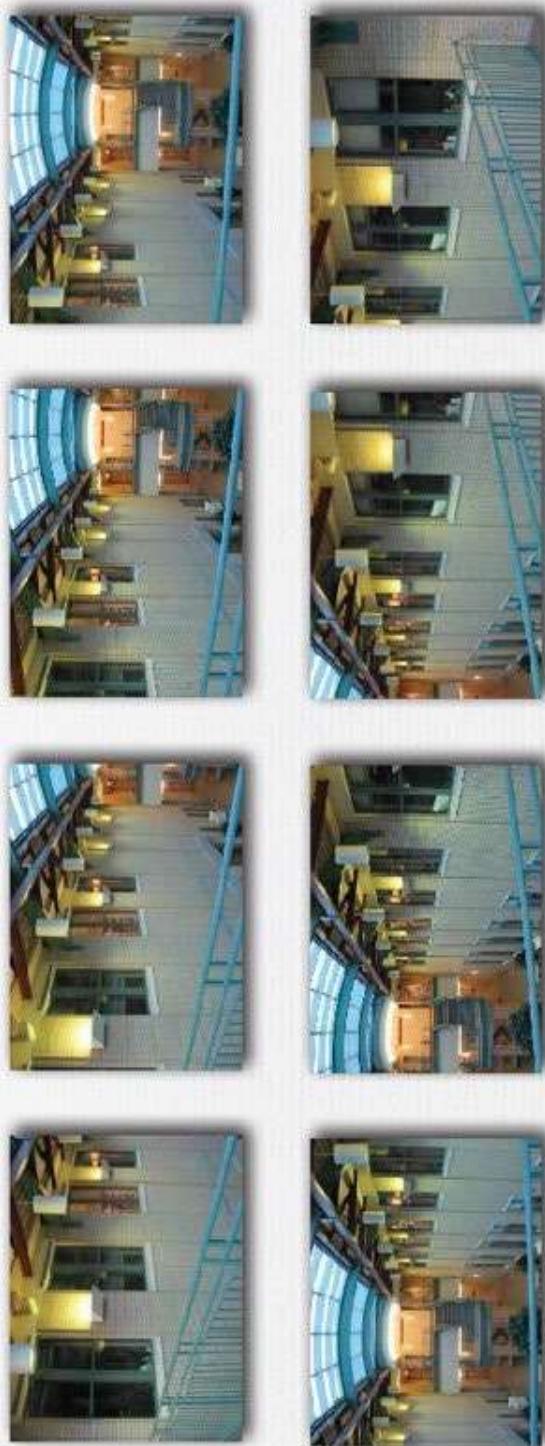
stereo & head rotation



horizontal-only  
parallax

Stitching images together to make a mosaic

---



18

© Marc Levoy

Slides from Marc Levoy

Panoramas

What kind of transformation do we need?

---



© Marc Levoy

61

Slides from Marc Levoy

Panoramas

## Stitching images together to make a mosaic



- ♦ step 1: find corresponding features in a pair of image
- ♦ step 2: compute perspective from 2<sup>nd</sup> to 1<sup>st</sup> image
- ♦ step 3: warp 2<sup>nd</sup> image so it overlays 1<sup>st</sup> image
- ♦ step 4: blend images where they overlap one another
- ♦ repeat for 3<sup>rd</sup> image and mosaic of first two, etc.

# Example: the Matterhorn



common  
picture  
plane of  
mosaic  
image

© Marc Levoy

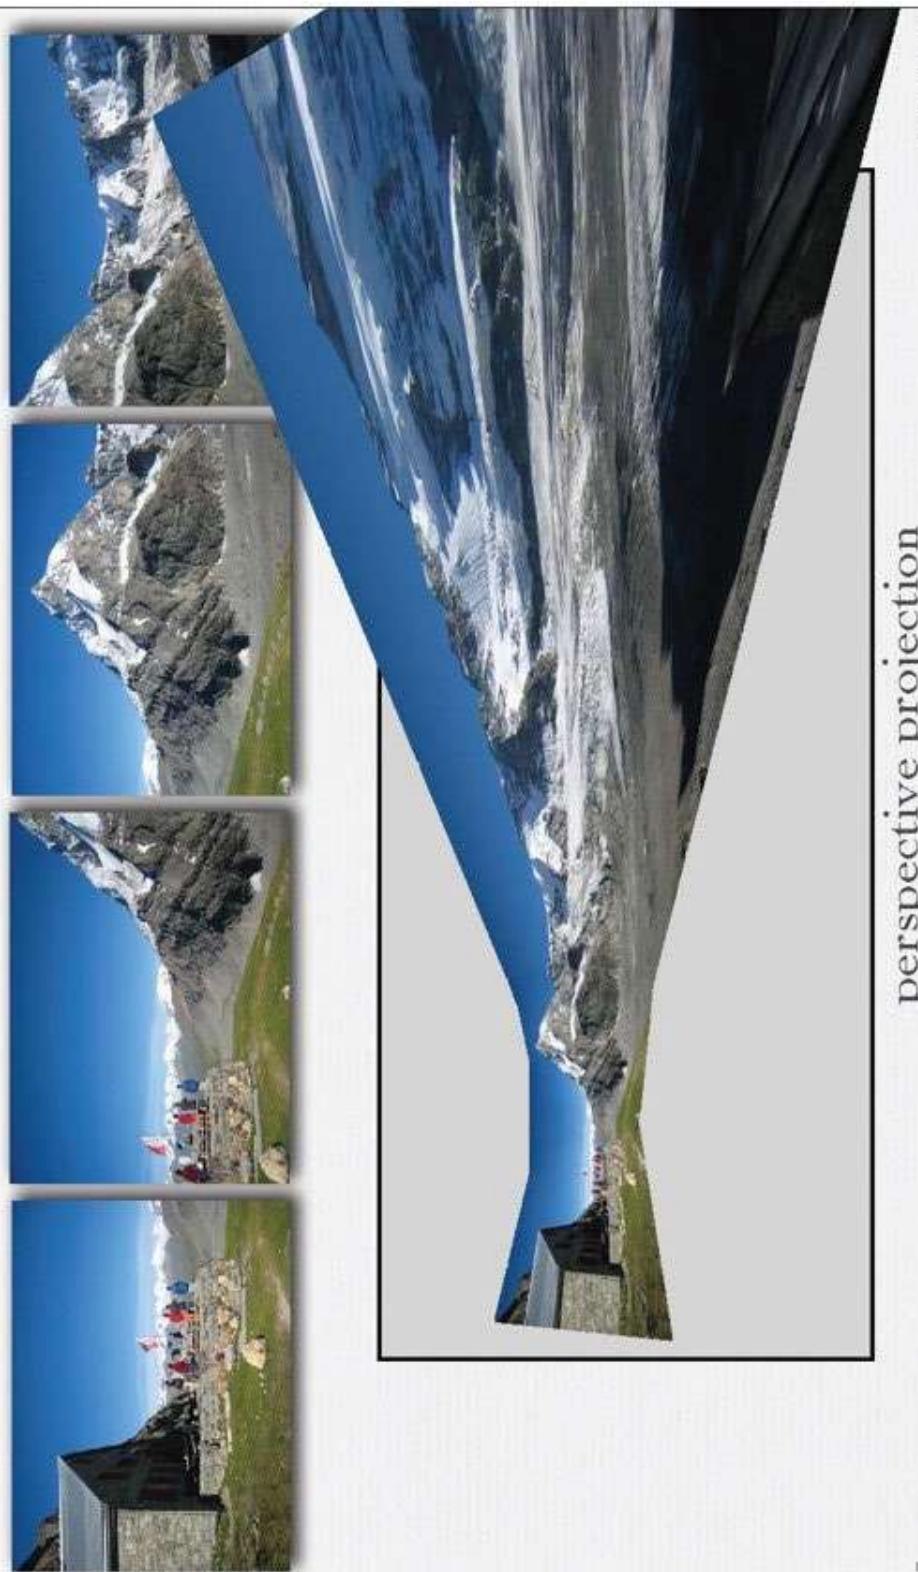
perspective projection

26

Slides from Marc Levoy

# Panoramas

Using 4 shots instead of 3



27

© Marc Levoy

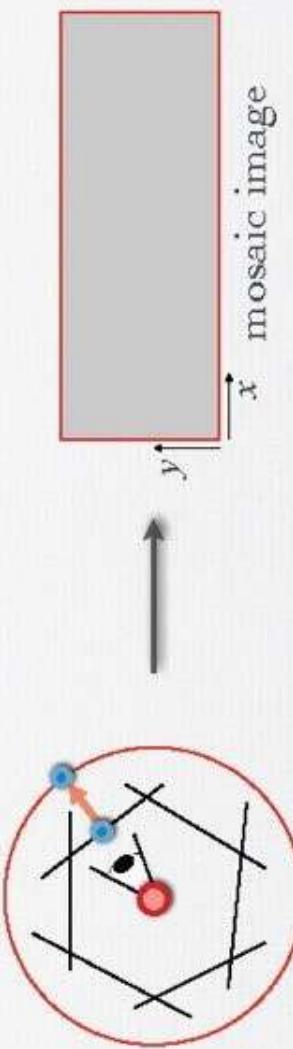
perspective projection

Slides from Marc Levoy

Panoramas

# Cylindrical panoramas

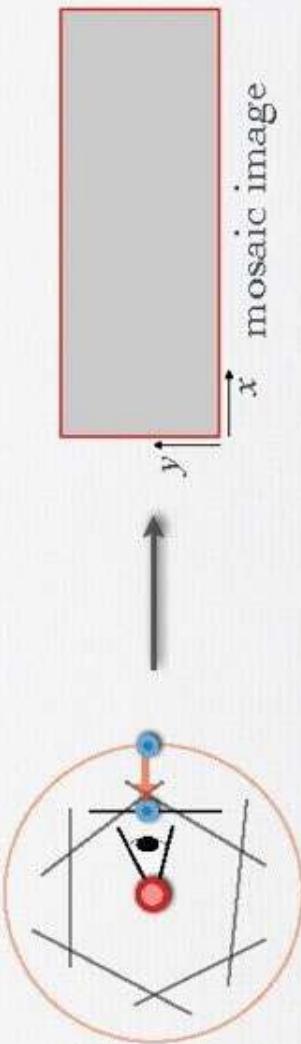
- even works for 360° panorama



- project each image onto a cylinder
- a cylindrical image can be stored as a rectangular image

# Cylindrical panoramas

- ♦ even works for 360° panorama



- ♦ project each image onto a cylinder
- ♦ a cylindrical image can be stored as a rectangular image
- ♦ to view without distortion, reproject part of the cylinder onto a picture plane representing the display screen
  - if your FOV is narrow, this view won't be too distorted

(FLASH DEMO)

<http://graphics.stanford.edu/courses/cs173/applets/projections.html>

# Back to the Matterhorn



© Marc Levoy

cylindrical projection

30

Slides from Marc Levoy

Panoramas

# Back to the Matterhorn



blended

© Marc Levoy

Slides from Marc Levoy

Panoramas

# Spherical panoramas



- ◆ projections are to a sphere instead of a cylinder
- ◆ can't store as rectangular image without extreme stretching

© Marc Levoy

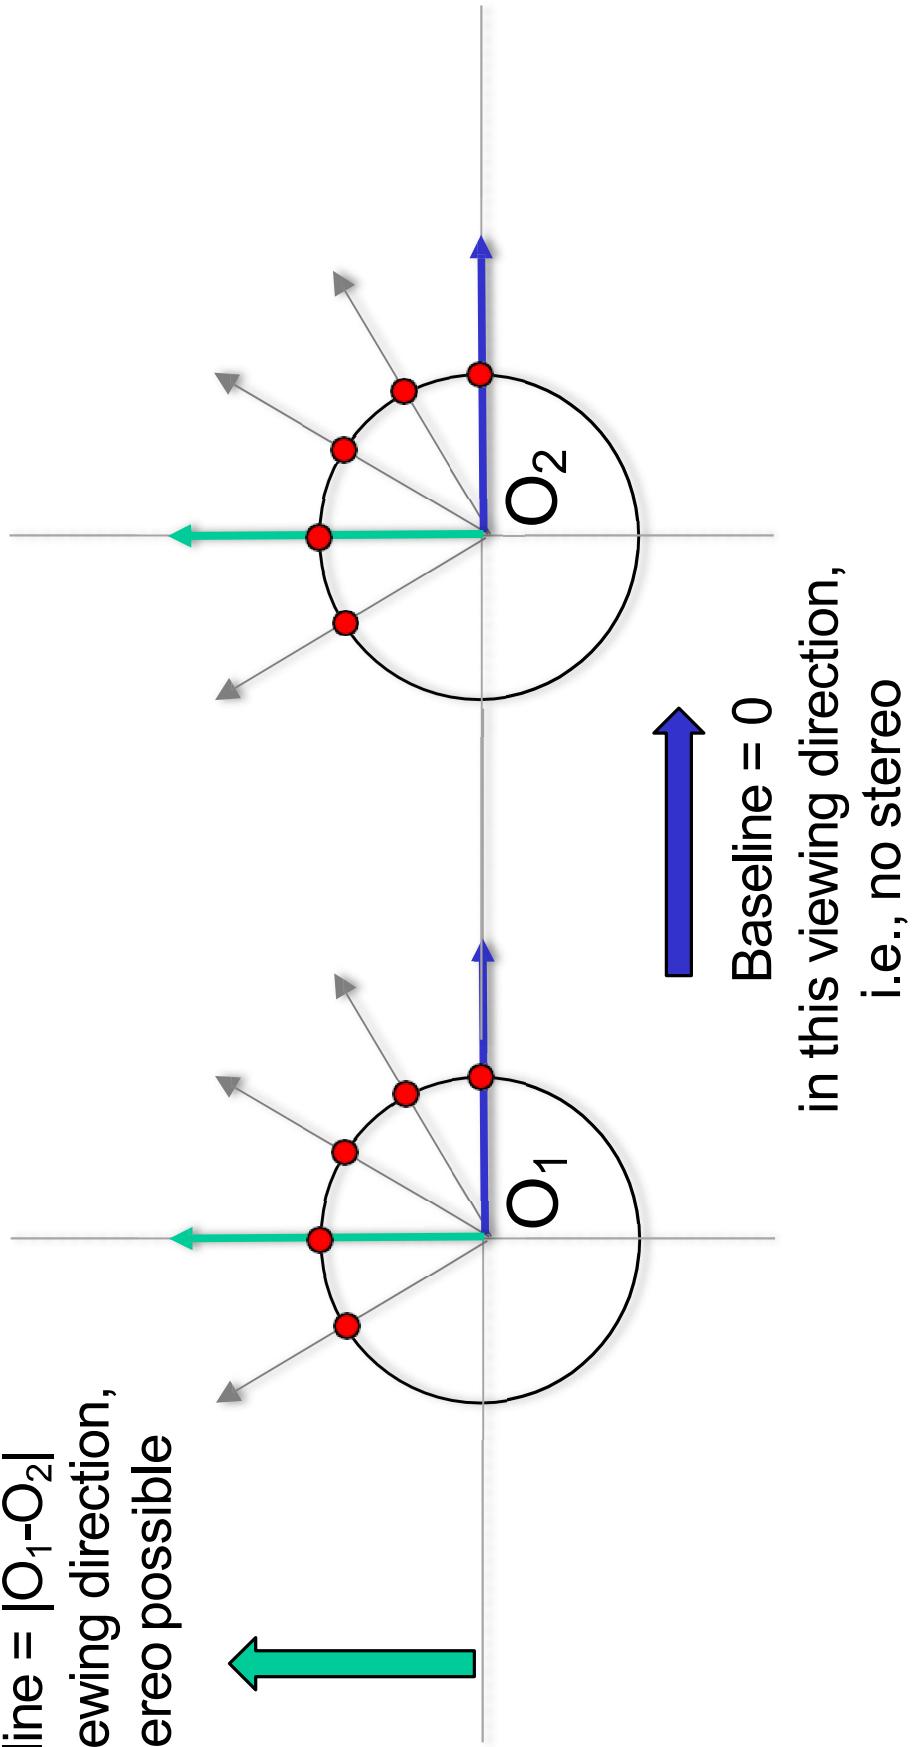
34

Slides from Marc Levoy

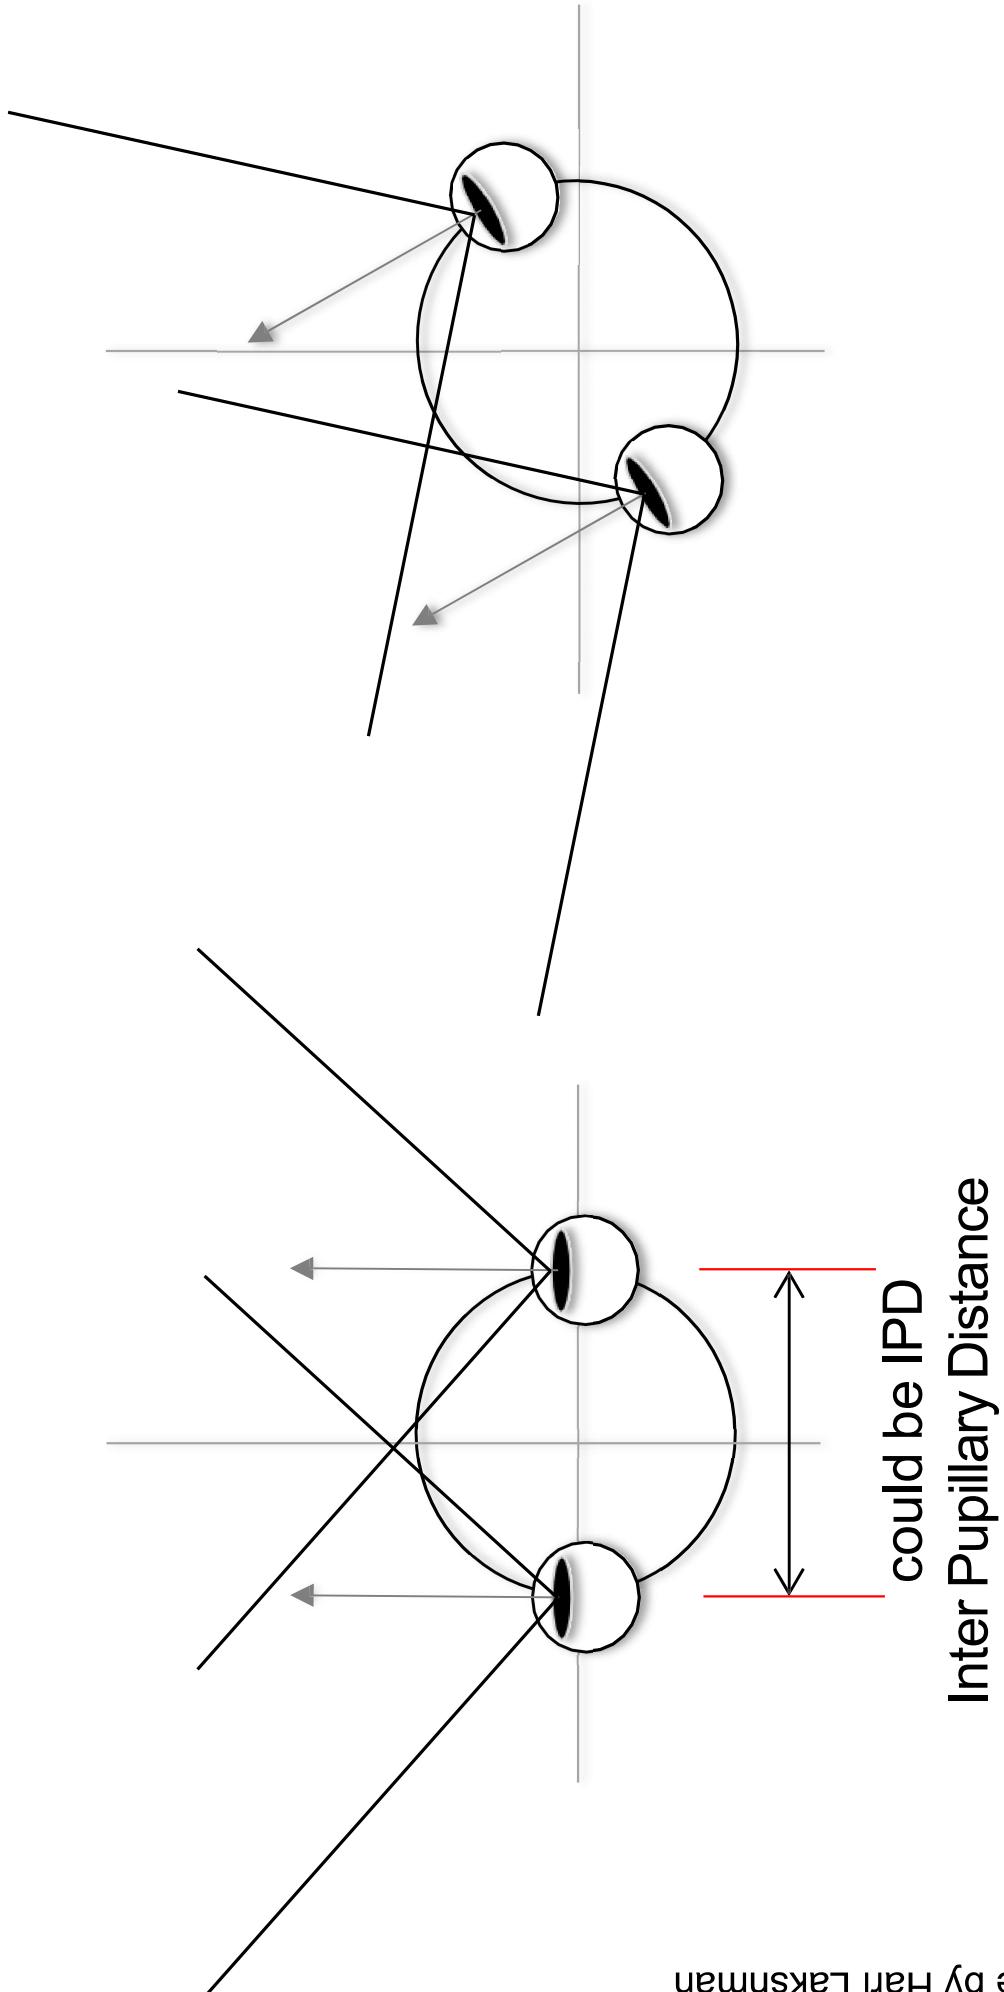
Panoramas

# A Pair of Mono Panoramas

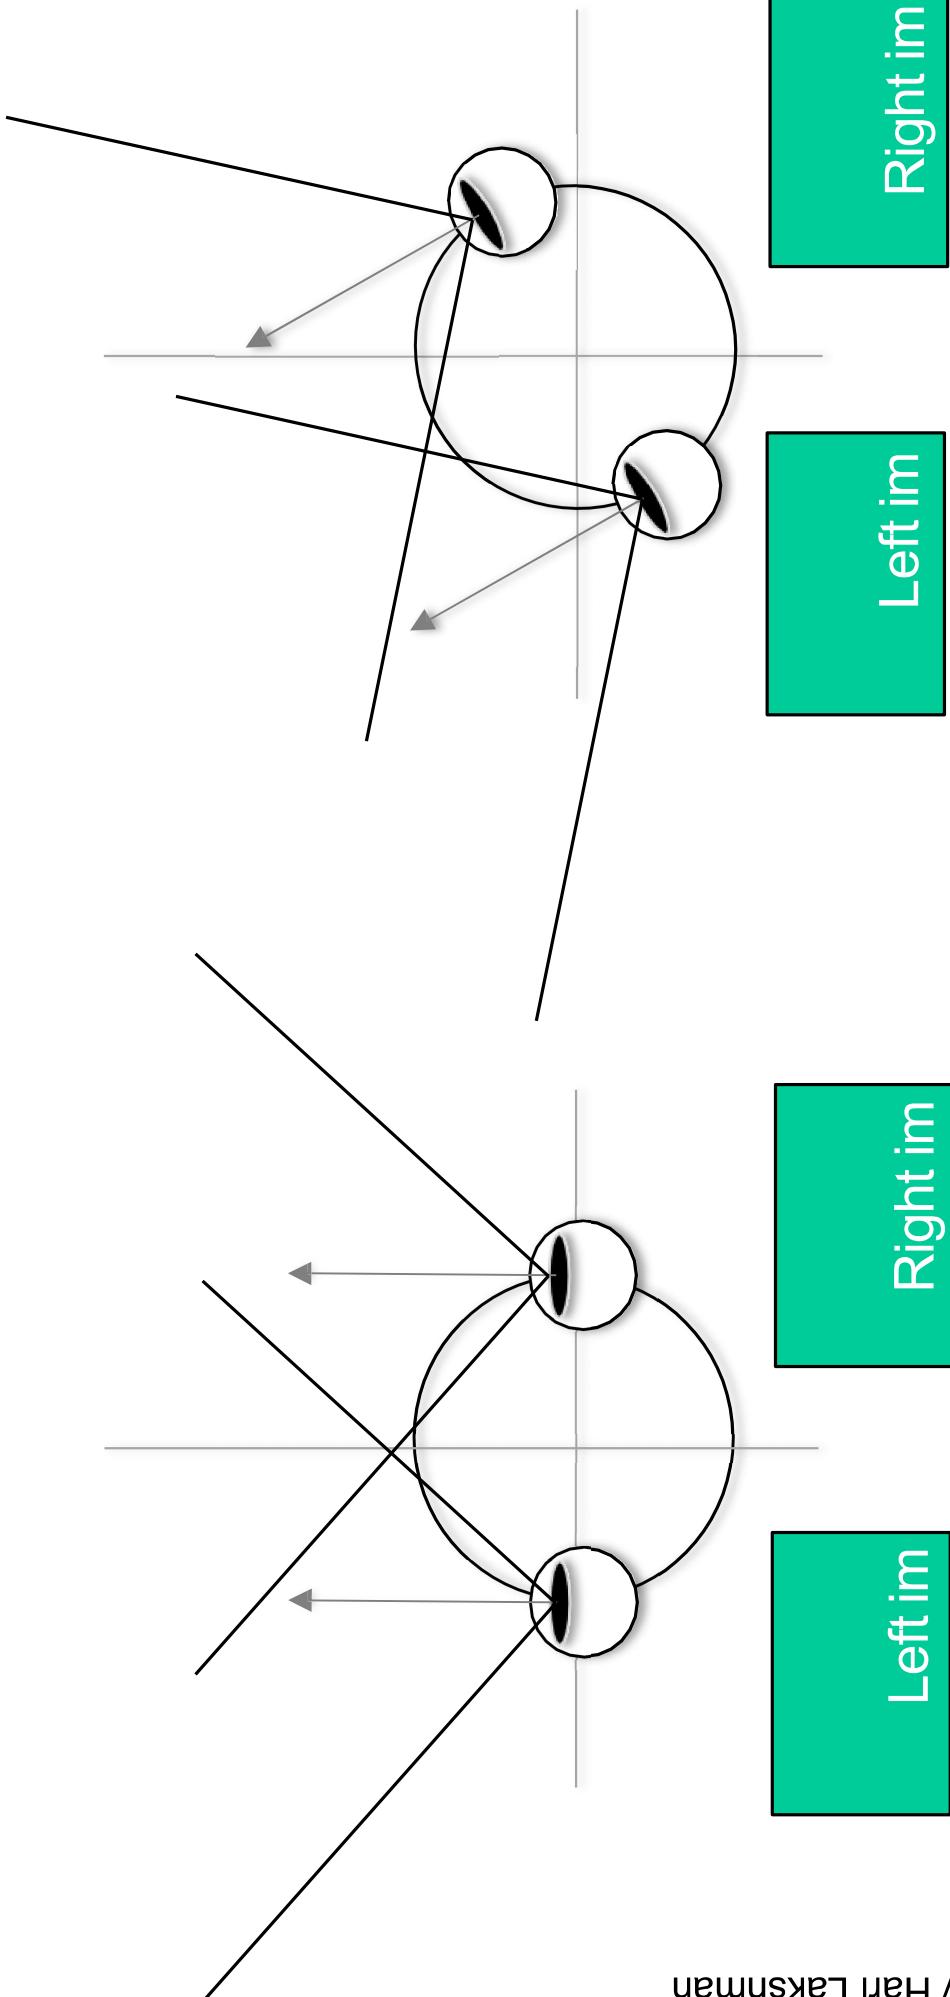
Baseline =  $|O_1 - O_2|$   
in this viewing direction,  
i.e., stereo possible



# Head Rotation

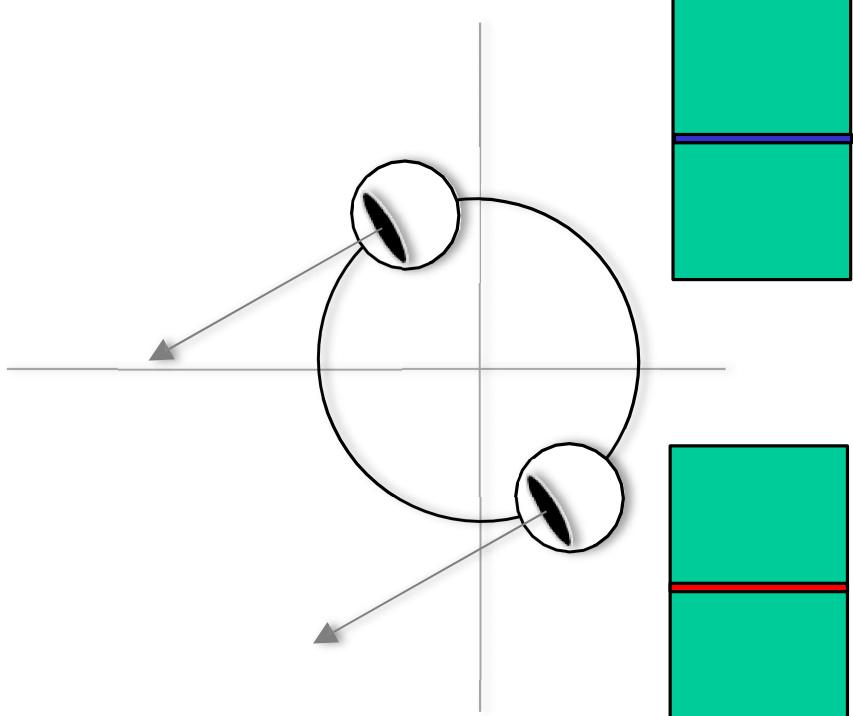


# Image Pair for Each Direction

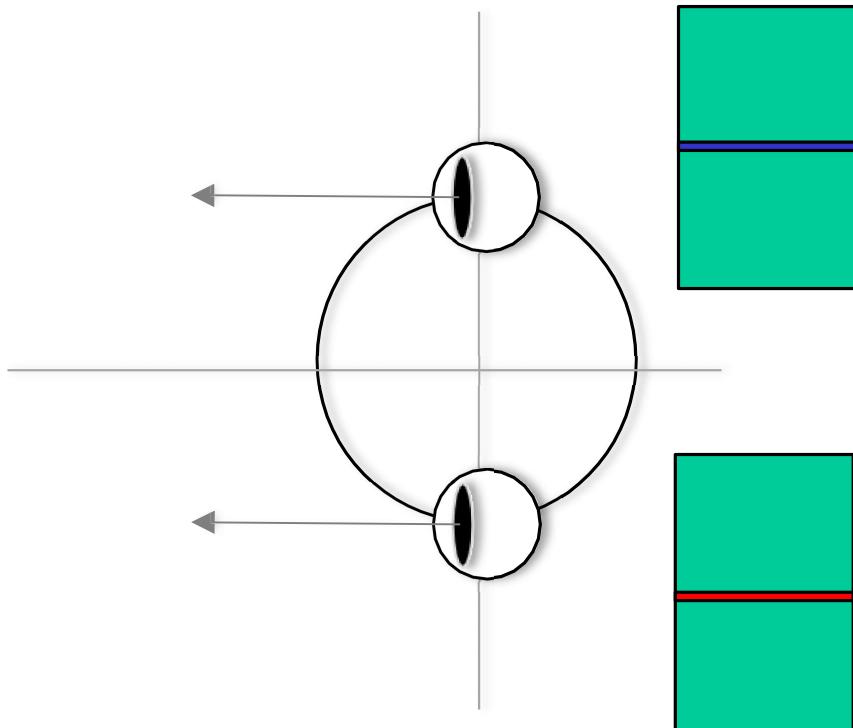


Store image pair for each direction ➔ Problem: Too much data

# Approximation: Store only Middle Ray

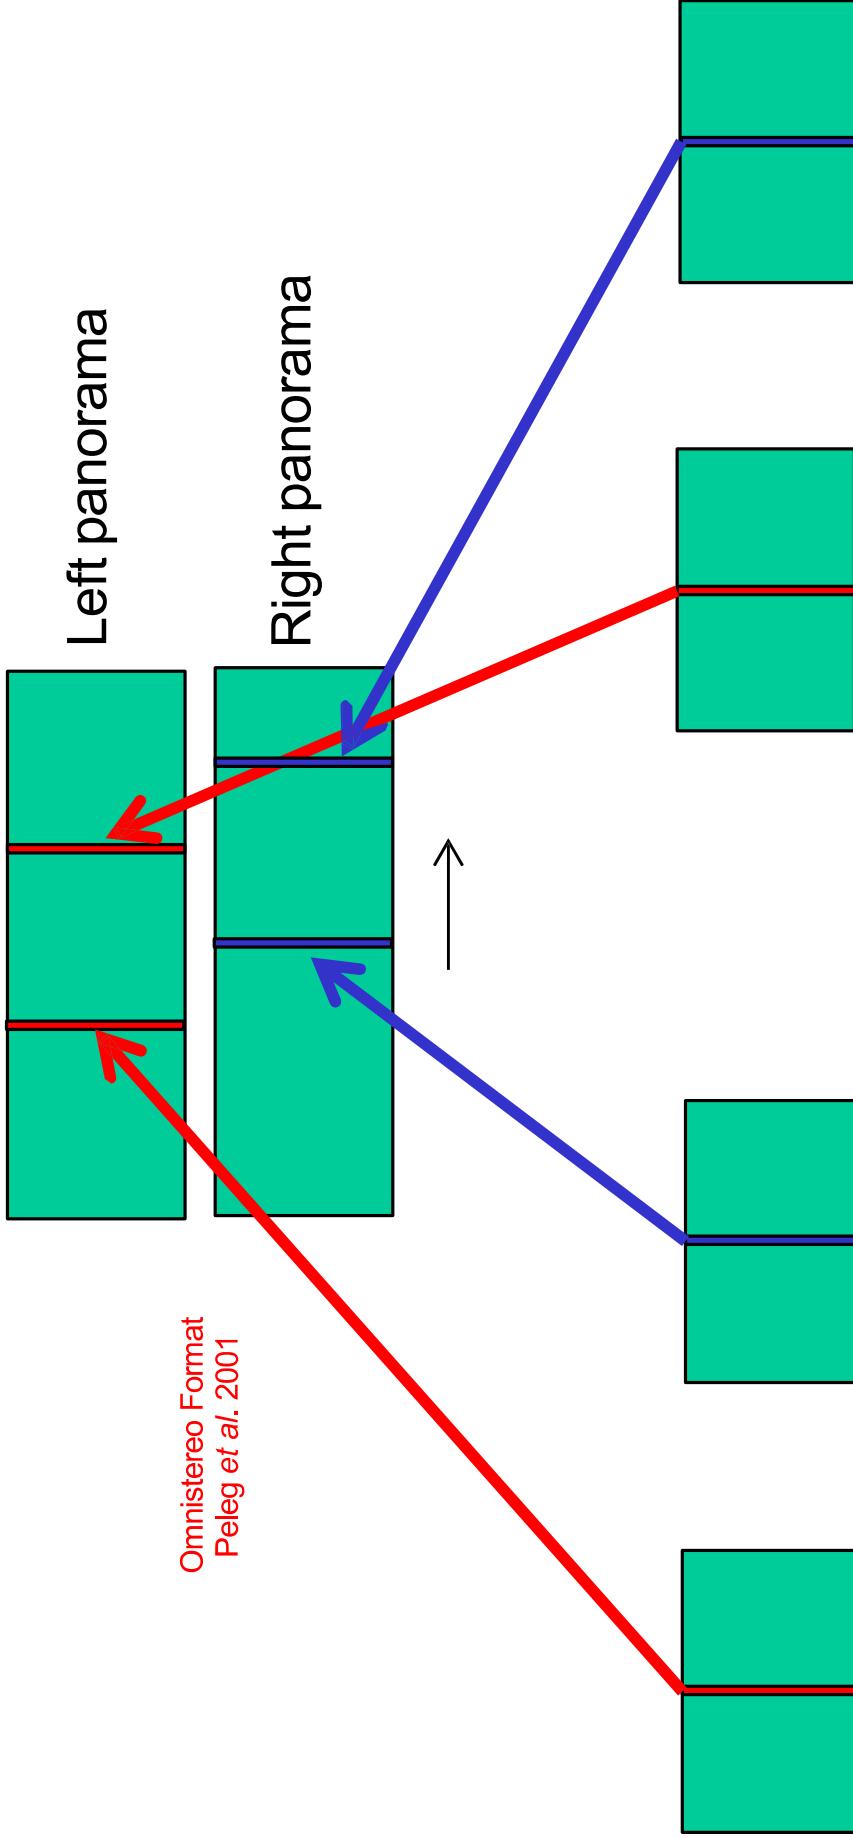


Omnistereo Format  
Peleg et al. 2001

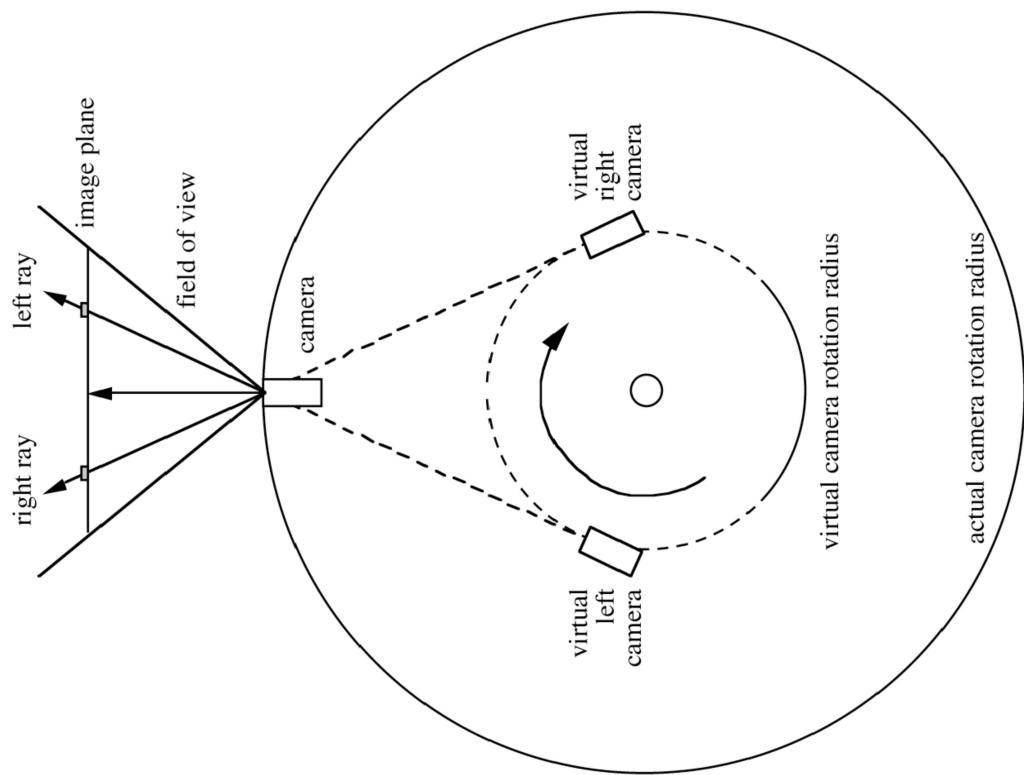


Approximation: store only middle ray for L and R eyes for each direction

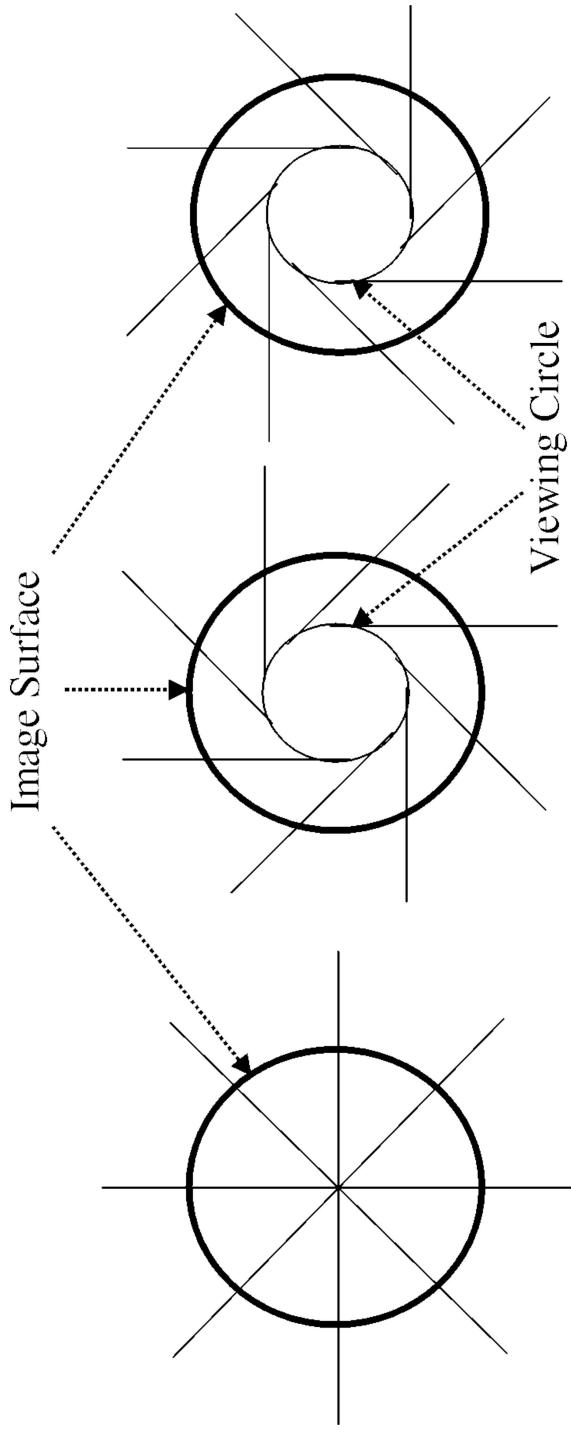
# Omnistereo Panoramas



# Capture using Single Camera



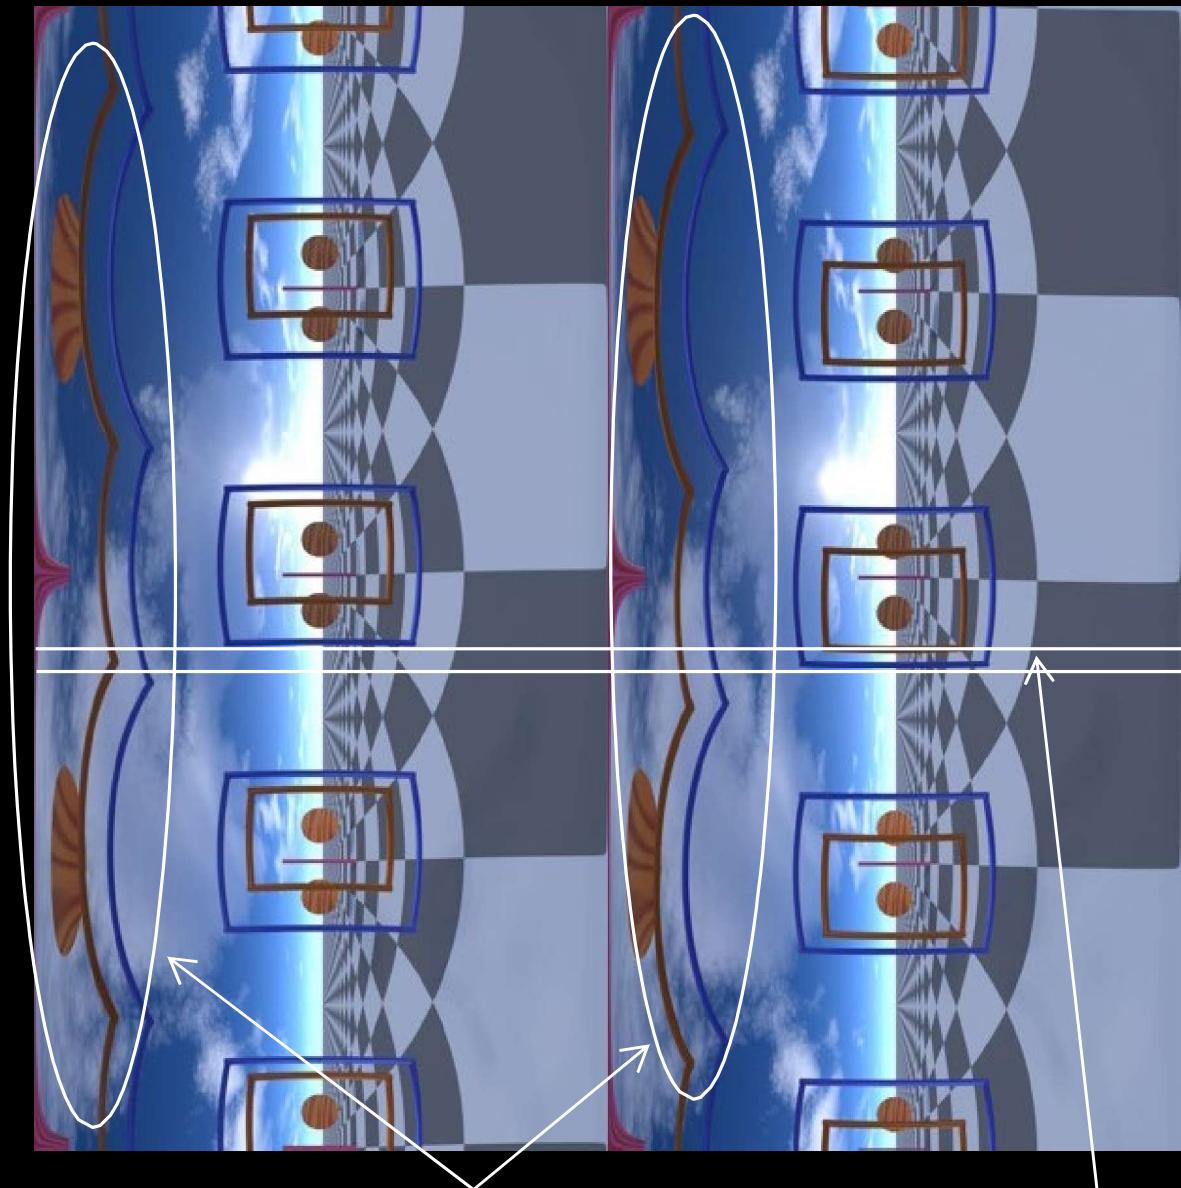
# Comparison: Mono and Stereo Panoramas



**Omnistereo,  
Multiperspective**

**Central,  
a.k.a. Mono**

Omnistereo example



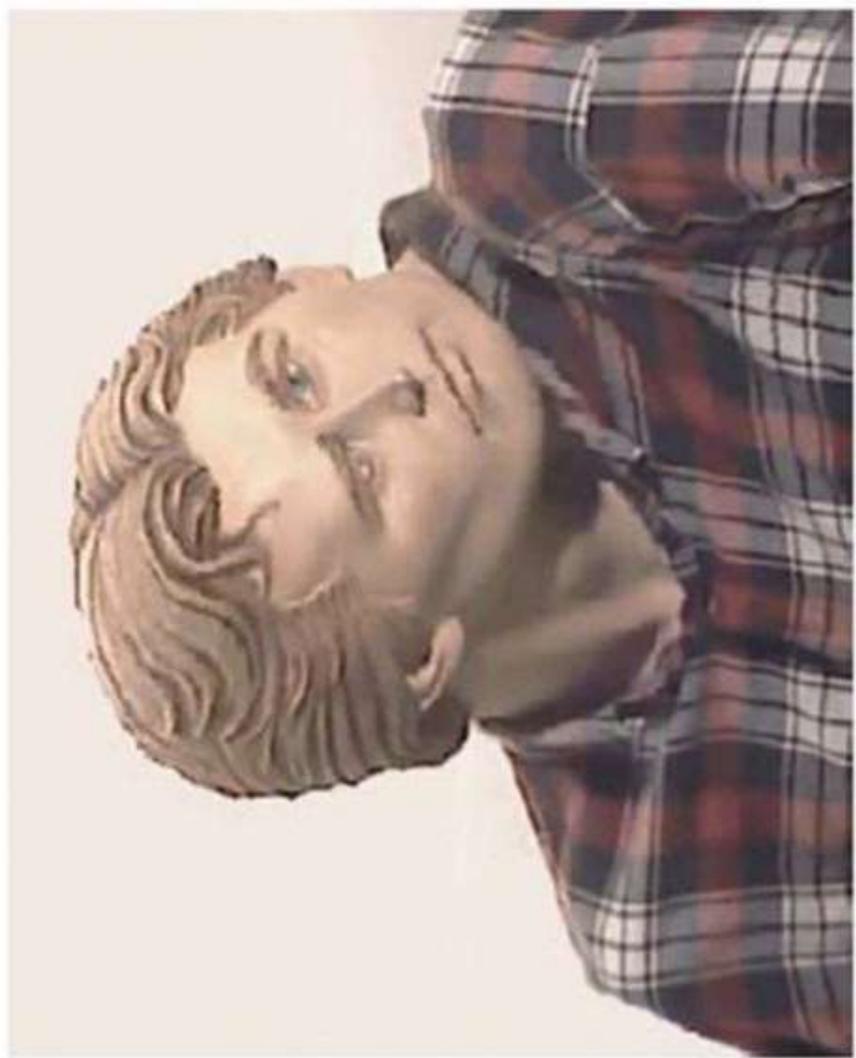
Left panorama

Right panorama

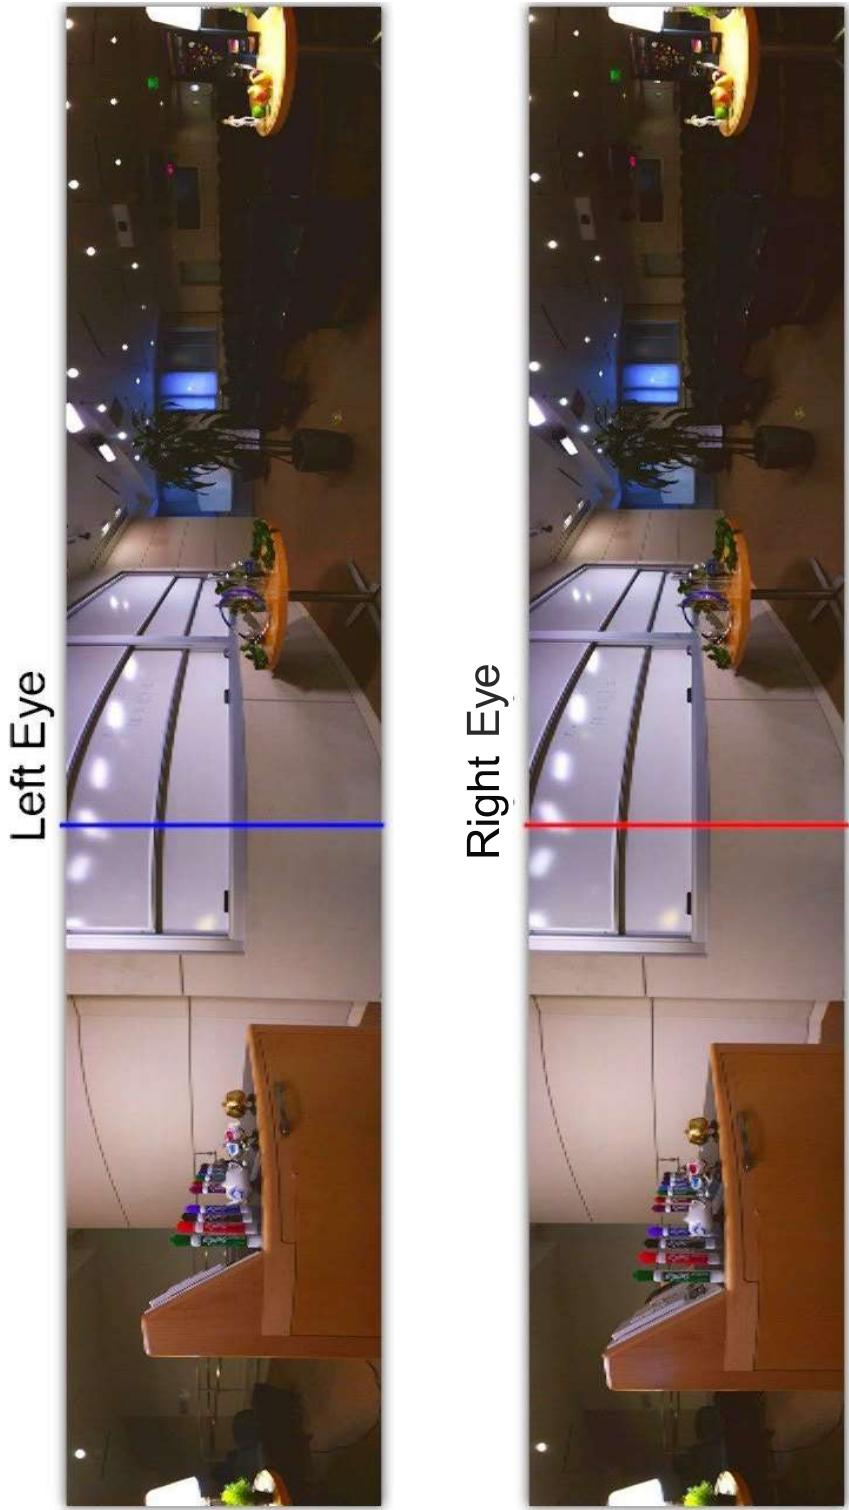
Sphere-to-plane  
distortions

Disparity

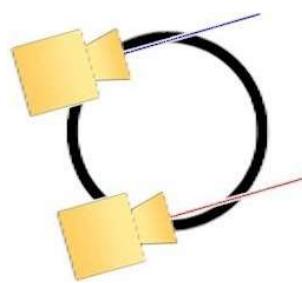
# Multiperspective Projection



# Omnidirectional Stereo



widely used by YouTube VR, Google Daydream, Facebook, ...



# Existing VR Cameras

Recorded Videos ~ 17 Gb/sec



# Facebook's Surround 360



RAW Data: 17 Gb/sec

Compute time: days to weeks on conventional computer,  
minutes to hours on data center

# Facebook's Surround 360



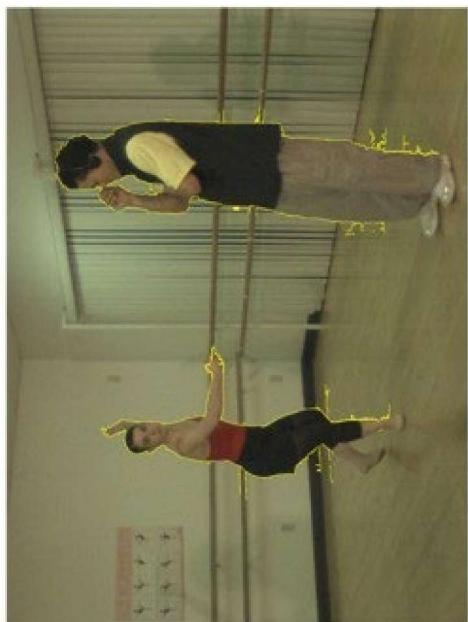
RAW Data: 17 Gb/sec

Compute time: days to weeks on conventional computer,  
minutes to hours on data center

# Biggest Problem of Panorama – 6DoF Viewing



(b)

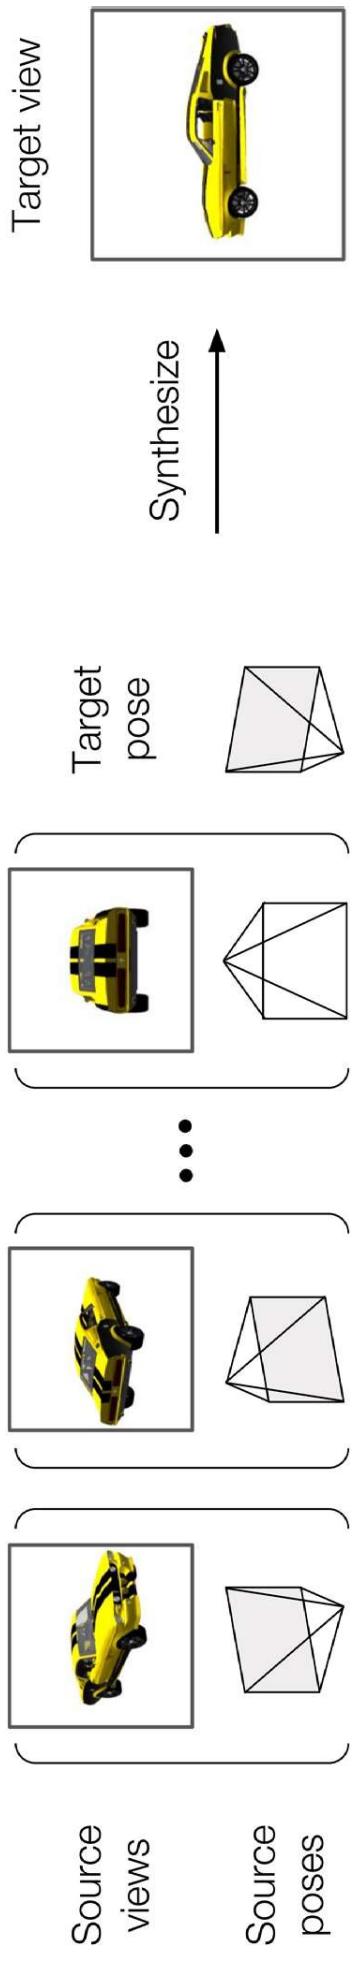


(a)



# Problem: Novel View Synthesis

- ❖ The process of using images of a scene and their camera poses to synthesize new images of the scene at arbitrary camera poses.



- ❖ Subset of *Neural Rendering*, “deep image or video generation approaches that enable explicit or implicit control of scene properties”. (Ayush Tewari et. al.)
  - ❖ E.g. illumination, camera parameters, pose, geometry, appearance, and semantic structure

Image: “Multi-view to Novel View: Synthesizing Novel Views with Self-Learned Confidence” (Sun et al., 2018)

# Applications in VR



NeX: Real-time View Synthesis with Neural Basis Expansion

Real-time rendering

# Related Work

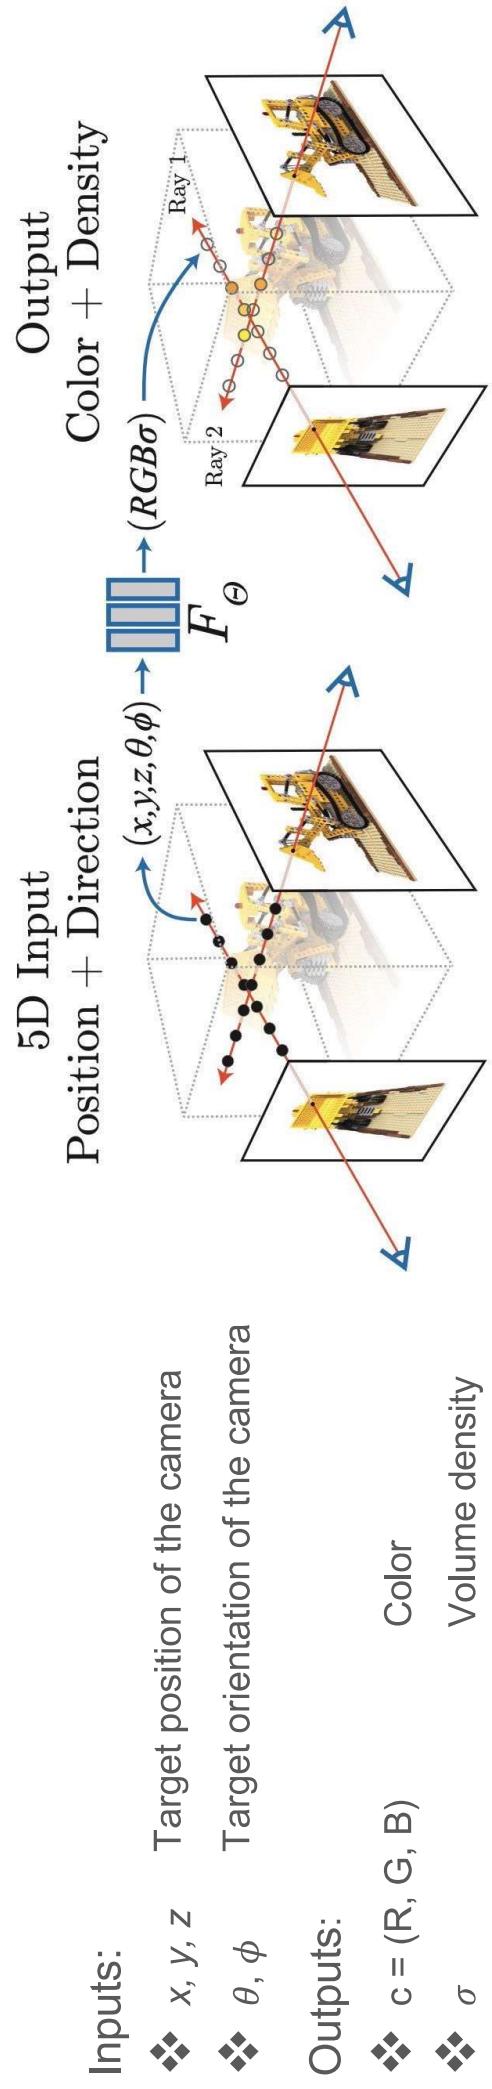
- ❖ Neural Volumes (NV), 2019
  - ❖ Deep 3D convolutional network architecture
  - ❖ Predicts a fixed-size discretized voxel grid
  - ❖ Limitation: discrete voxel grids do not scale well and lose fine detail at high resolutions
  - ❖ Limitation: requires a bounded volume and knowledge of the background
- ❖ Scene Representation Networks (SRN), 2019
  - ❖ Uses a recurrent neural network to model a rendering function
  - ❖ Limited to simple shapes with low geometric complexity
- ❖ Local Light Field Fusion (LLFF), 2019
  - ❖ 3D convolutional network architecture
  - ❖ Predicts multiplane images and fuses them to create new views
  - ❖ Fast to train (<10 minutes) at the cost of large storage requirements (~GB for each scene)

# NeRF: Key Insights

- ❖ Represent static scenes in a **continuous** space.
- ❖ Encode a continuous radiance field **within the parameters** of a fully-connected neural network.
- ❖ Regress directly from viewing location and direction to color and transparency

# Problem Setting

- ❖ Given a dataset containing RGB images of a static scene, their corresponding camera poses, and intrinsic parameters,
- ❖ Predict the color and volume density for every viewing location and direction

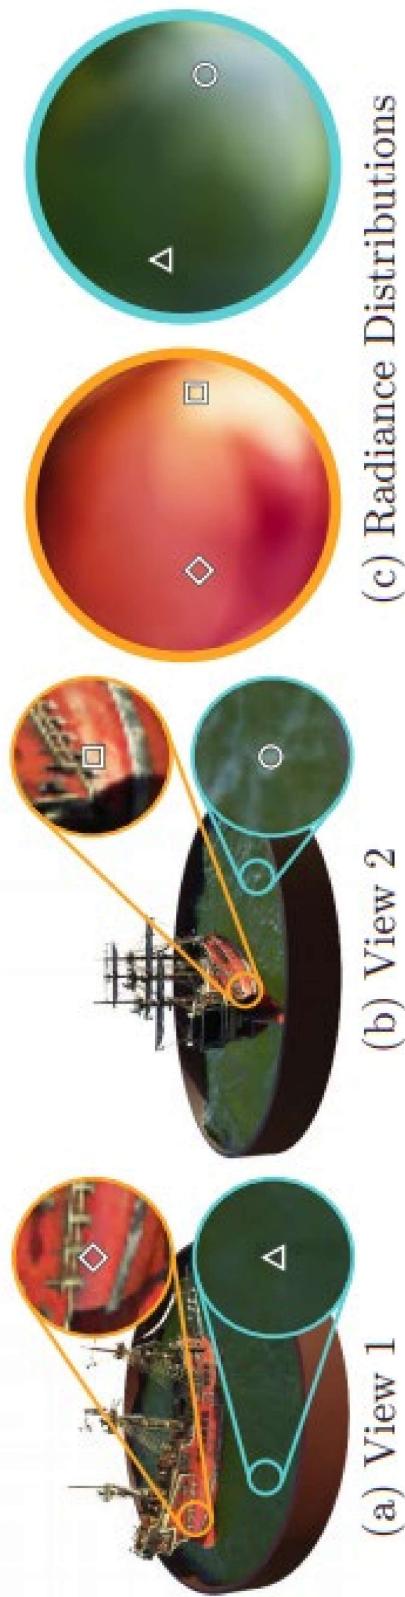


# NeRF: Inspired by Volume Rendering



# Definition of radiance field

- Radiance field is a 5-dimensional function which maps a 3D location  $\underline{x}, \underline{y}, \underline{z}$  and a direction in 3D sphere  $\underline{d}$  to a color  $(r, g, b)$ :



[8] Mildenhall20

# NeRF: Volume Rendering

- ❖ Generating a view from NeRF requires rendering all rays that pass through each pixel of the desired virtual camera

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Expected color of a camera ray

Predicted Volume Density

Predicted Color

Probability that nothing has blocked the ray up to this point

- ❖ Numerically estimated using quadrature and stratified sampling

$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i, \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

The relative contribution of this segment

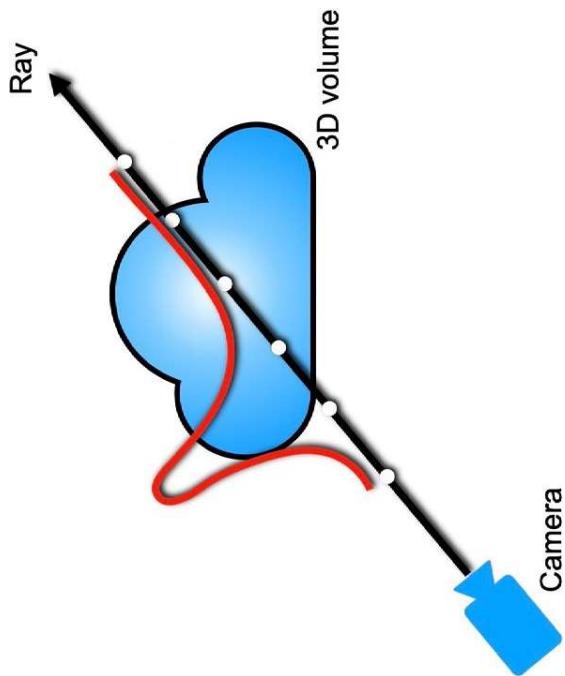
Predicted Color

Probability that nothing has blocked the ray up to this point

- ❖ Differentiable: allows optimization using gradient descent

# NeRF: Hierarchical Sampling

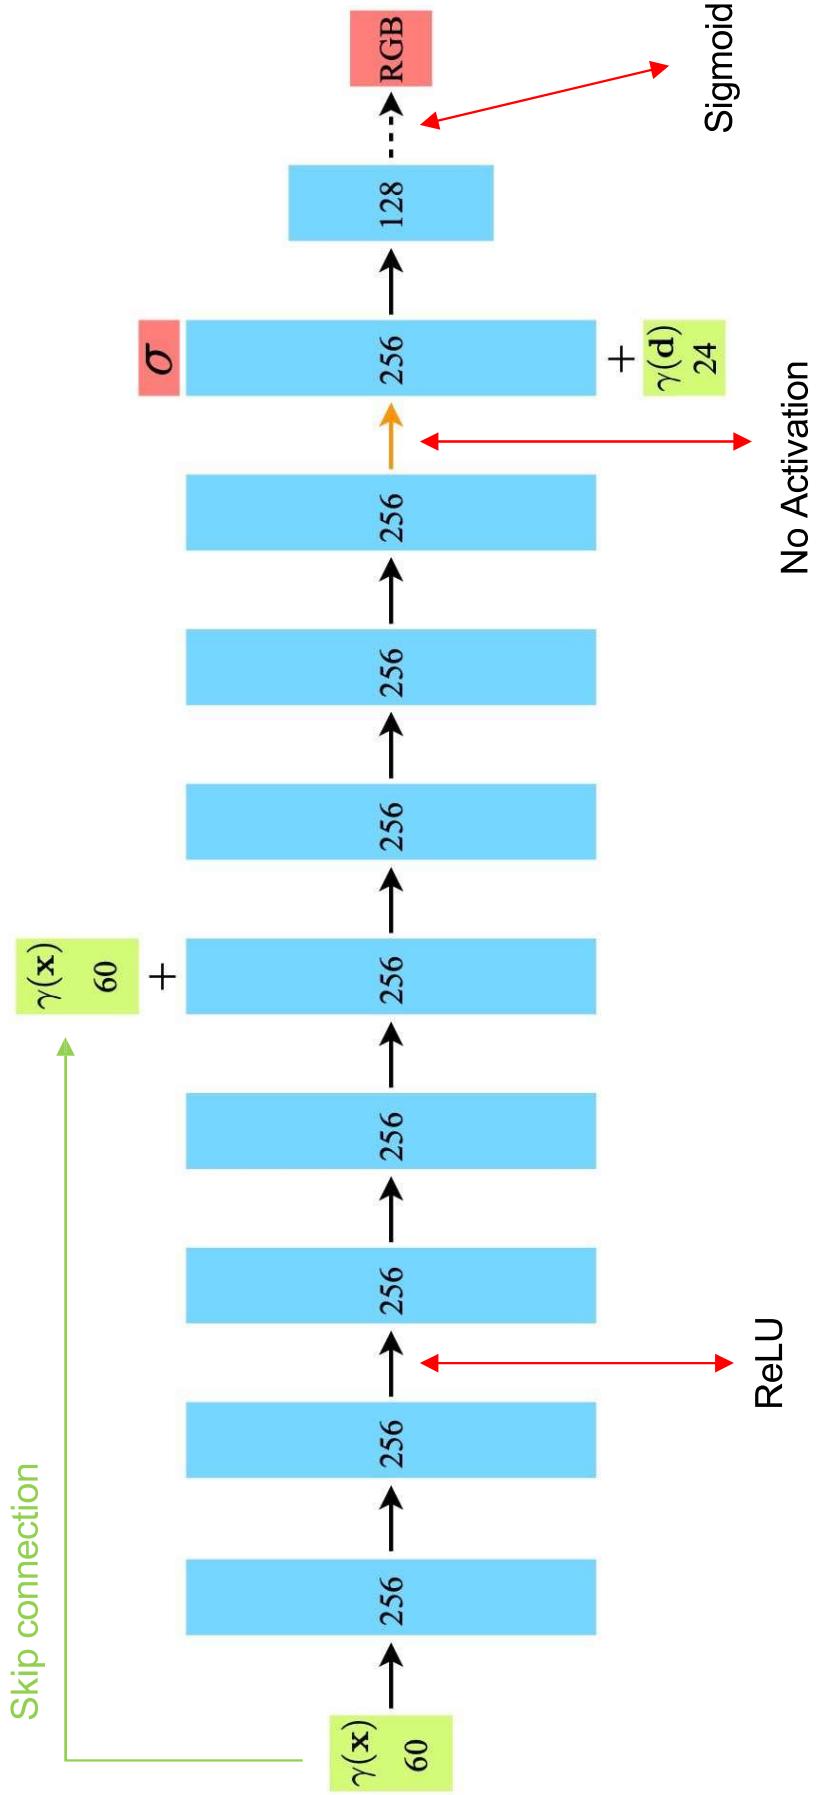
- ❖ Problem: It is inefficient to integrate over empty and occluded spaces in a scene
- ❖ Solution: Allocate samples proportionally to their expected effect on the final rendering.
- ❖ Evaluate a “coarse” network on a set of  $N_c$  locations along a ray to produce a PDF along the ray
- ❖ Evaluate a “fine” network on  $N_c$  and a second a set of  $N_f$  locations sampled from the PDF



# NeRF: Positional Encoding

- ❖ Map individual components of position and direction vectors to a higher dimensional space
  
$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$
- ❖ Similar concept as positional encoding in Transformer Architectures
- ❖ Empirically improves the preservation of high-frequency **geometry** and **texture**
- ❖ Surprising result, explored further in a follow-up work by the same authors
  - ❖ “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains” (Tancik, Srinivasan, Mildenhall, et al., 2020)

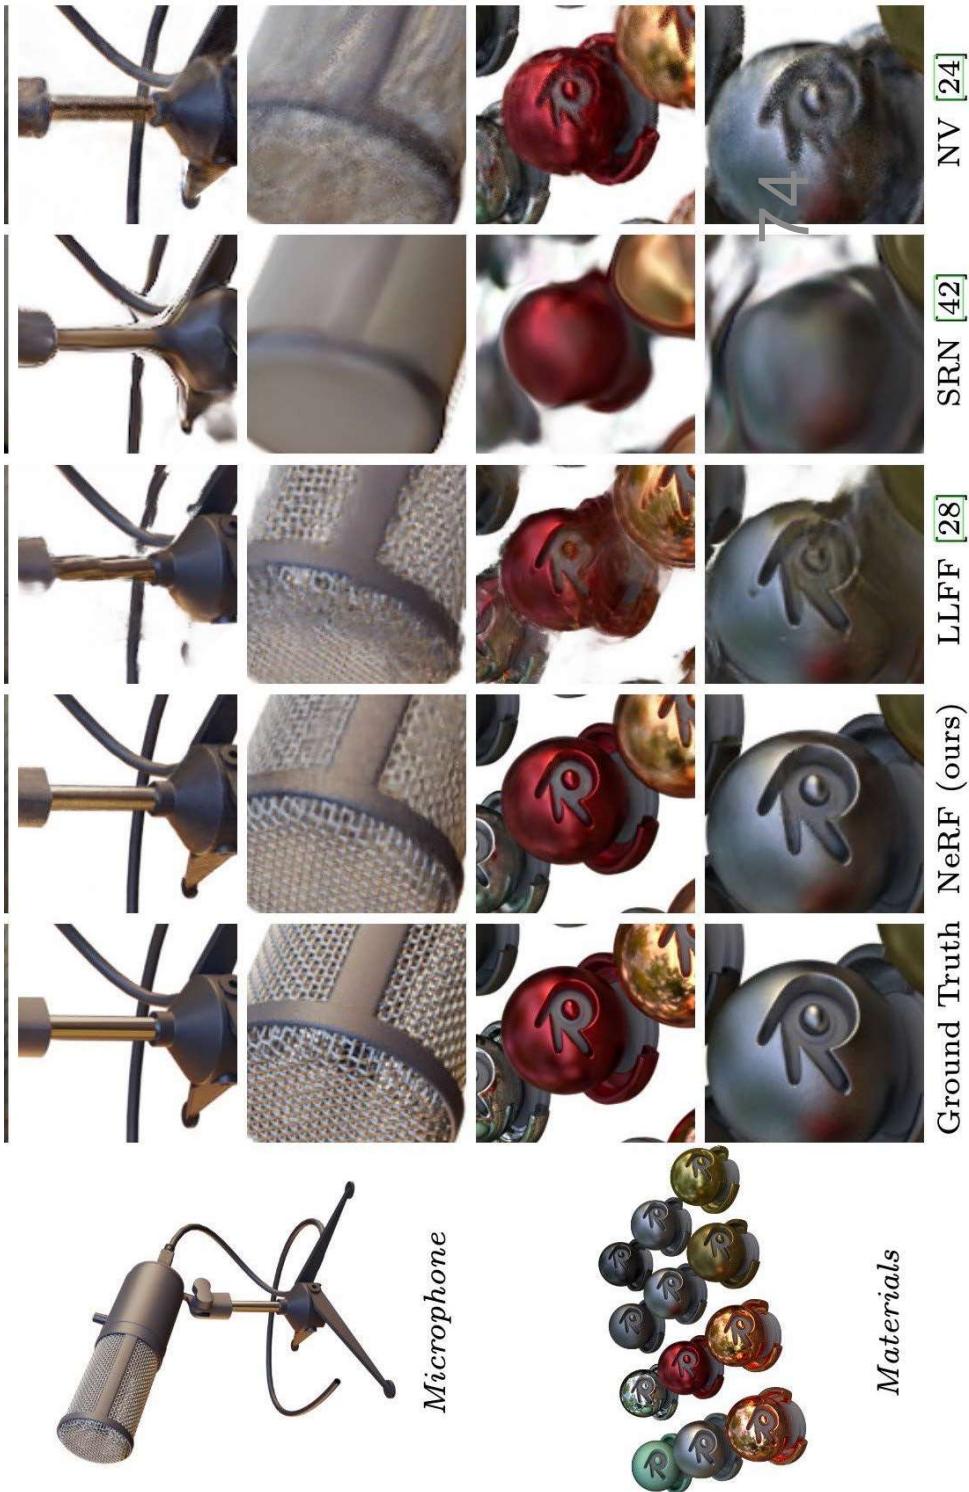
# Network Architecture



# Training Summary

- ❖ Sample a batch of camera rays from the dataset (bs=4096)
  - ❖ Use hierarchical sampling to query coarse and fine points
  - ❖ Use the volume rendering equation to calculate the color of the ray
  - ❖ Compute the squared error between rendered and true pixel colors
- $$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left[ \left\| \hat{C}_c(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 + \left\| \hat{C}_f(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 \right]$$
- ❖ Optimize network parameters using Adam

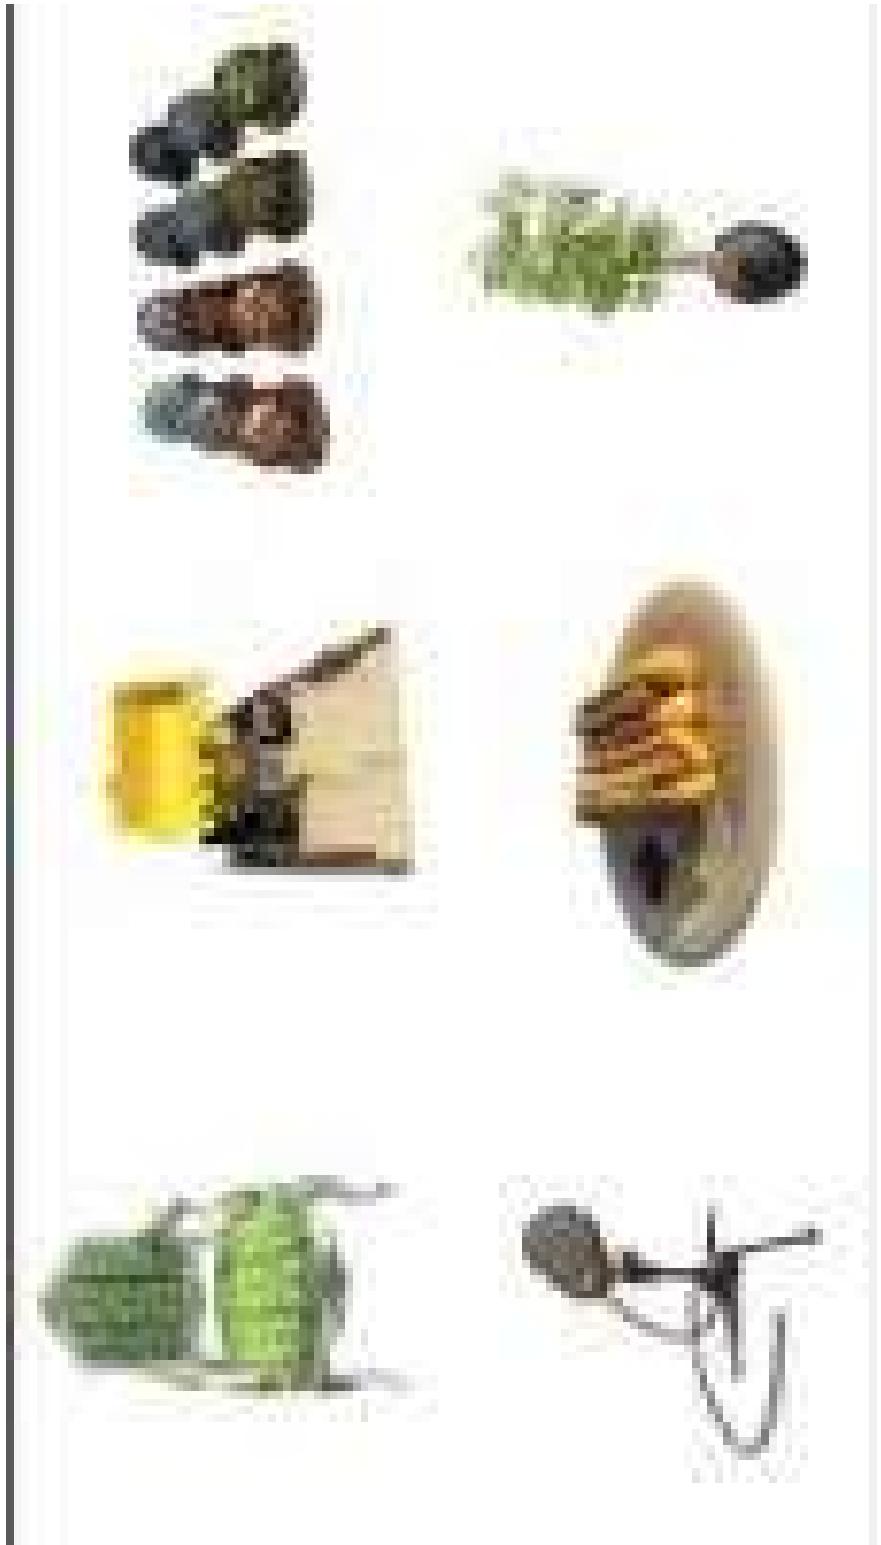
# Qualitative Results: Realistic Synthetic Objects



# Qualitative Results: Real-world scenes



# Animated Results



Video: <https://www.matthewtancik.com/nerf>

# Instant Neural Graphics Primitives



<https://github.com/NVlabs/instant-ngp>

# Limitations

- ❖ NeRF only works with static scenes
- ❖ A trained NeRF model does not generalize to more than one scene
- ❖ Computationally expensive: 1-2 days to train each individual scene on a modern GPU
- ❖ Inference is slow: each pixel in a synthesized image requires volume rendering

# Summary

- ❖ Novel view synthesis generates images of scenes at previously unseen viewpoints.
- ❖ Prior works are limited to simple shapes and do not scale well to high-resolution images
- ❖ NeRF encodes a static scene within the parameters of a feedforward neural network.
- ❖ The authors show very impressive qualitative results and show state-of-the-art performance with quantitative metrics and different scene types.

# D-NeRF: Neural Radiance Fields for Dynamic Scenes

## NeRF

- Only applicable to rigid scenes
- 5D continuous function
- Requiring multiple views of a rigid scene

## D-NeRF

- + Applicable for rigid and non-rigid scenes
- + 6D continuous function by considering time-component as an additional input
- + Requiring a single view per time instant for non-rigid scenes.

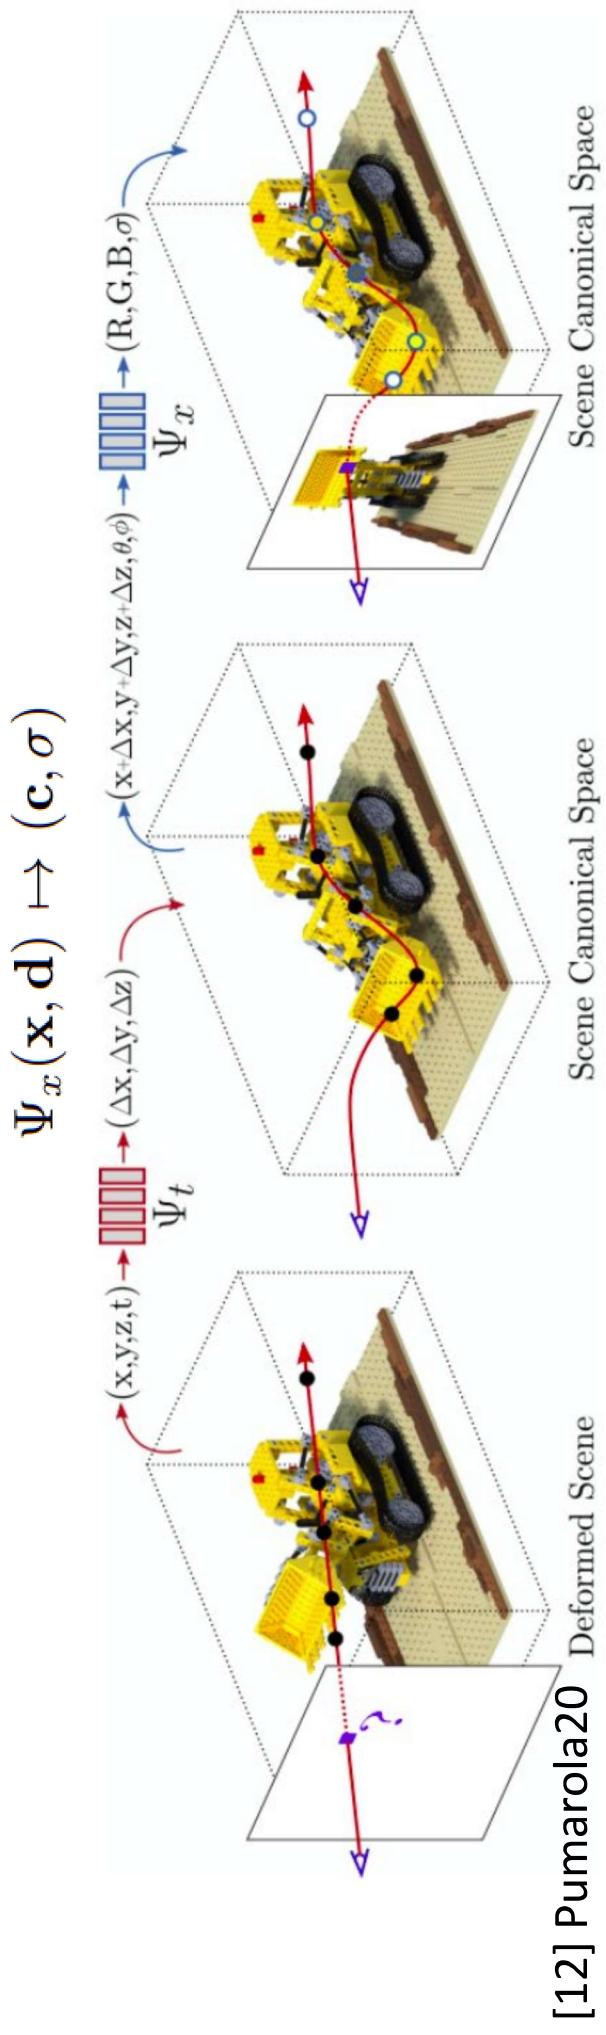


[12] Pumarola20

Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC

# D-NeRF: Neural Radiance Fields for Dynamic Scenes

- **Deformation network**  $\Psi_t$  predict deformation field between the scene at time instant  $t$  and the scene in canonical space ( $t=0$ )  
$$\Psi_t(\mathbf{x}, t) = \begin{cases} \Delta \mathbf{x}, & \text{if } t \neq 0 \\ 0, & \text{if } t = 0 \end{cases}$$
- **Canonical network**  $\Psi_x$  predict color and density in canonical configuration



# D-NeRF: Neural Radiance Fields for Dynamic Scenes

Volumetric rendering is the same as NeRF in **canonical space**:



Predicted colors

$$C(p, t) = \int_{h_n}^{h_f} \mathcal{T}(h, t) \sigma(\mathbf{p}(h, t)) \mathbf{c}(\mathbf{p}(h, t), \mathbf{d}) dh$$

where  $\mathbf{p}(h, t) = \mathbf{x}(h) + \Psi_t(\mathbf{x}(h), t)$ ,

$$[\mathbf{c}(\mathbf{p}(h, t), \mathbf{d}), \sigma(\mathbf{p}(h, t))] = \Psi_x(\mathbf{p}(h, t), \mathbf{d}),$$
$$\text{and } \mathcal{T}(h, t) = \exp\left(-\int_{h_n}^h \sigma(\mathbf{p}(s, t)) ds\right).$$

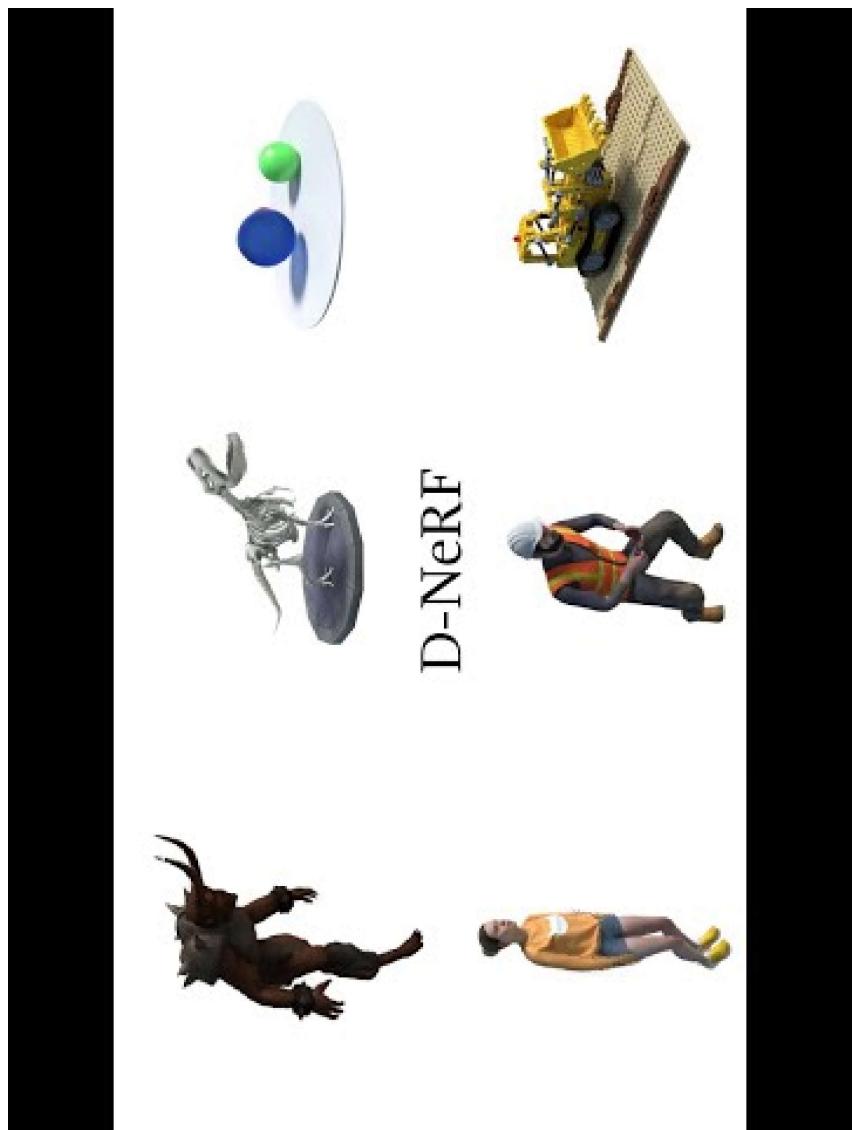
Opacity

$$C(p, t) = \int_{h_n}^{h_f} \mathcal{T}(h, t) \sigma(\mathbf{p}(h, t)) \mathbf{c}(\mathbf{p}(h, t), \mathbf{d}) dh$$

Volume density

[12] Pumarola20

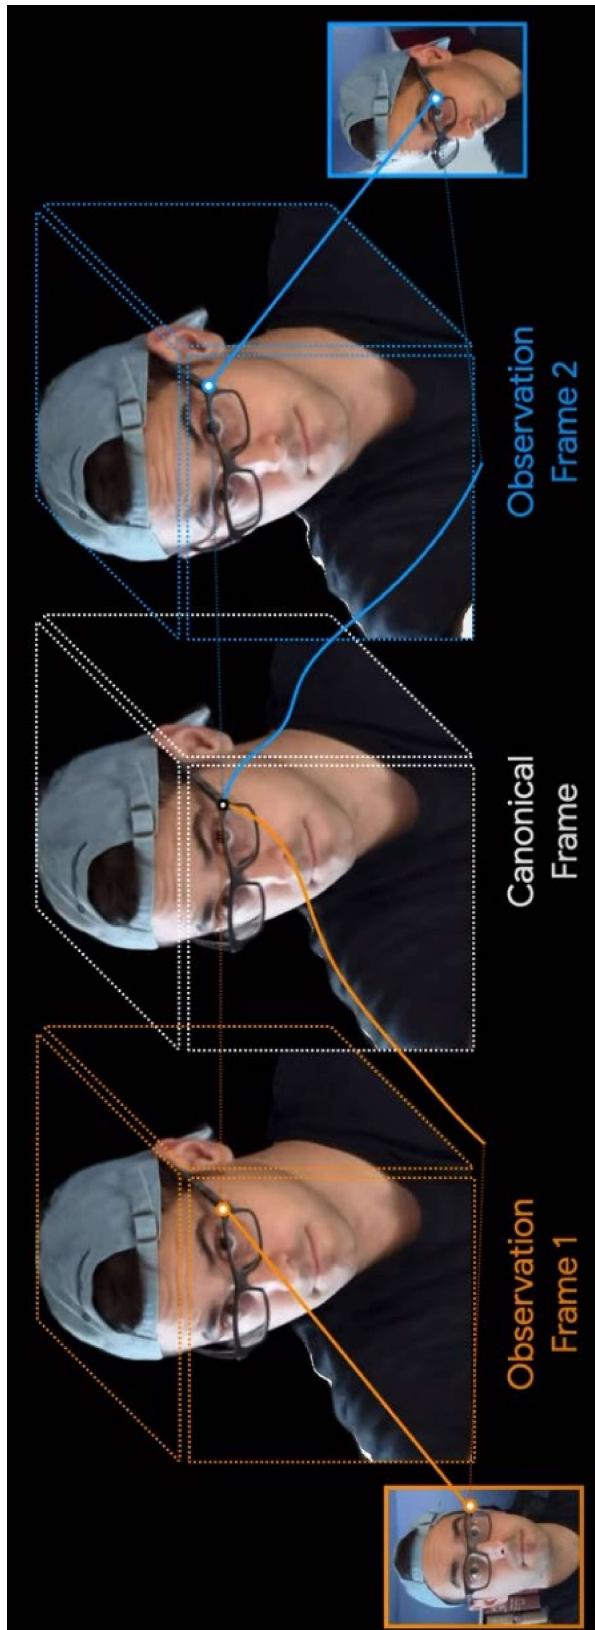
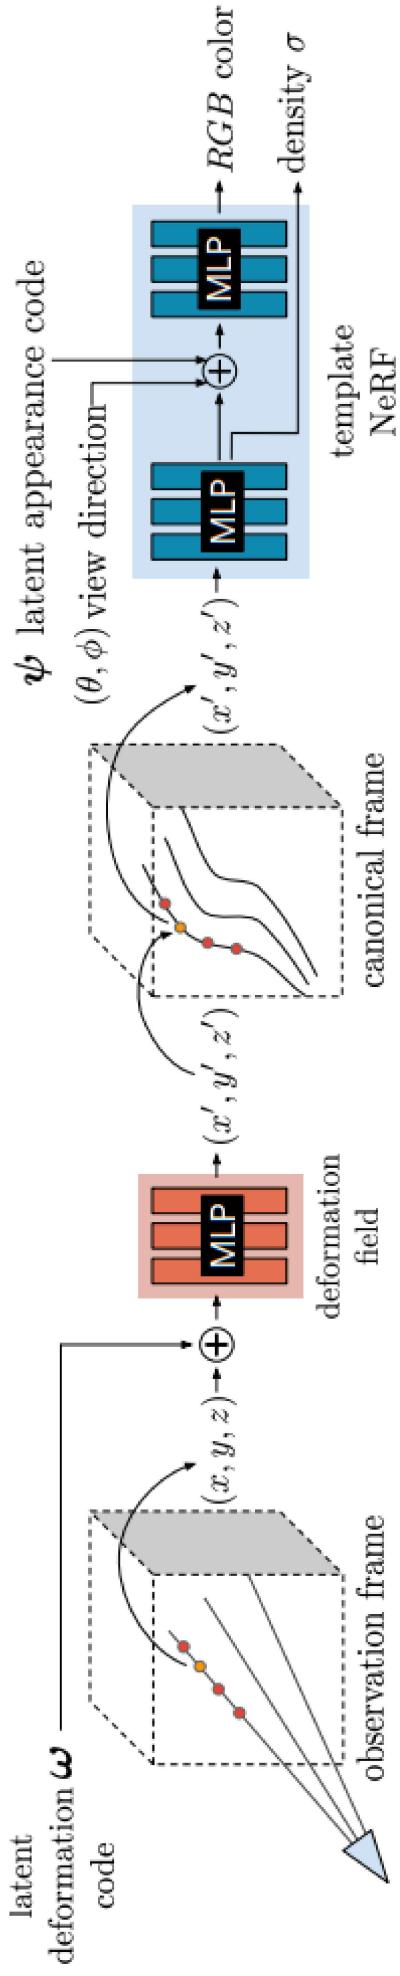
# D-NeRF: Neural Radiance Fields for Dynamic Scenes



[12] Pumarola20

Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC

# Deformable Neural Radiance Fields



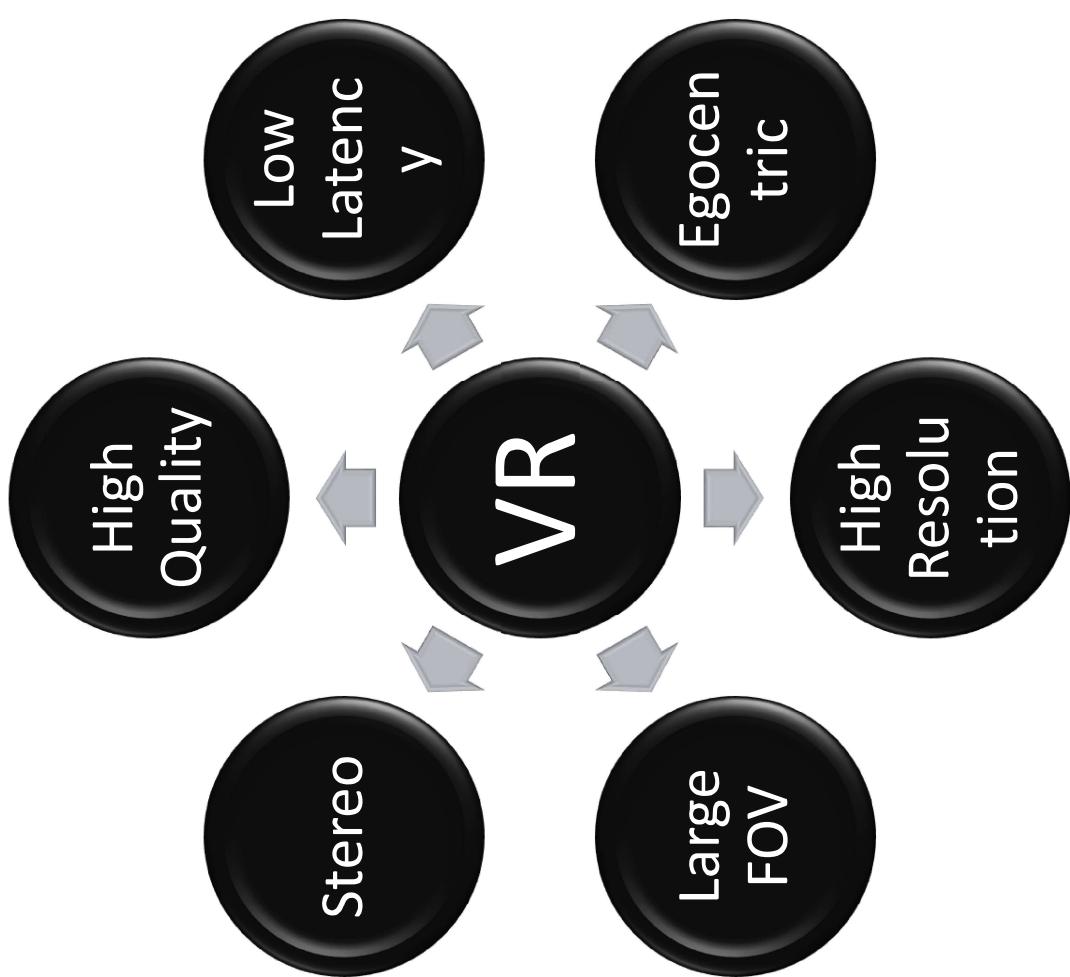
Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC

# Deformable Neural Radiance Fields



[13] Park20

Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC



NeRF

High Latency

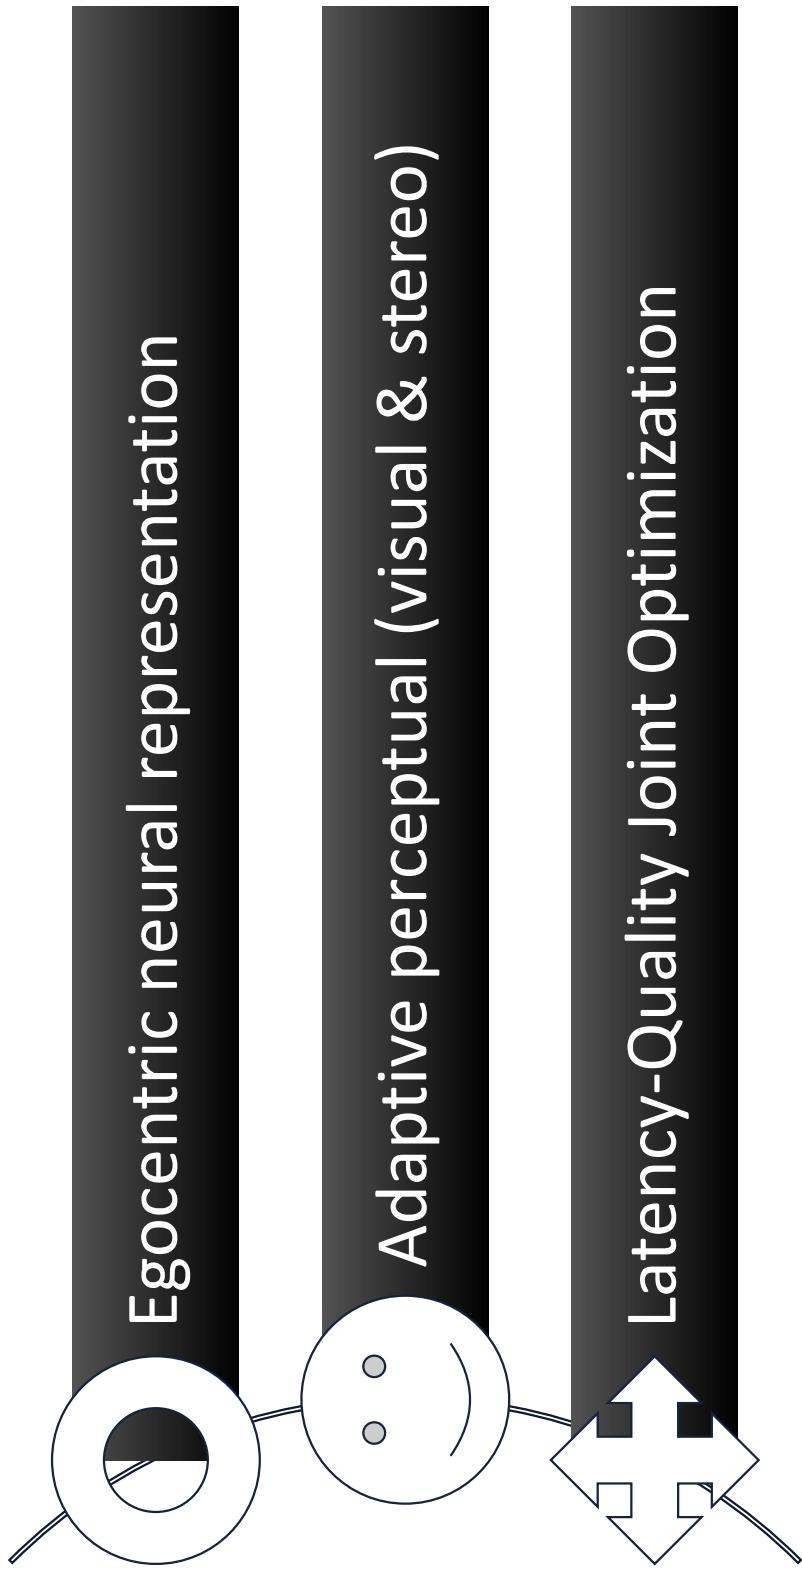
Low Quality



Egocentric

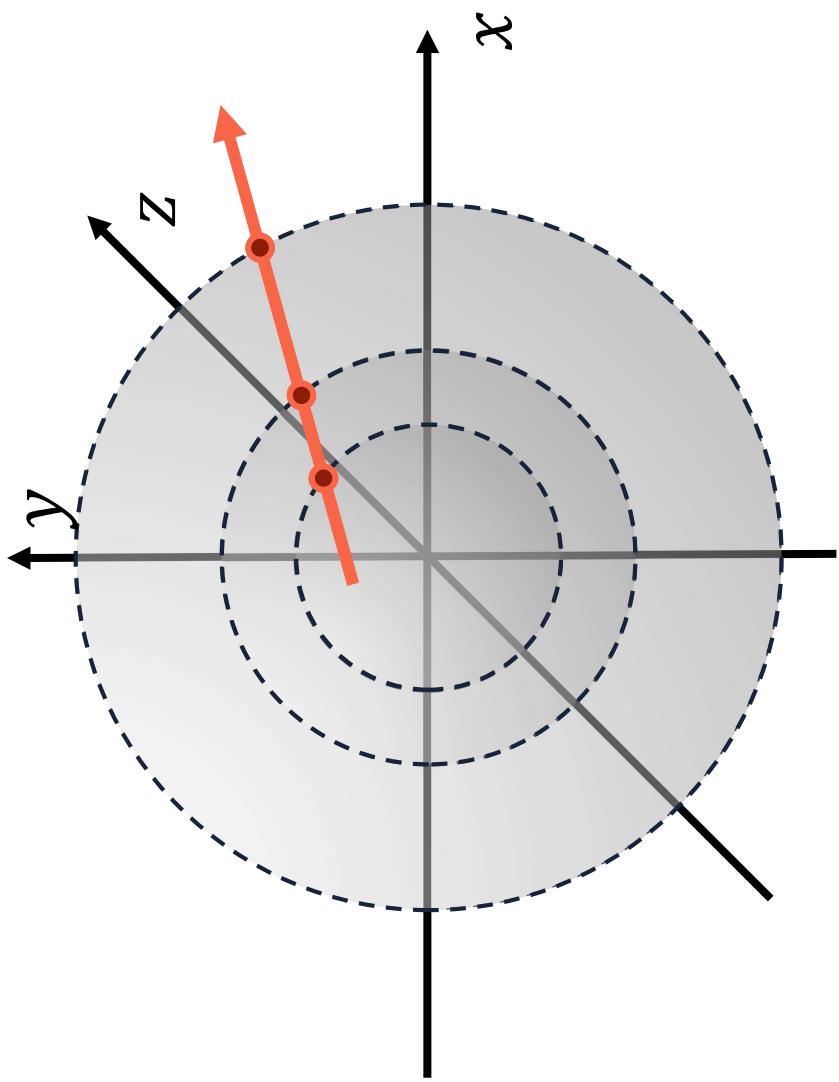
VR

# Gaze-Contingent Neural Rendering for VR

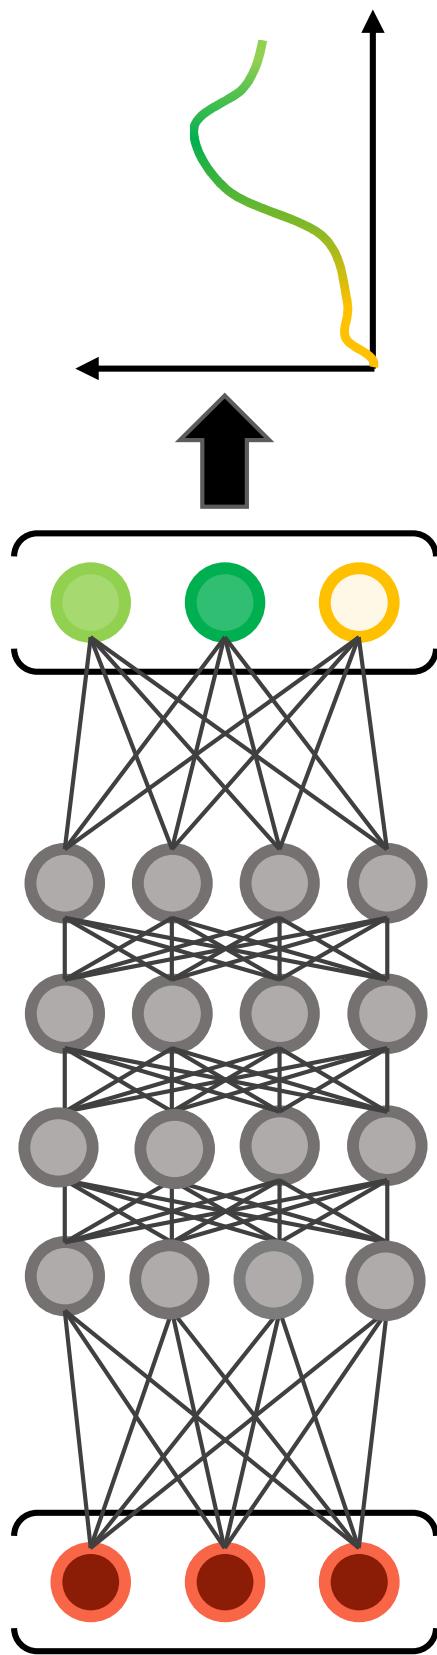




# #1 Egocentric Neural Representation



## #1 Egocentric Neural Representation



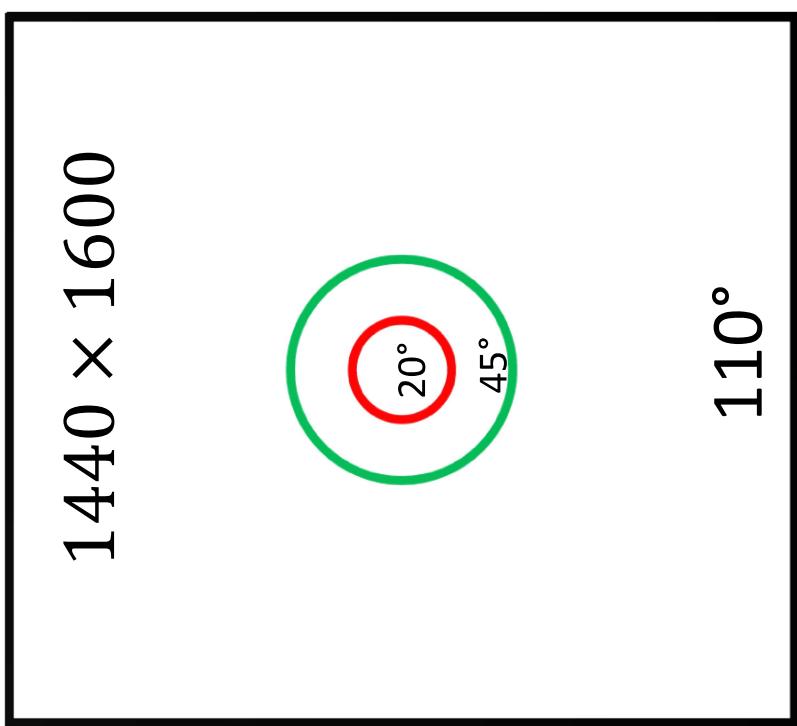
Higher quality

Faster

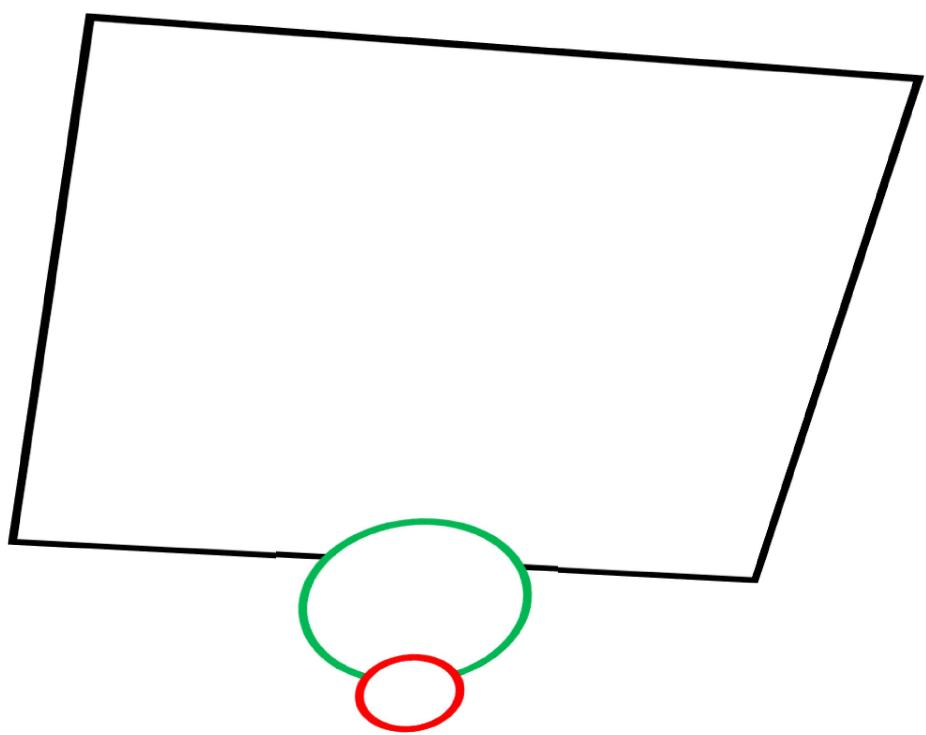
## #2. Adaptive Visual Acuity



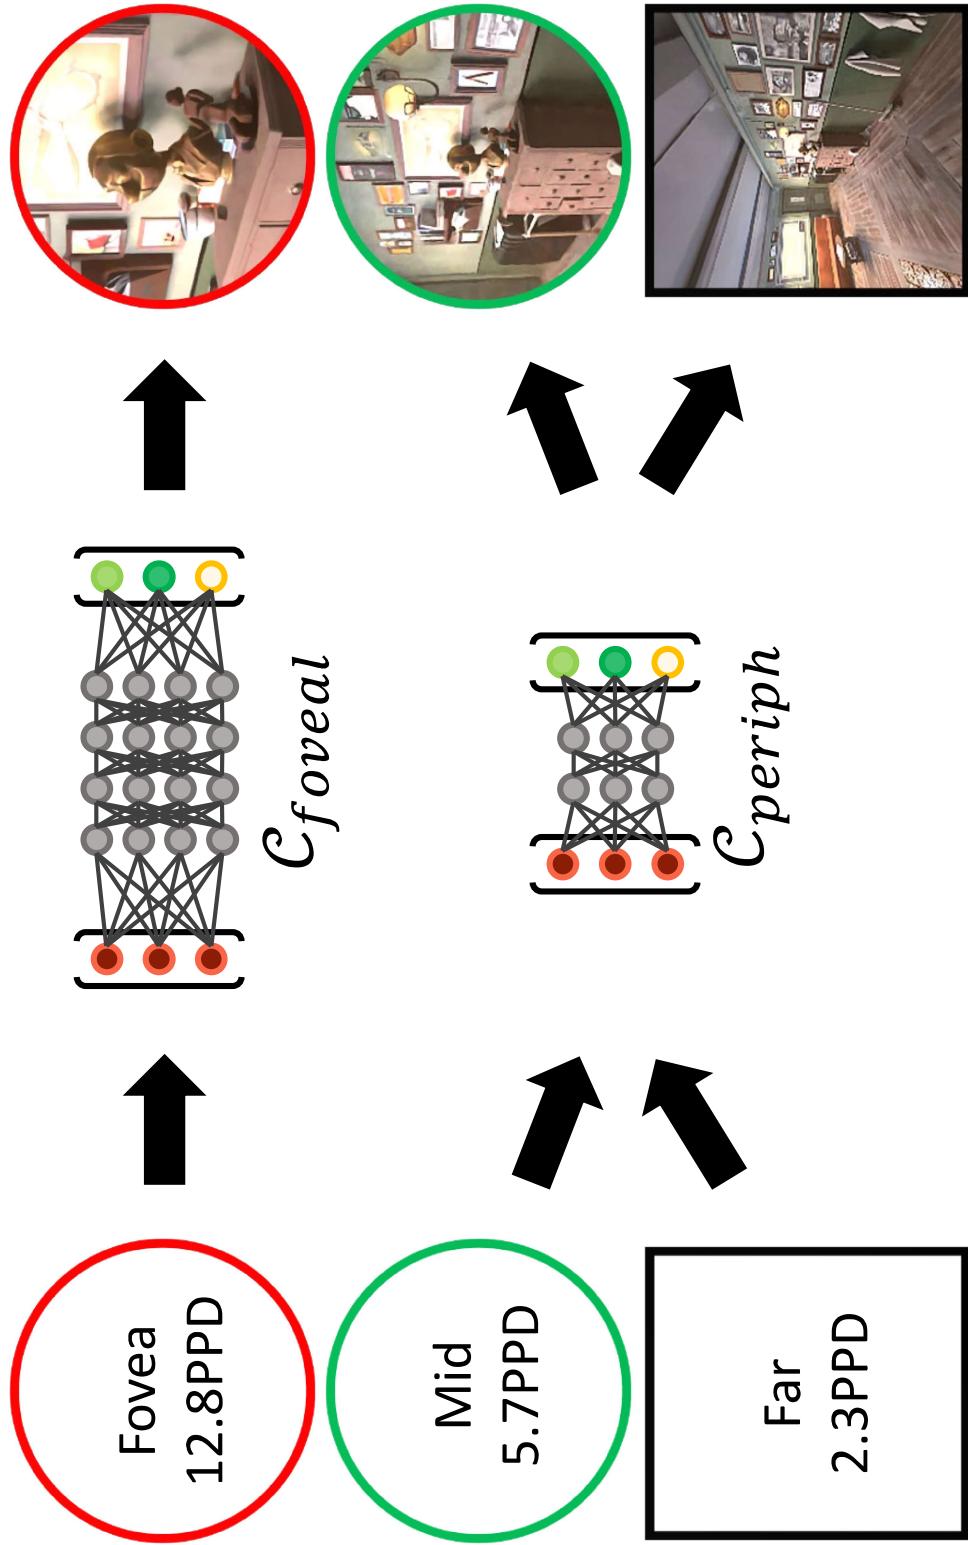
## #2. Adaptive Visual Acuity



## #2 Adaptive Visual Acuity



## #2 Adaptive visual Acuity



## #2 Adaptive Visual Acuity

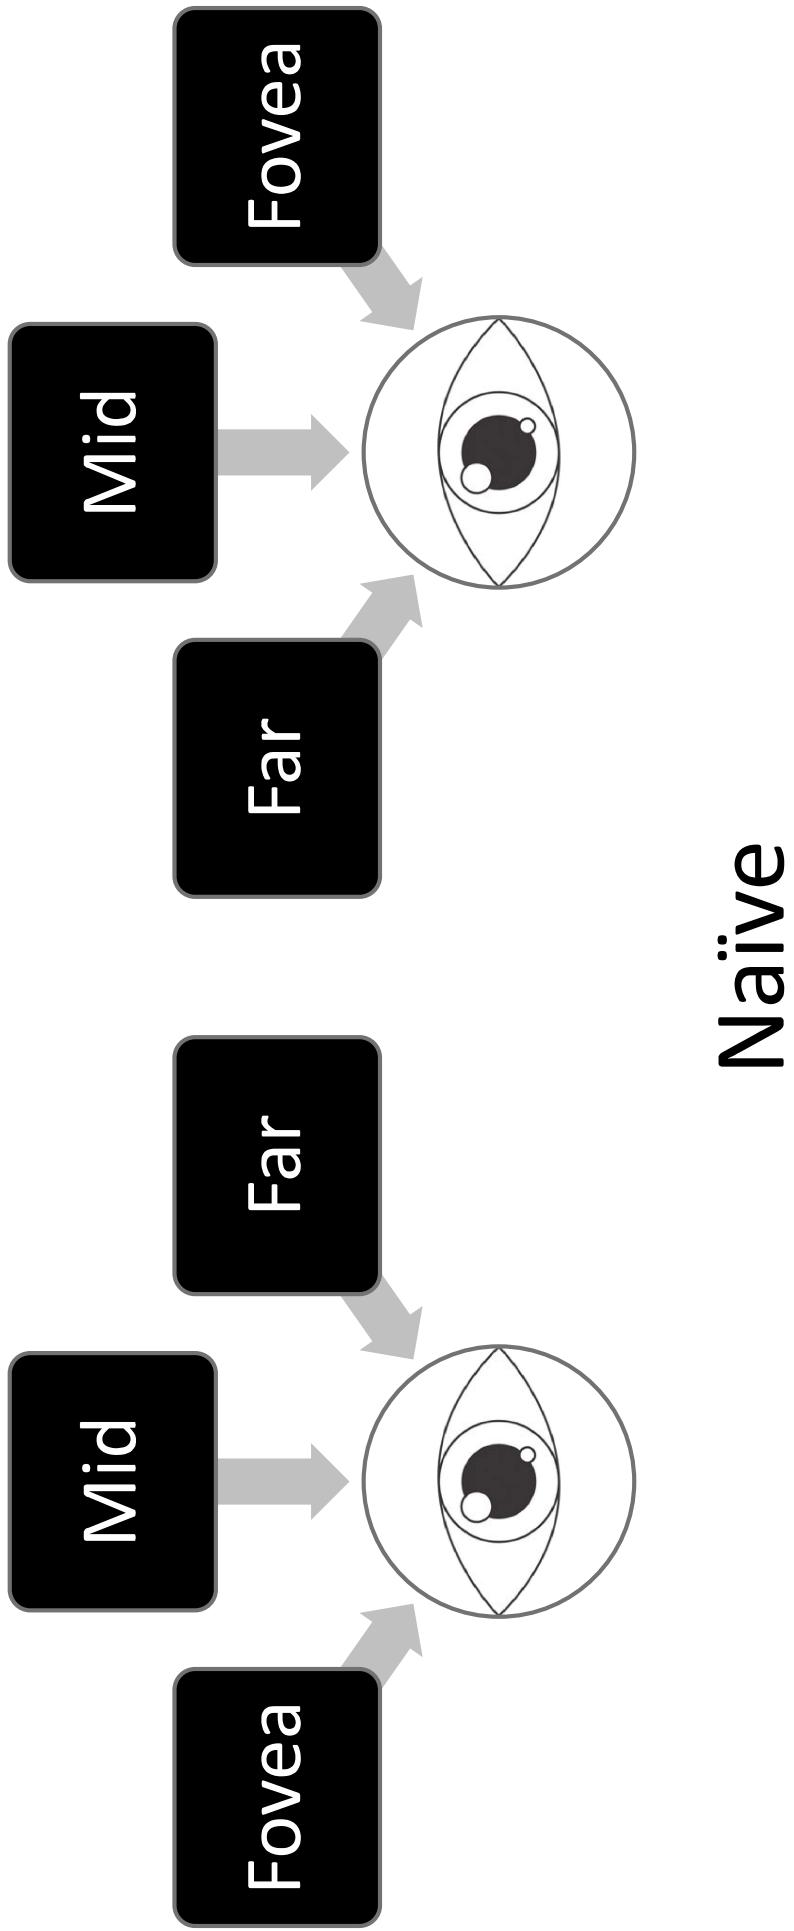


Save

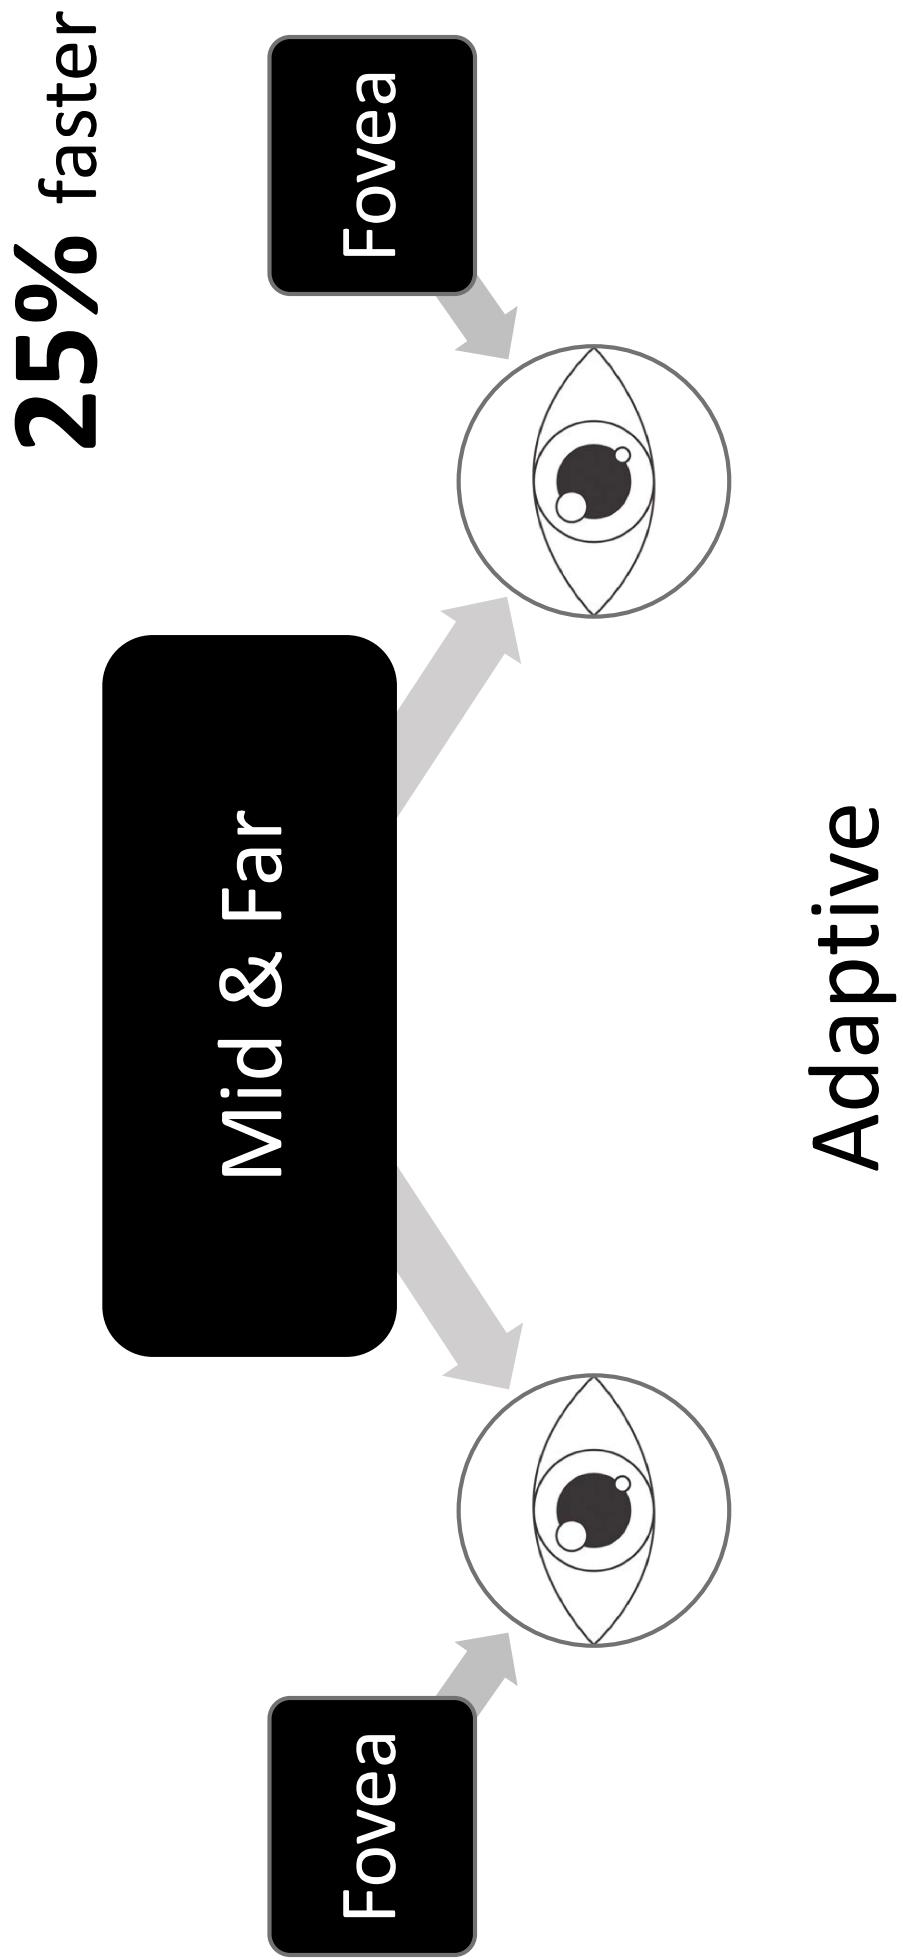
~98%

computational cost

### #3 Adaptive Stereoscopic Acuity



### #3 Adaptive Stereoscopic Acuity



## #4 Latency-Quality Joint Optimization

Hyper parameters:

$N$ : # of networks

$N_m$ : # of layers per network

$N_c$ : # of channels per layer

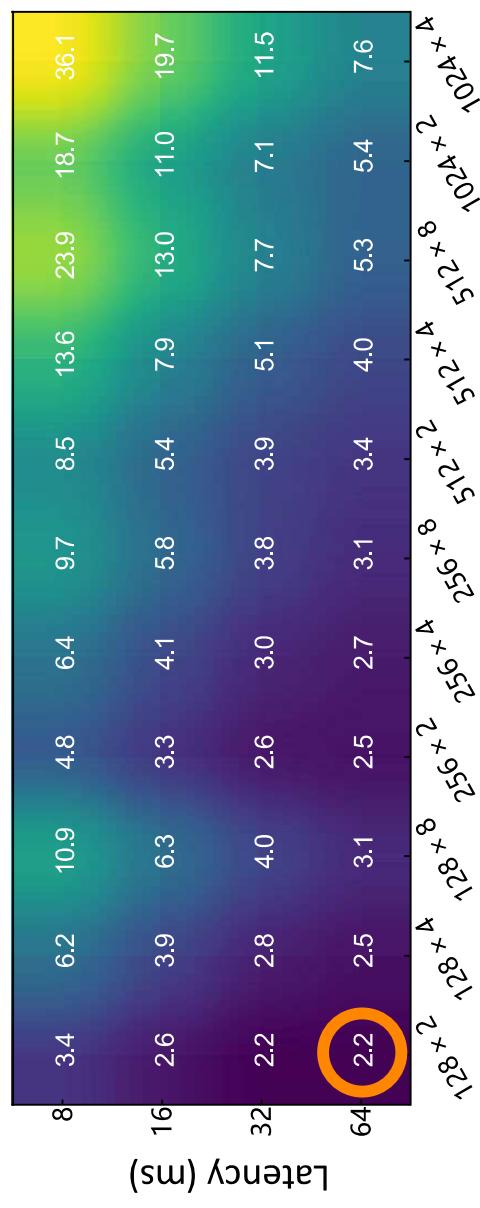
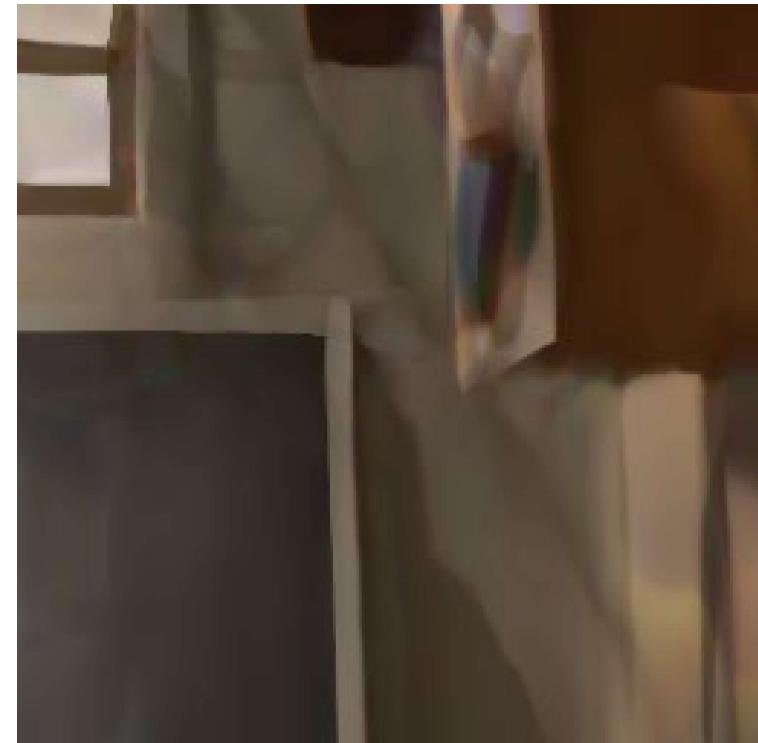
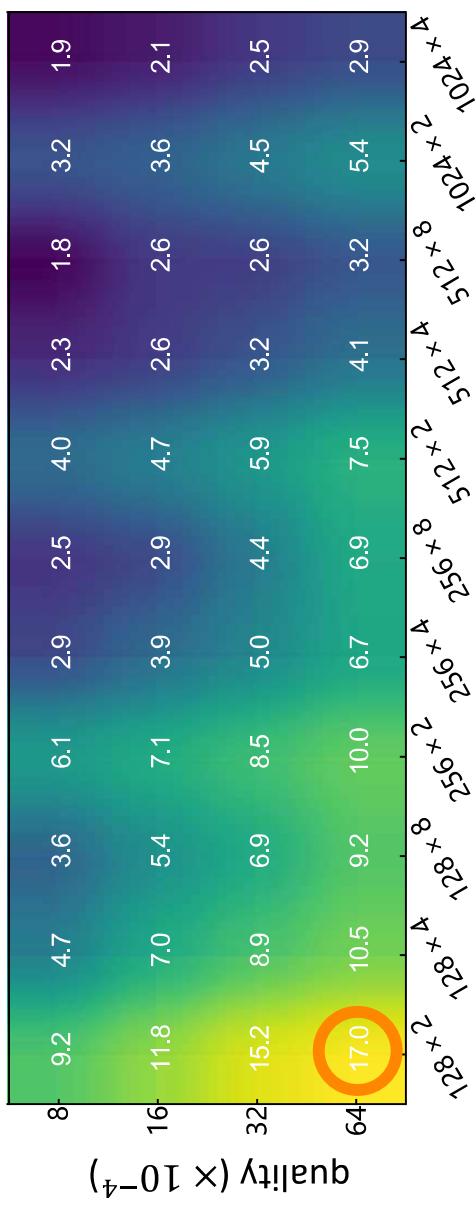
Spatial-temporal modeling:

$$E(N, N_m, N_c) = \sum_t \int \mathcal{C}_{N_m, N_c}(q) \times \\ \| \mathbf{M}(\mathbf{x}_t, \mathbf{R}_t) \cdot \mathbf{q} - \mathbf{M}(\mathbf{x}_{t-l}, \mathbf{R}_{t-l}) \cdot \mathbf{q}(r_k, \mathbf{x}_{t-l}, \mathbf{v}_{t-l}) \| \mathbf{d}\mathbf{q}$$

Objective function:

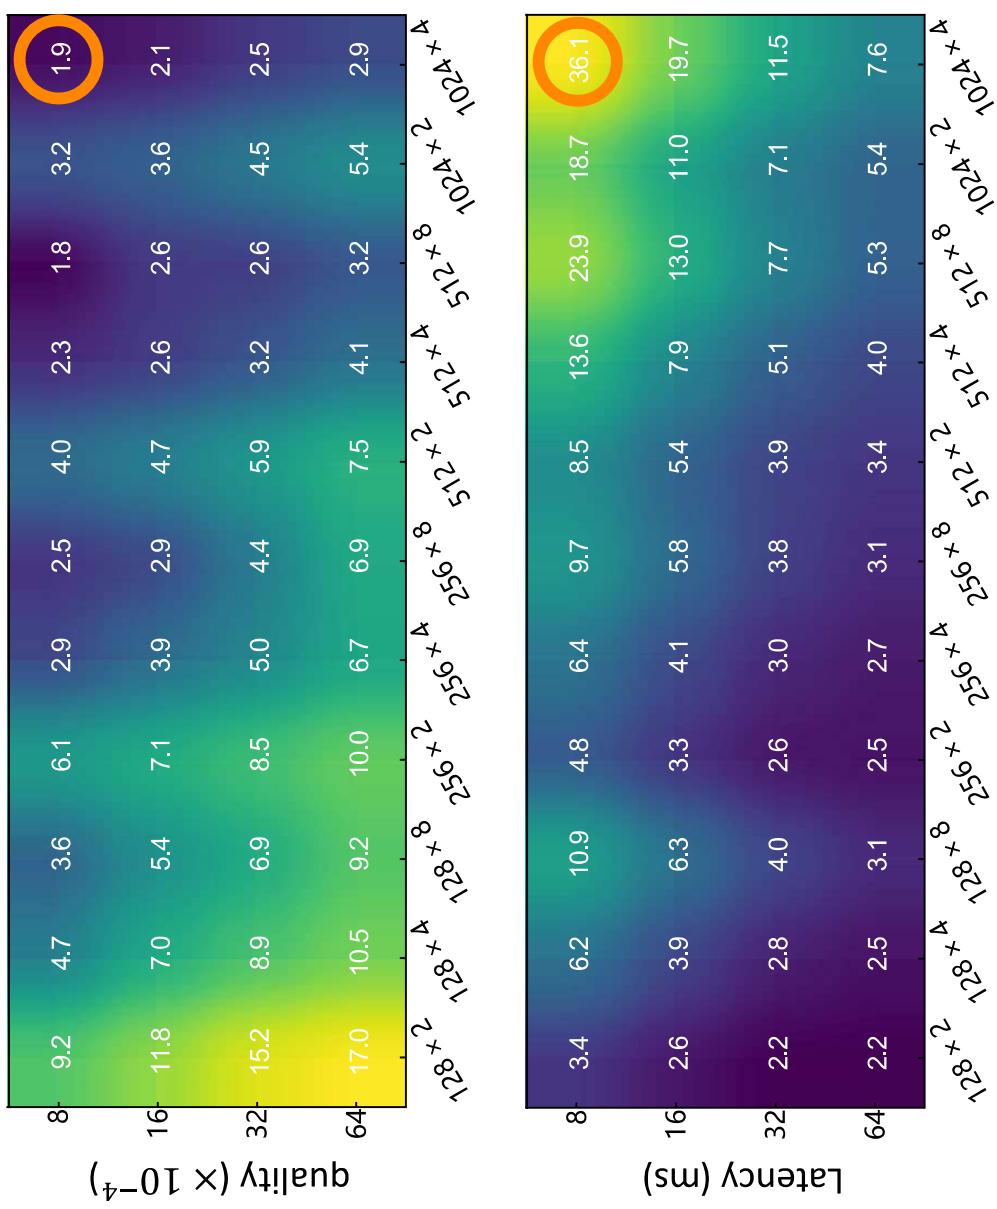
$$\underset{N, N_m, N_c}{\operatorname{argmin}} E(N, N_m, N_c), \quad \text{s.t. } l(N, \mathbf{r}) < L_0$$

## #4 Latency-Quality Joint Optimization

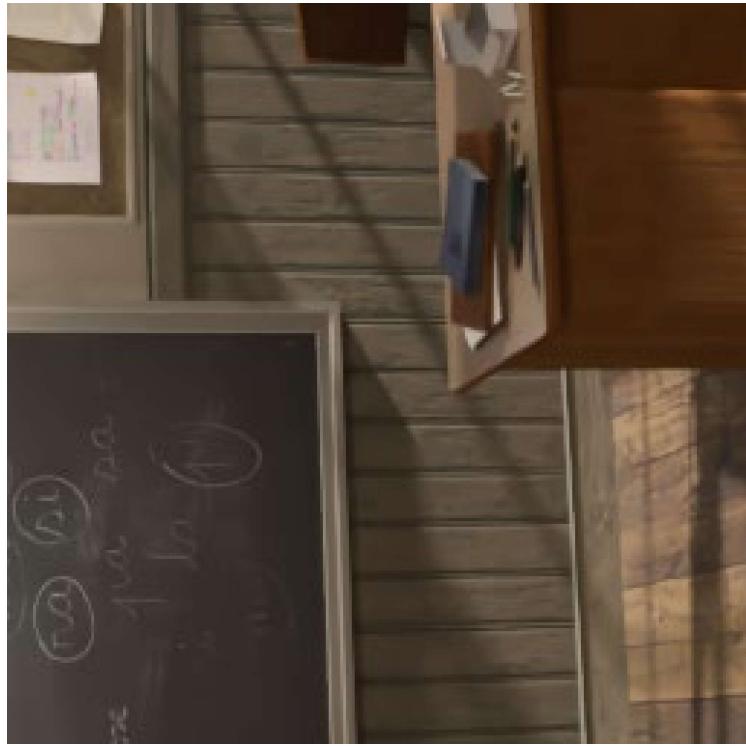


Best Performance

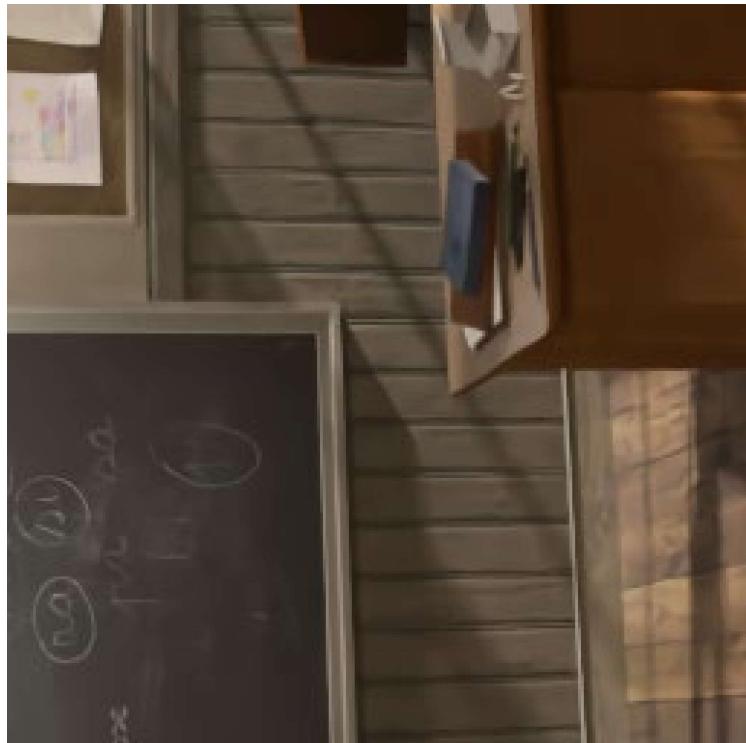
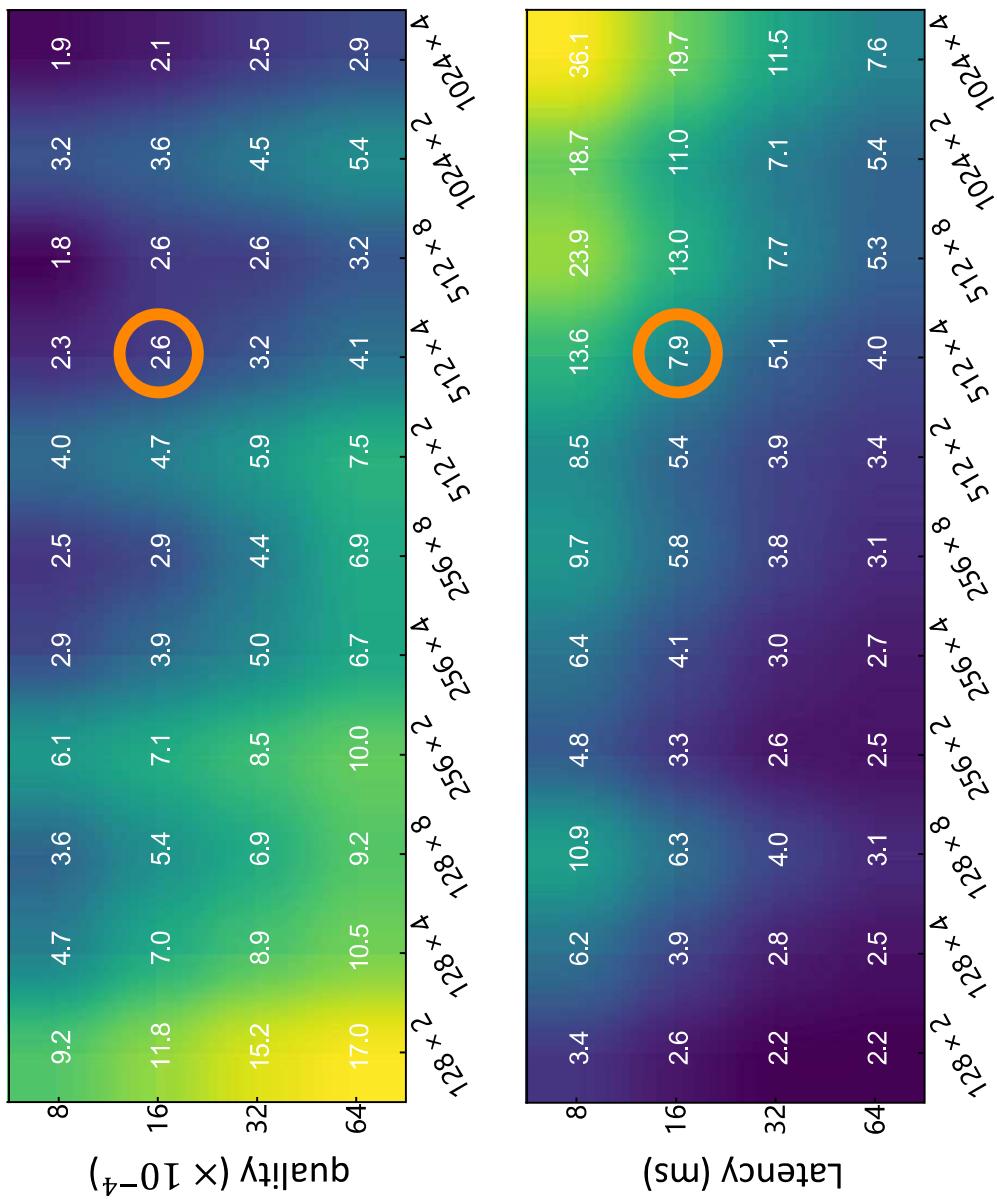
## #4 Latency-Quality Joint Optimization



Best Quality



## #4 Latency-Quality Joint Optimization



Ours



# Explicit vs Implicit Representation (2D)

Explicit:

$$\mathbf{f}(\alpha) = (r \cos(\alpha), r \sin(\alpha))^T$$

Domain:  $[0, 2\pi]$

Implicit:

$$\mathbf{f}([0, 2\pi])$$

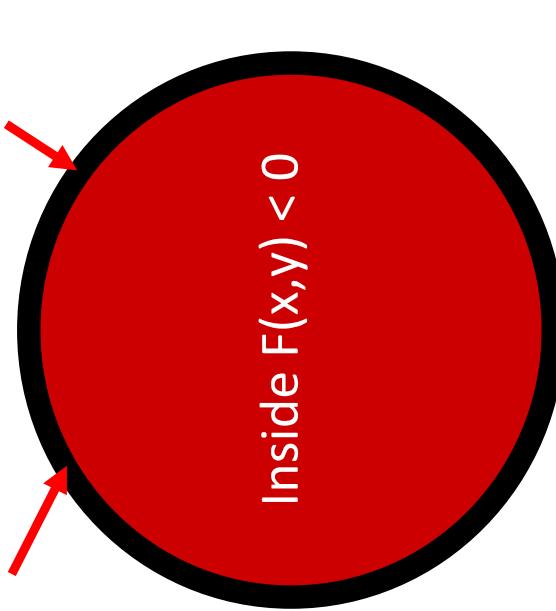
Domain:  $(x, y) \in \mathbb{R}^2$

Implicit:

$$F(x, y) = \sqrt{x^2 + y^2} - r$$

Domain:  $(x, y) \in \mathbb{R}^2$

$\Rightarrow$  Circle is implicitly defined by  $\{(x, y) | F(x, y) = 0\}$



Outside  $F(x, y) > 0$

$\mathbf{f}(\alpha)$  defines the interface  
 $F(x, y)$  defines the **Signed Distance Function** of the circle

# Explicit vs Implicit Representation (3D)

Explicit:

$$\mathbf{f}(\alpha, \beta) = (r \sin(\alpha) \cos(\beta), -r \cos(\alpha) \sin(\beta), r \sin(\alpha) \sin(\beta))$$

Domain:  $\alpha \in [0; 2\pi]$ ,  $\beta \in [0; \pi]$

Implicit

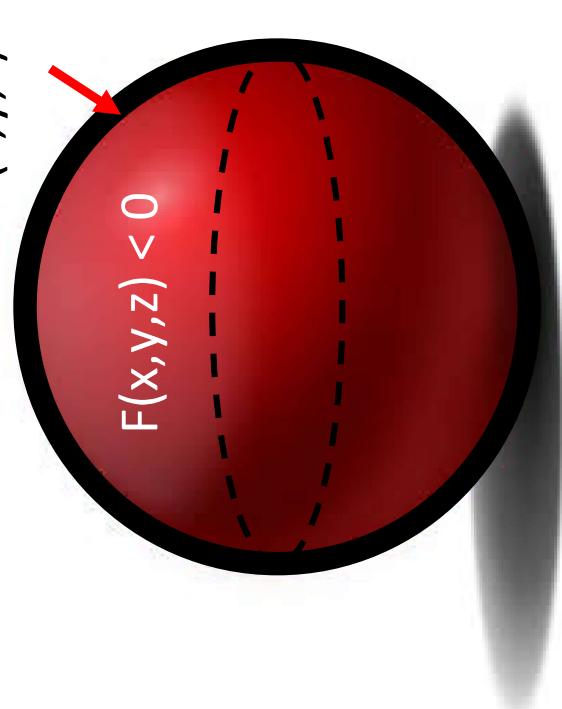
$$F(x, y, z) = \sqrt{x^2 + y^2 + z^2} - r$$

Domain:  $(x, y, z) \in \mathbb{R}^3$

$\Rightarrow$  Sphere is implicitly defined by  $\{(x, y, z) | F(x, y, z) = 0\}$

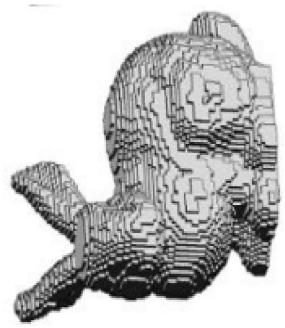
$\mathbf{f}(\alpha, \beta)$  defines the 3D surface

$F(x, y, z)$  defines the **Signed Distance Function** of the sphere

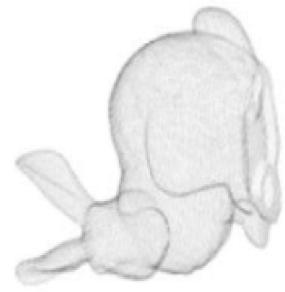


# Representing 3D surfaces

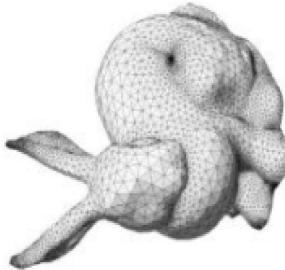
Explicit:



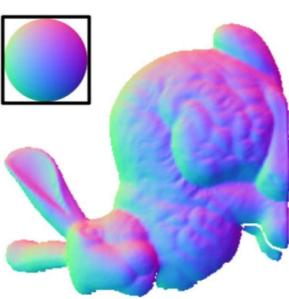
Voxels



Point  
clouds

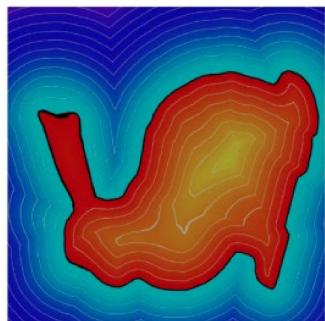


Mesh



Surface  
Normals

Implicit:



Signed distance field  
(e.g gaussian mixtures)

Mixture of primitives  
(e.g gaussian mixtures)

Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC

# Signed Distance Field (SDF)

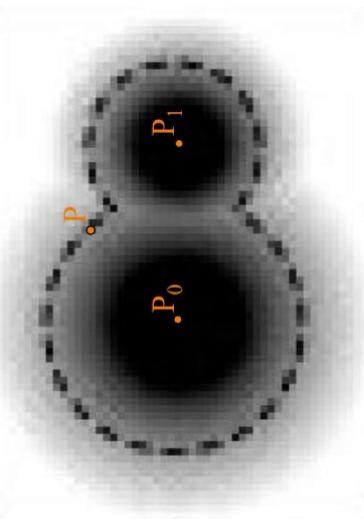
- Maps each 3D points  $p$  to it's signed distance to the object surface  $S$ . The sign is positive if the  $p$  is inside the object, and negative otherwise.
- $SDF(p) = \text{sign}(p) \cdot \min_{q \in S} \|p - q\|$
- Sign indicates whether the point  $p$  is inside (-) or outside (+) of the shape
- Shape's boundary as the zero-level-set of SDF
- Allows for Constructive Solid Geometry (CSG) through boolean operations



Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC

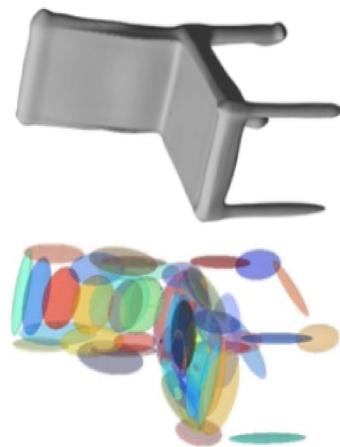
# Mixture of Gaussians

- Represents a shape as a mixture of local implicit functions (3D gaussians)



$$F(\mathbf{x}, \Theta) = \sum_{i \in [N]} f_i(\mathbf{x}, \theta_i)$$
$$f_i(\mathbf{x}, \theta_i) = c_i \exp \left( -\frac{-(\mathbf{p}_{i,d} - \mathbf{x}_d)^2}{2r_{i,d}^2} \right)$$

- Shape's boundary is defined as an iso-level of the **global implicit function**



[1] Genova19  
[2] Genova20

Michaël Ramamonjisoa, Van Nguyen Nguyen, Imagine - ENPC

# Representing 3D surfaces with Implicit Functions

Pros:

- Compared to point clouds: clearly defines the (iso-)surface
- Compared to meshes: can continuously adapt to arbitrary topology
- Compared to voxels: can be represented with few parameters (e.g. mixture of simple implicit functions)
- They are continuous in 3D
- Can give analytic normals, can be applied with boolean operations, etc

# Representing 3D surfaces with Implicit Functions

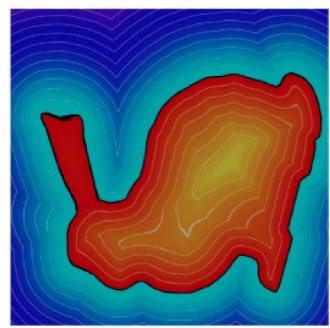
## Pros:

- Compared to **point clouds**: clearly defines the (iso-)surface
- Compared to **meshes**: can continuously adapt to arbitrary topology
- Compared to **voxels**: can be represented with few parameters (e.g. mixture of simple implicit functions)
- They are **continuous** in 3D
- Can give analytic normals, can be applied with boolean operations, etc

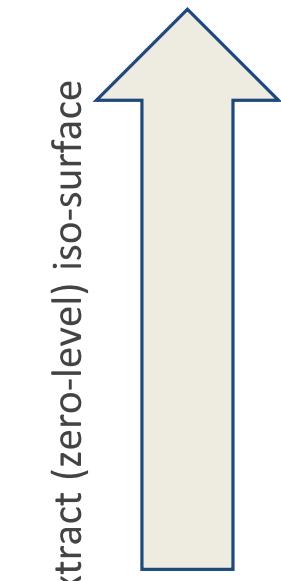
## Cons:

- SDF is well-defined for only watertight meshes (there is an interior and an exterior)
- Need extra steps to visualize

# Converting Implicit Surfaces to meshes

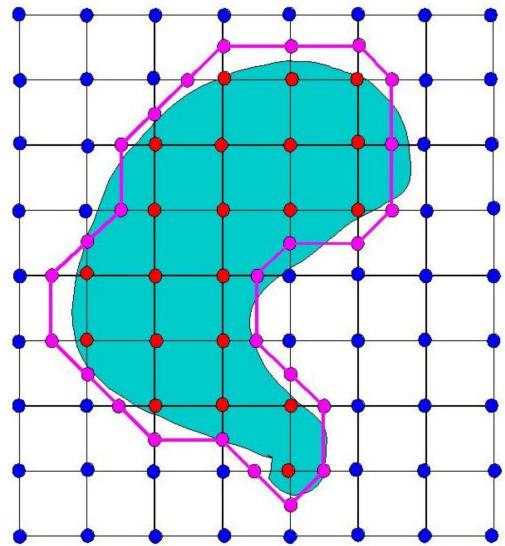


Extract (zero-level) iso-surface

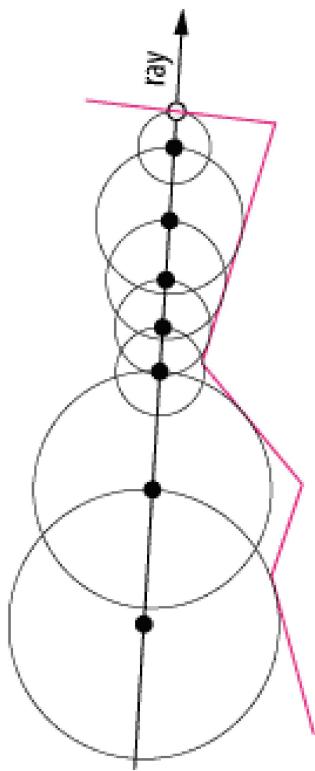


Implicit function

Marching Cubes



Ray marching



Mesh

# Representing 3D surfaces with Implicit Functions

## Pros:

- Compared to point clouds: clearly defines the (iso-)surface
- Compared to meshes: can continuously adapt to arbitrary topology
- Compared to voxels: can be represented with few parameters (e.g. mixture of simple implicit functions)
- They are continuous in 3D
- Can give analytic normals, can be applied with boolean operations, etc

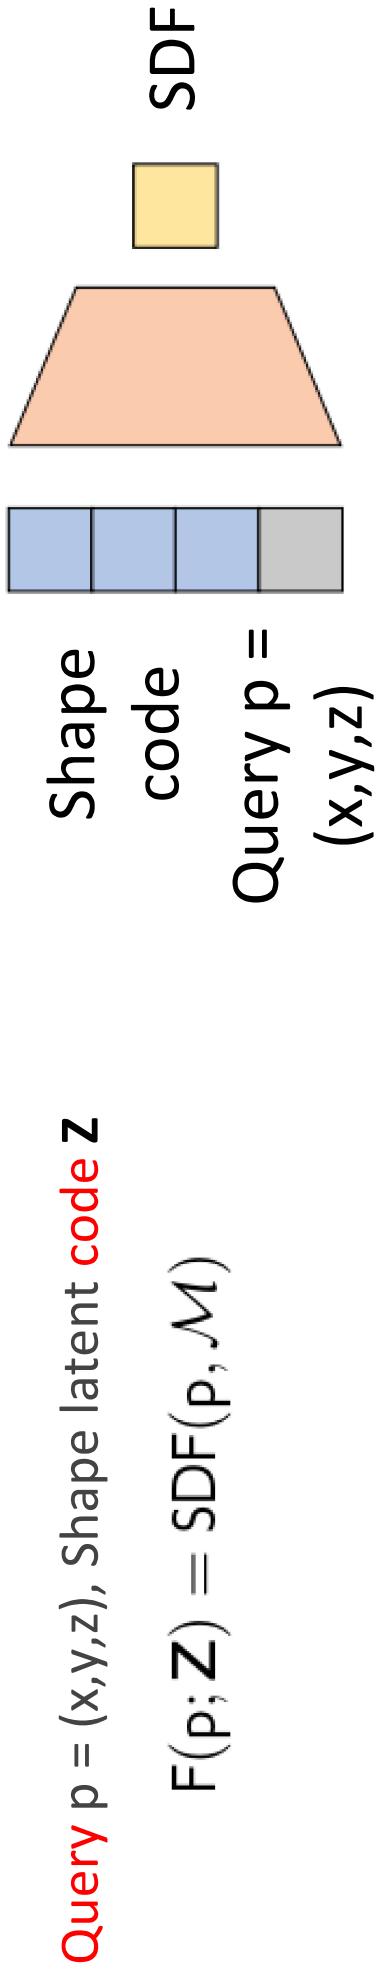
## Cons:

- Implicit functions is well-defined for only watertight meshes (there is an interior and an exterior)
- Need extra steps to visualize
- Not all complex shapes can be efficiently / accurately represented with simple primitives

# Representing 3D surfaces

DeepSDF: Efficiently representing complex shapes by learning their SDF

Idea: Learn a *continuous* representation of 3D implicit surfaces

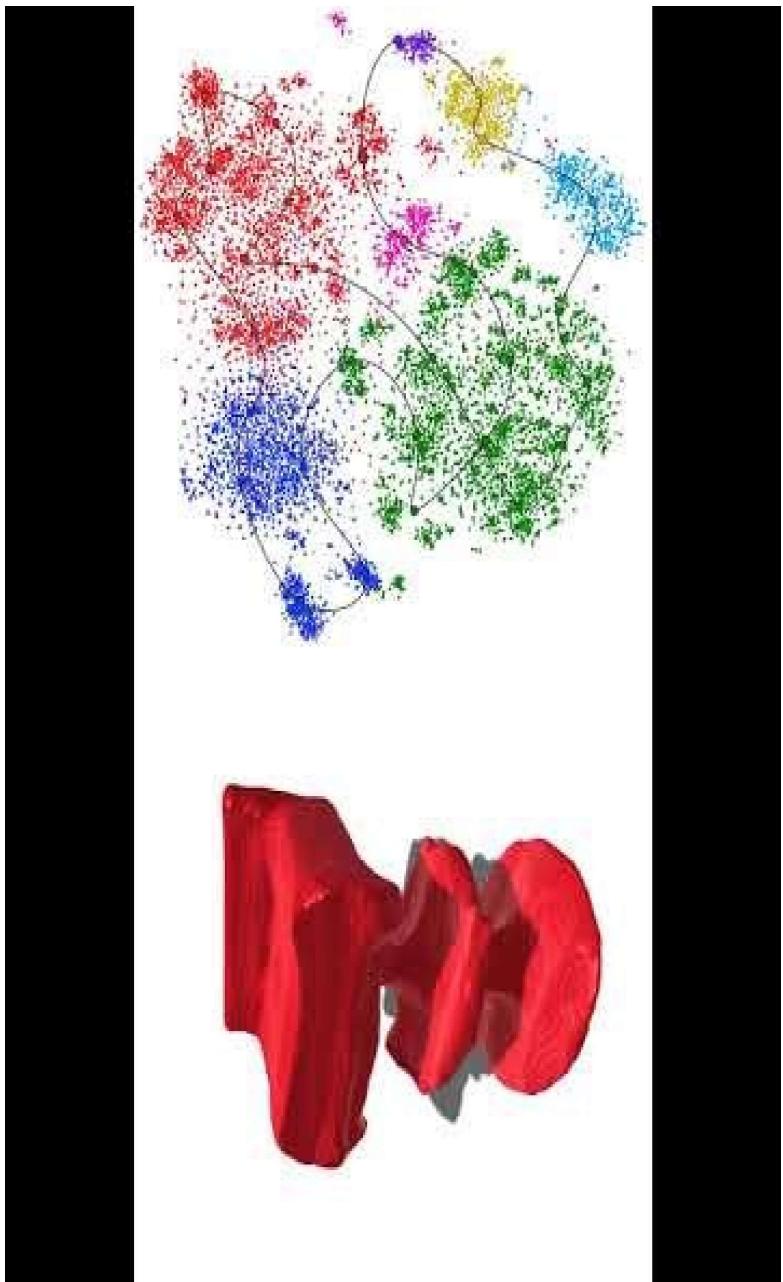


=> Continuity in 3D space **AND** shapes space

[3] Park19

# Representing 3D surfaces

DeepSDF: Representing complex shapes by learning their SDF



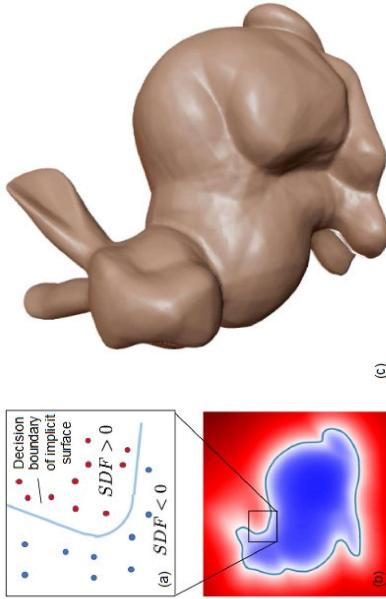
[3] Park19

# Take home message on Implicit Functions

Representation of a continuous field

**Learned implicit functions:**

- Can represent complex shapes
- Are **continuous mappings** because they use **MLPs**
- Are applicable to N-D data: 2D images, 3D shapes, radiance fields

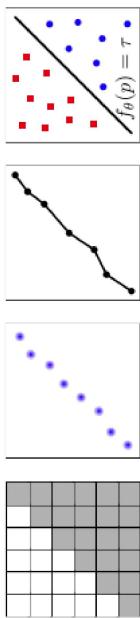


Visualization of implicit functions is done by extracting iso-surfaces:

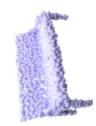
1. Running inference for multiple queries in input space
2. Rendering the result by combining the queries

# More works on Implicit Functions for 3D shape

- Occupancy Networks



[4] Mescheder19



[5] Saito19  
[6] Saito20

- PiFu and PiFuHD



# References

- [1] Genova et al., [Learning Shape Templates with Structured Implicit Functions](#), ICCV 2019
- [2] Genova et al., [Local Deep Implicit Functions for 3D Shape](#), CVPR 2020
- [3] Park et al., [DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation](#), CVPR 2019
- [4] Mescheder et al., [Occupancy Networks: Learning 3D Reconstruction in Function Space](#), CVPR 2019
- [5] Saito, Huang, Natsume et al., [PIFu: Pixel-Aligned Implicit Function for High-Resolution Clothed Human Digitization](#), ICCV 2019
- [6] Saito et al., [PIFuHD: Multi-Level Pixel-Aligned Implicit Function for High-Resolution 3D Human Digitization](#), CVPR 2020

## Courses and Seminars

Lecture on [Implicit geometry](#)

Lecture on [Implicit surface](#)

Lecture on [Explicit & Implicit Surfaces](#)

[Thomas Funkhouser's talk at 3DGV seminar](#)

[Princeton COS 426, Spring 2014 on Implicit Surfaces & Solid Representations](#)