

Lecture3 Engine Neural Network

1. How does learning happen?

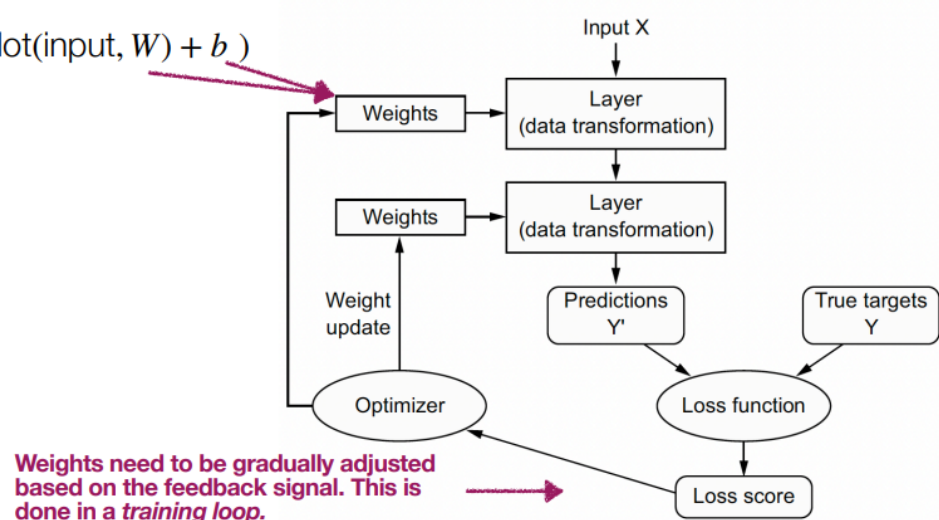
Kernel and Bias

```
[8] from tensorflow import keras
    from tensorflow.keras import layers
    model = keras.Sequential([
        layers.Dense(512, activation="relu"),
        layers.Dense(10, activation="softmax")
    ])
```

$$\text{output} = \text{relu}(\text{dot}(\text{input}, W) + b)$$

- We saw that dense layers are these combination of operations, where W and b are **parameters** of the layer
- Those two matrices are called the **kernel** and **bias** attributes of the dense layer
- Initially, they are **randomized**
 - i.e. assigned random values

$$\text{output} = \text{relu}(\text{dot}(\text{input}, W) + b)$$



Training loop

1. Draw a batch of training samples x , and corresponding labels y_{pred} .



```
[11] model.fit(train_images, train_labels, epochs=5, batch_size=128)
```

Epoch 1/5

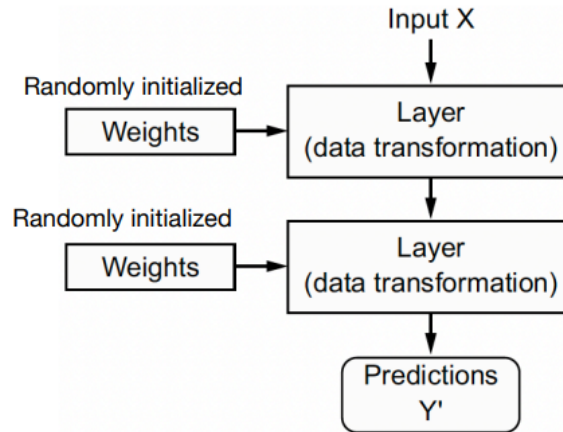
469/469 [=====] - 16s 31ms/step - loss: 0.2553 - accuracy: 0.9262

```

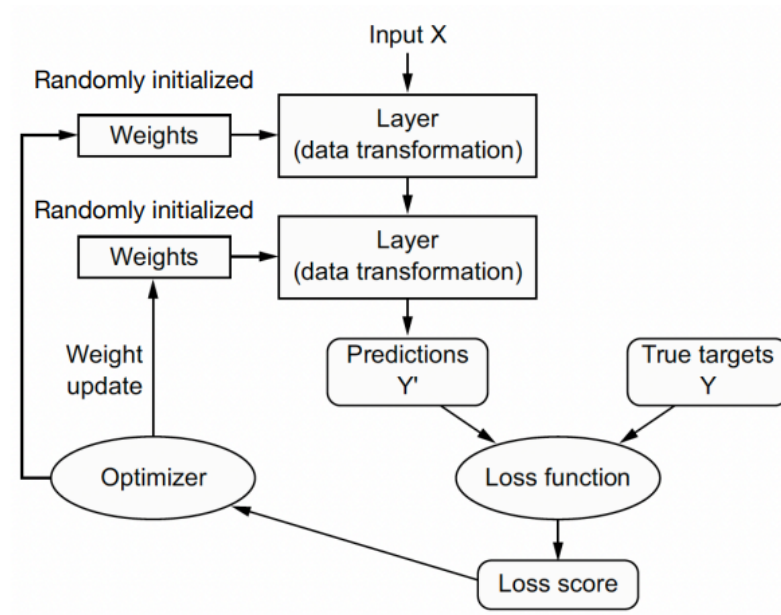
1 batch = train_images[:128]
2 batch = train_images[128:256]
3 n = 3
4 batch = train_images[128*n:128*(n+1)]

```

2. Run the model on `x` (**forward pass**) to obtain predictions `y_pred`



4. Update all weights of the model in a way that slightly reduces the loss on this batch.



2. How to Update Weight?

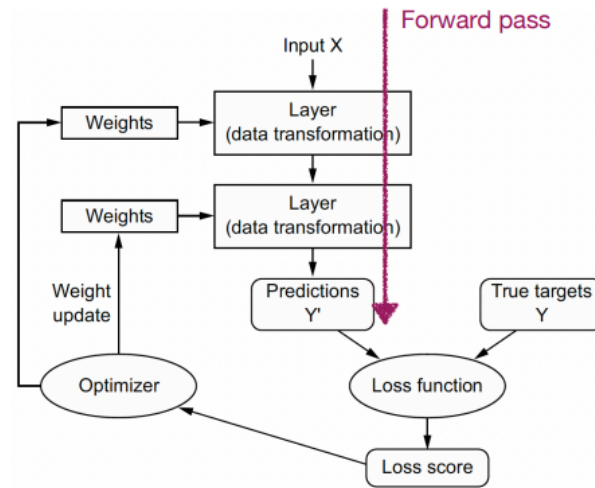
Weight initialization: Randomly

$$W = \begin{bmatrix} 0.30 & 0.45 & 0.10 \\ 0.70 & 0.56 & 0.32 \end{bmatrix} \quad b = \begin{bmatrix} 0.50 \\ 0.21 \end{bmatrix}$$

- A linear layer with
 - 3 inputs
 - 2 outputs

- its weight (w) is a (2×3) matrix (and a bias term (b))
- **This layer wouldn't do anything useful.**
 - (i.e., a randomly initialized model for MNIST classification would show about 10% accuracy)

Update the weights - A naive approach (X)



- Change one coefficient at a time,
 - e.g. increase its value by a little.
- Do the forward pass again, see if the loss improved.


Weight Adjustment	Loss Computation
Weight Initialization	$W = \begin{bmatrix} 0.30 & 0.45 & 0.10 \\ 0.70 & 0.56 & 0.32 \end{bmatrix} \longrightarrow \ell = 0.5$ $b = \begin{bmatrix} 0.50 \\ 0.21 \end{bmatrix}$
Increase the first element in W , loss increases	$W = \begin{bmatrix} 0.35 & 0.45 & 0.10 \\ 0.70 & 0.56 & 0.32 \end{bmatrix} \longrightarrow \ell = 0.6$ $b = \begin{bmatrix} 0.50 \\ 0.21 \end{bmatrix}$
Decrease the first element in W , loss decreases	$W = \begin{bmatrix} 0.25 & 0.45 & 0.10 \\ 0.70 & 0.56 & 0.32 \end{bmatrix} \longrightarrow \ell = 0.4$ $b = \begin{bmatrix} 0.50 \\ 0.21 \end{bmatrix}$

- This is terribly inefficient
- For a middle size network, e.g. 400k parameters, this is 800k forward passes
- In other words, we'd like a method that can **update all the weights** (slightly increase or decrease the value in a way that the model gets **better** (show a lower loss))

Gradient Descent


Invented in 1847 by French mathematician Louis-Augustin Cauchy

It is an optimization technique that powers modern neural networks






- Core idea: The “**gradient**” encodes information about what **small change to each coefficient** (i.e. positive or negative) will **decrease the loss the most**
- Using that information, you can change the coefficients iteratively until finding a **minimum value** for the loss

How



How do we choose in which direction to move?





The rate of change rc is defined as:


$$rc = \frac{\ell(w_0 + \epsilon) - \ell(w_0)}{w_0 + \epsilon - w_0}$$

In this case, $rc > 0$ and we should have decreased the value of w_0 , i.e. subtract a positive number.

What happens if we go the other way (downhill)?

$$rc = \frac{\ell(w_0 + \epsilon) - \ell(w_0)}{w_0 - \epsilon - w_0} = \frac{\ell(w_0) - \ell(w_0 - \epsilon)}{\epsilon}$$

In this case $\ell(w_0) > \ell(w_0 - \epsilon)$ so $rc > 0$. We are going in the right direction, and we should indeed decrease w_0 .




Naive Gradient Descent

```
# initialize iterative process
w_iter = w0
if not w_iter:
    w_iter = np.random.uniform(-2,2,1)[0]

# at each step, subtract the derivative * the learning rate
w_list = [w_iter]
for i in range(0,nsteps):
    w_iter = w_iter - lr* rate_change(w_iter)
    w_list.append(w_iter)

return np.array(w_list)
```



Given a function $\ell(w)$ that you want to minimize with respect to w , do:

1. Initialize $w_{iter} = w_0$ in a random point
2. In a loop, do:

- $w_{iter} = w_{iter} - lr \cdot rc$
- where $rc = \frac{l(w_{iter}+\epsilon) - l(w_{iter})}{\epsilon}$
- lr : **learning rate** (or **step size** in optimization)

3. Repeat step 2 until reach N iterations


3. Gradient

What is Gradient?

$$rc = \frac{\ell(w_0 + \epsilon) - \ell(w_0)}{\epsilon} \xrightarrow{\epsilon \rightarrow 0} \nabla \ell(w_0)$$

When ϵ is very small (almost 0), rc becomes the gradient of $\ell(w)$.

1D → N-dim parameters



- The 2D plots so far
 - parameter space on x-axis (1-dim param)
 - loss function on y-axis
- Actual neural network models have **many parameters** (like, thousands or millions or billions)
 - parameter space is **N-dim** (i.e. w is a N-dim array)
 - loss function $y=f(\text{input}, w)$

GD in Tensors

When working with tensors, GD is applied to each element of the tensor in the same way as before.

$$W_{iter} = W_{iter} - lr \cdot \nabla \ell(W_{iter})$$

To know the value of $l(W_{iter})$ we need to do a forward pass, so actually $l(W_{iter}) = l(W_{iter}, X)$ where X is the input data.

Stochastic Gradient Descent

We can't fit all the data in memory at once to compute the gradient. We use **batches of data** instead

Stochastic Gradient Descent (SGD) is **GD over batches**.

$$W_{iter} = W_{iter} - lr \cdot \nabla \ell(X_{iter}, W_{iter})$$


- where $l(X_{iter}, W_{iter})$ is the loss over one batch of data X_{iter}

Variants and Extensions of SGD

Many variants and extensions of SGD incorporating many tricks

- SGD with momentum
- RMSprop
- Adam, AdamW
- Adagrad


Learning Rate (aka step size)



$$W_{iter} = W_{iter} - lr \cdot \nabla \ell(X_{iter}, W_{iter})$$

- If **too big**, SGD might not converge (find the minimum)


- If **too small**, SGD will take a lot of iterations (and time) to converge. If number of iterations is not enough, it won't find the solution.
- Safe bet: **start with a bigish and a few iterations, observe a batch of validation data and tune it.**



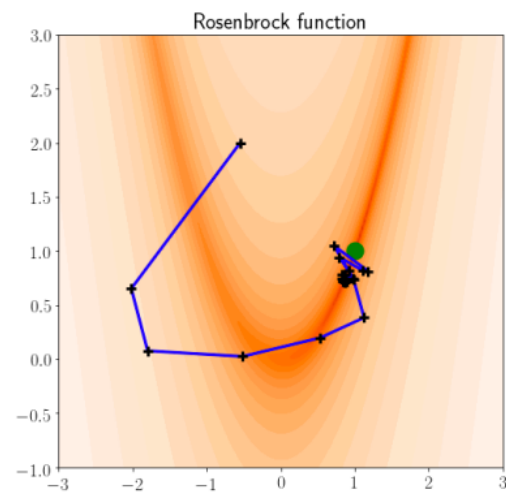
Momentum

In DNNs there are many “valleys” where SGD can get stuck, and momentum helps it to “keep going”.


$$W_{iter} = W_{iter} - lr \cdot \nabla \ell(X_{iter}, W_{iter}) + \beta(W_{iter} - W_{iter-1})$$




Gradient Descent



Gradient Descent with momentum



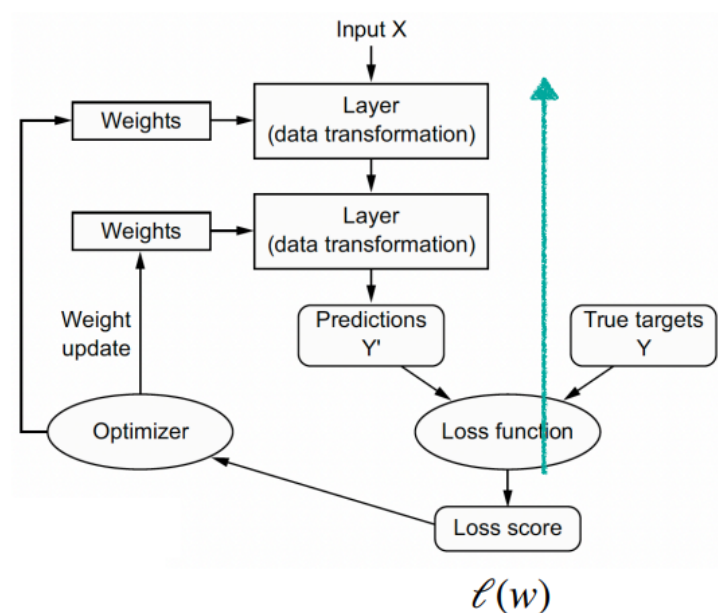
Local Minimum



- SGD is not infallible
- Momentum typically helps in this cases, but also not infallible
- Good news: local minimums in DNNs are not that bad

Backpropagation

Backpropagation is a clever trick to compute the gradient of the loss with respect to the parameters through multiple layers (=deep learning models!)



- **Main idea:** all the **gradients** of the tensor operations in the network **are known** in advance by TensorFlow, Keras, PyTorch. Backpropagation combines the gradients of each layer to form the gradient of the loss.

Takeaways

- SGD is very sensitive to the learning rate. Momentum often helps
- What is important when choosing an optimizer?
 - Choose an optimizer with momentum and some trick for adjusting the learning rate (e.g. Adam)
 - Choose a **big learning rate to start with** (e.g. 0.01 with Adam).

- If network doesn't converge, make it smaller.
- Bigger lr speeds things up.
- Always monitor the loss!.
- Choose **a small number of epochs to start** (e.g. 10) and a small part of the data.
 - If SGD (or variants) converges, augment the data and/or epochs until you reach the size of your problem.
 - If it doesn't converge, you need more epochs

4. Remember these terms

- Training loop
- Batches
- Optimizer
- Loss function
- Gradient
- Gradient Descent
- Stochastic Gradient Descent (SGD)
- Learning rate
- Number of iterations
- Momentum
- Back propagation
- Convergence
- Global minimum
- Local minimum

Training in a nutshell

```
def train():
    # typical torch training loop

    # define an optimizer
    optimizer = torch.optim.SGD(params=model.parameters(), lr=0.0001, momentum=0.2)

    for epoch_idx in range(num_epoch):
        for i, data in enumerate(training_loader):
            # Every data instance is an input + label pair
            inputs, labels = data

            # Make predictions for this batch
            preds = model(inputs)

            # Zero your gradients for every batch! otherwise, it remembers the previous gradients
            optimizer.zero_grad()

            # Compute the loss and its gradients
            loss = loss_fn(preds, labels)
            loss.backward() # backpropagation.
            # after this, we know the gradients of all the parameters.
            # in Torch/TF, the gradient values are stored in each parameter

            # Update the weights (by little bit (=by learning rate, 0.0001))
            # internally, the optimizer knows how to perform SGD. so..
            optimizer.step()
            # now the model.parameters() are a little better than before.
```